

Security & Privacy

Data Poisoning: add a backdoor (through a physical key)
(add poison training point)

Backdoor with trigger / triggerless

Expressed as a bilevel optimization problem:

$$X_p^* = \underset{X_p}{\operatorname{argmin}} L_{\text{adv}}(X_+, Y_{\text{adv}}; \theta^*(X_p))$$

X_p : poisoned data that we add

L_{adv} : how well we do at attacking our targets X_+

$$\theta^*(X_p) = \underset{\theta}{\operatorname{argmin}} L_{\text{train}}(X_c \cup X_p, Y; \theta)$$

Approximating solutions to bilevel optimization problems

"Metapoison" attack:

unroll stochastic gradient descent updates

$$\theta_1 = \theta_0 - \alpha \nabla_{\theta} L_{\text{train}}(X_c \cup X_p, Y; \theta_0)$$

$$\theta_2 = \theta_1 - \alpha \nabla_{\theta} \mathcal{L}_{\text{train}}(X_c \cup X_p, Y; \theta_1)$$

$$X_p^{i+1} = X_p^i - \beta \nabla_{X_p} \mathcal{L}_{\text{adv}}(X_t, Y_{\text{adv}}; \theta_2)$$

θ : a differentiable function of X_p

and we can take gradients

Approximating solutions to bilevel opt problems

How can we solve this?

Idea: instead of the argmin, write down the gradient descent updates and 'unroll' stochastic gradient descent updates.

$$\begin{aligned}\theta_1 &= \theta_0 - \alpha \nabla_{\theta} \mathcal{L}_{\text{train}}(X_c \cup X_p, Y; \theta_0) \\ \theta_2 &= \theta_1 - \alpha \nabla_{\theta} \mathcal{L}_{\text{train}}(X_c \cup X_p, Y; \theta_1) \\ X_p^{i+1} &= X_p^i - \beta \nabla_{X_p} \mathcal{L}_{\text{adv}}(x_t, y_{\text{adv}}; \theta_2),\end{aligned}$$

Now θ is a (differentiable) function of X_p and we can take gradients.

This is called the "Metapoison" attack

[Huang+ 2020]

Poison Type:

{ No Overlap
With Overlap

Aside: What's the state of empirical results in data poisoning? (vision)

Data poisoning is actually pretty brittle: what breaks data poisoning attacks

- Data augmentation / changing to SGD / transfer / ResNets
- Constraining for imperceptibility via l_∞
- Black box attacks
- Flipping the target image

Attack	CIFAR-10			TinyImageNet		
	Transfer		From Scratch	Transfer		From Scratch
	WB	BB		WB	BB	
FC	22.0	7.0	1.33	49.0	2.0	4.0
CP	33.0	7.0	0.67	14.0	1.0	0.0
BP	85.0	8.5	2.33	100.0	10.5	44.0
WiB	-	-	26.0	-	-	32.0
CLBD	5.0	6.5	1.00	3.0	1.0	0.0
HTBD	10.0	9.5	2.67	3.0	0.5	0.0

Attacks are viable, but not as good as we had seen

[Schwarzchild+ 2020]

Provable methods for data poisoning mitigation

$$P = (1 - \epsilon) P_{\text{clean}} + \epsilon Q$$

Data poisoning \Leftrightarrow

An adversary arrives and adds samples from an arbitrary distribution Q with the number of samples up to ϵ times the clean dataset

Recap and future threats

Practical, easy poisoning attacks exist for downstream, fine-tuned models

Metapoisn style attacks work for fine-tuned models

Defenses (provable and otherwise) are still an open problem

Data poisoning LMs – not yet seen, but likely in the future

LMs : privacy risk

Aggregation: combine multiple, public sources of information

Accessibilty: make sensitive, public information more available

Privacy Attacks!

{ memorization of public facts
aggregation



provable guarantees


Differential Privacy:

A formal privacy guarantee for a randomized algorithm



Differential privacy with deep learning (DP-SGD)

Q: How can we apply this to deep neural networks?

SGD:




Compute gradients





Sum and update



Differentially private SGD



Compute gradients



Clipping



Sum, noise and update

Mixed results for DP w/ deep neural nets in NLP

Prior attempts to apply DP to large neural models in NLP (via DPSGD) have often failed.

Example: Kerrigan et al – trained language generation models on reddit data

Input: “Bob lives close to the..”

Non-private outputs: “station and we only have two miles of travel left to go”

Private output ($\epsilon = 100$): “along supply am certain like alone before decent exceeding”

Why did things fail? (The dimensionality hypothesis)

1. Large language models have ~ 300 million parameters. That is *a lot* of things to privatize
2. Theory says differential privacy performance should degrade with dimension \sqrt{d}/n
3. Most (if not all) successful DP methods relied on low-dimensional statistics.


Differential privacy with large language models

Training large language models from scratch with DP

Open problem – large model size poses statistical + computational issues

Using a public language model to build a private downstream model

This is possible!

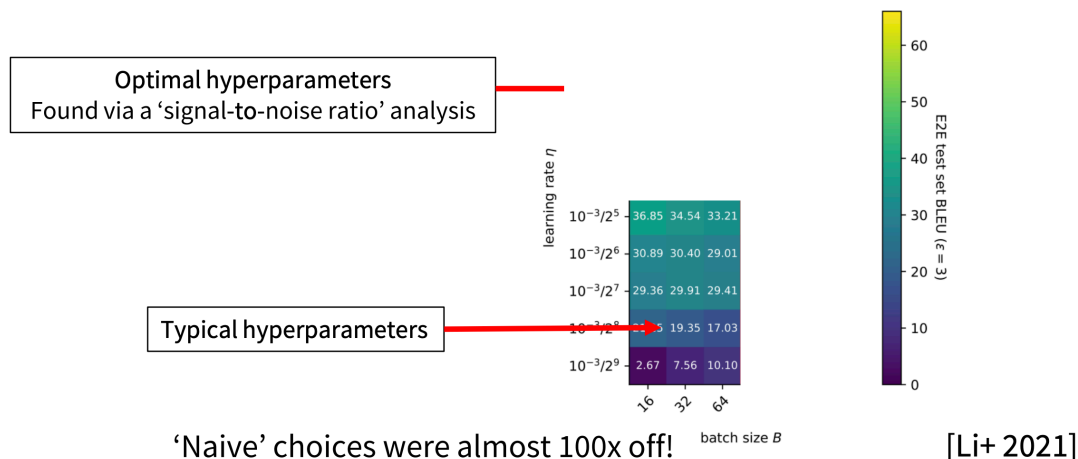


Fine if tuned right

Language model performance – fine if tuned right

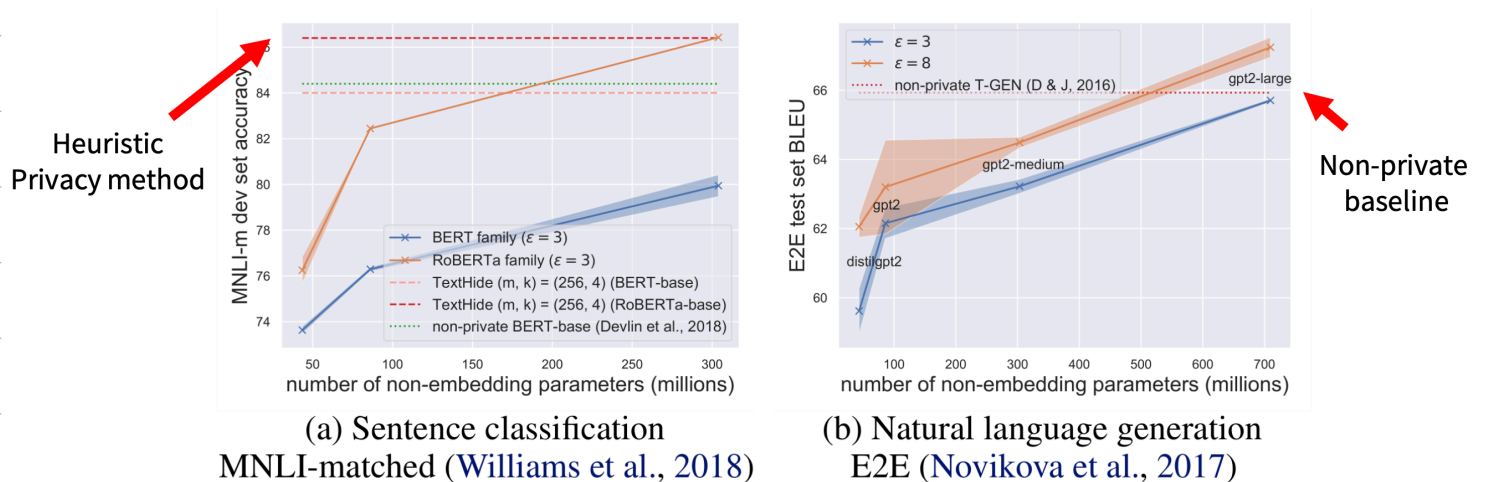
Identifying the problem: using *non-private* hyperparameters for *private* optimization

Solution: a way of predicting DP-SGD performance via ‘signal-to-noise’ ratios



Bigger models are better private learners

DP-SGD (which people ruled out) beats nonprivate baselines + heuristic privacy notions



Privacy:

Pre-trained, large language models are key to privacy

In the non-private case, pre-training is a small gain (5 BLEU points on E2E)

Metric	DP Guarantee	Gaussian DP + CLT	Compose tradeoff func.	full	LoRA	Method prefix	RGP	top2	retrain
BLEU	$\epsilon = 3$	$\epsilon \approx 2.68$	$\epsilon \approx 2.75$	61.519	58.153	47.772	58.482	25.920	15.457
	$\epsilon = 8$	$\epsilon \approx 6.77$	$\epsilon \approx 7.27$	63.189	63.389	49.263	58.455	26.885	24.247
	non-private	-	-	69.463	69.682	68.845	68.328	65.752	65.731
ROUGE-L	$\epsilon = 3$	$\epsilon \approx 2.68$	$\epsilon \approx 2.75$	65.670	65.773	58.964	65.560	44.536	35.240
	$\epsilon = 8$	$\epsilon \approx 6.77$	$\epsilon \approx 7.27$	66.429	67.525	60.730	65.030	46.421	39.951
	non-private	-	-	71.359	71.709	70.805	68.844	68.704	68.751

For private learning, the difference is **huge**:

- unusable (15 BLEU) when trained from scratch
- usable (61.5 BLEU) when privately fine-tuning a base LM.