

A Survey on Visual Transformer

Kai Han¹, Yunhe Wang^{1*}, Hanting Chen^{1,2}, Xinghao Chen¹, Jianyuan Guo¹, Zhenhua Liu^{1,2}, Yehui Tang^{1,2},
An Xiao¹, Chunjing Xu¹, Yixing Xu¹, Zhaohui Yang^{1,2}, Yiman Zhang¹, Dacheng Tao^{3*}

¹Noah's Ark Lab, Huawei Technologies

²Peking University ³University of Sydney

{kai.han, yunhe.wang, xinghao.chen, jianyuan.guo, xiaoan1, xuchunjing, yixing.xu, zhangyiman1}@huawei.com

{htchen, liu-zh, yhtang, zhaohuiyang}@pku.edu.cn, dacheng.tao@sydney.edu.au

Abstract

Transformer is a type of deep neural network mainly based on self-attention mechanism which is originally applied in natural language processing field. Inspired by the strong representation ability of transformer, researchers propose to extend transformer for computer vision tasks. Transformer-based models show competitive and even better performance on various visual benchmarks compared to other network types such as convolutional networks and recurrent networks. With high performance and without inductive bias defined by human, transformer is receiving more and more attention from the visual community. In this paper we provide a literature review of these visual transformer models by categorizing them in different tasks and analyze the advantages and disadvantages of these methods. In particular, the main categories include the basic image classification, high-level vision, low-level vision and video processing. The self-attention in computer vision is also briefly revisited as self-attention is the base component in transformer. Efficient transformer methods are included for pushing transformer into real applications on the devices. Finally, we give a discussion about the challenges and further research directions for visual transformers.

1. Introduction

Deep neural networks have become the fundamental infrastructure in modern artificial intelligence system. There have been various network types proposed for addressing different tasks. Multi-layer perception (MLP) or say fully connected (FC) network is the classical neural network which is stacked of multiple linear layers and nonlinear activations [125, 126]. Convolutional neural networks (CNNs) introduce convolutional layers and pooling layers for pro-

*Corresponding authors. All authors are in alphabetical order of last name (except the first and the corresponding authors).




Figure 1. Milestones of transformer. The visual transformer models are in red.

cessing shift-invariant data such as images [81, 78]. Recurrent neural networks (RNNs) utilize the recurrent cells to process sequential data or time series data [127, 59]. Transformer is a kind of newly proposed neural network which mainly utilizes self-attention mechanism [6, 108] to extract intrinsic feature [149]. Among these networks, transformer is a recently invented neural network and show great potential for extensive artificial intelligence applications.

Transformer is originally applied on natural language processing (NLP) tasks and brings in significant improvement [149, 34, 12]. For example, Vaswani *et al.* [149] first propose the transformer based solely on attention mechanisms for machine translation and English constituency parsing tasks. Devlin *et al.* [34] introduce a new language representation model called BERT, which pre-trains a transformer from unlabeled text by jointly conditioning on both left and right context. BERT obtains state-of-the-art results on eleven NLP tasks of the time. Brown *et al.* [12] pre-train a gigantic transformer based model GPT-3 with 175 billion parameters on 45TB compressed plaintext data and achieve strong performances on different types of downstream natural language tasks without fine-tuning. These Transformer based models show strong representation capacity and have obtained breakthrough in NLP area.

Inspired by the power of transformer in NLP, recently researchers extend transformer for computer vision (CV) tasks. CNNs used to be the fundamental component in vision applications [56, 123], but transformer is showing its ability as an alternative of CNN. Chen *et al.* [21] train a sequence transformer to auto-regressively predict pixels and

Table 1. Representative works of visual transformers.

Subject	Secondary Subject	Method	Keypoints	Publication
Image classification	Image classification	iGPT [21]	Pixel prediction self-supervised learning, GPT model	ICML 2020
		ViT [36]	Image patches, standard transformer	ICLR 2021
High-level vision	Object detection	DETR [16]	Set-based prediction, bipartite matching, transformer	ECCV 2020
		Deformable DETR [187]	DETR, deformable attention module	ICLR 2021
		ACT [184]	Adaptive clustering transformer	arXiv 2020
		UP-DETR [33]	Unsupervised pre-training, random query patch detection	arXiv 2020
		TSP [140]	New bipartite matching, encoder-only transformer	arXiv 2020
Low-level vision	Image enhancement	Max-DeepLab [152]	PQ-style bipartite matching, dual-path transformer	arXiv 2020
		VisTR [155]	Instance sequence matching, instance sequence segmentation	arXiv 2020
	Image generation	IPT [20]	Multi-task, ImageNet pre-training, transformer model	arXiv 2020
Video processing	Image generation	TTSR [163]	Texture transformer, RefSR	CVPR 2020
	Video inpainting	Image Transformer [110]	Pixel generation using transformer	ICML 2018
	Video Captioning	STTN [173]	Spatial-temporal adversarial loss	ECCV 2020
Efficient transformer	Decomposition	Masked Transformer [186]	Masking network, event proposal	CVPR 2018
	Distillation	ASH [101]	Number of heads, importance estimation	NeurIPS 2019
	Quantization	TinyBert [72]	Various losses for different modules	EMNLP Findings 2020
	Architecture Design	FullyQT [115]	Fully quantized transformer	EMNLP Findings 2020
		ConvBert [71]	Local dependence, dynamic convolution	NeurIPS 2020

achieve competitive results with CNNs on image classification task. ViT is a recently proposed vision transformer model by Dosovitskiy *et al.* [36] which applies a pure transformer directly to sequences of image patches and attains state-of-the-art performance on multiple image recognition benchmarks. Apart from the basic image classification, transformer has been utilized to address more computer vision problems such as object detection [16, 187], semantic segmentation, image processing and video understanding. Due to the excellent performance, more and more transformer-based models are proposed for improving various visual tasks.

The transformer-based vision models are springing up like mushrooms, which leads to difficulty to keep pace with the rate of new progress. Thus, a survey of the existing works is urgent and can be beneficial for the community. In this paper, we focus on providing a comprehensive overview of the recent advances in visual transformers and discuss about the potential directions for further improvement. To have a better archive and be convenient to researchers on different topics, we category the transformer models by their application scenarios as shown in Table 1. In particular, the main subjects include basic image classification, high-level vision, low-level vision and video processing. High-level vision deals with the interpretation and use of what is seen in the image [147] such as object detection, segmentation and lane detection. There have been a number of transformer models addressing these high-level vision tasks, such as DETR [16], deformable DETR [187] for object detection and Max-DeepLab [152] for segmentation. Low level image processing is mainly concerned with extracting descriptions from images (that are usually represented as images themselves) [42], whose typical applications include super-resolution, image denois-

ing and style transfer. Few works [20, 110] in low-level vision use transformers and more investigation is required. Video processing is an import part in computer vision in addition to image-based tasks. Transformer can be applied in video naturally [186, 173] due to the sequential property of video. Transformer is beginning to show competitive performance on these tasks compared to the conventional CNNs or RNNs. Here we give a survey on these works of transformer-based visual models to keep pace with the progress in this field. The development timeline of visual transformer is shown in Figure 1 and we believe more and more excellent works will be engraved in the milestones.

The rest of the paper is organized as follows. Section 2 first formulates the self-attention mechanism and the standard transformer. We describe the methods of transformer in NLP in Section 3 as the research experience may be beneficial for vision tasks. Next, Section 4 is the main part of the paper, in which we summarize the visual transformer models on image classification, high-level vision, low-level vision and video tasks. We also briefly revisit self-attention mechanism for CV and efficient transformer methods as they are closely related to our main topic. Finally, we give a conclusion and discuss about several research directions and challenges.

2. Formulation of Transformer

Transformer [149] was firstly applied on the machine translation task in neural language processing (NLP). As shown in Fig. 2, it consists of an encoder module and a decoder module with several encoders/decoders of the same architecture. Each encoder is composed of a self-attention layer and a feed-forward neural network, while each decoder is composed of a self-attention layer, an encoder-

decoder attention layer and a feed-forward neural network. Before translating sentences with transformer, each word in the sentence will be embedded into a vector with $d_{model} = 512$ dimensions.

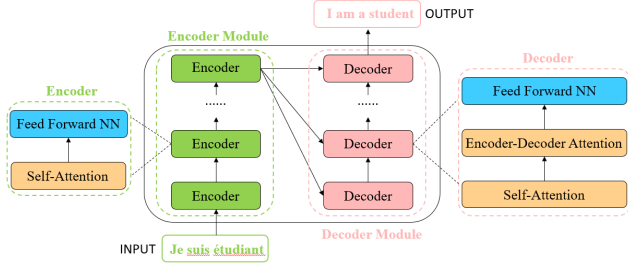


Figure 2. Pipeline of vanilla transformer.

2.1. Self-Attention Layer

In self-attention layer, the input vector is firstly transformed into three different vectors, *i.e.*, the query vector q , the key vector k and the value vector v with dimension $d_q = d_k = d_v = d_{model} = 512$. Vectors derived from different inputs are then packed together into three different matrices Q , K and V . After that, the attention function between different input vectors is calculated with the following steps (as shown in Fig. 3 left):

- **Step 1:** Compute scores between different input vectors with $S = Q \cdot K^T$;
- **Step 2:** Normalize the scores for the stability of gradient with $S_n = S / \sqrt{d_k}$;
- **Step 3:** Translate the scores into probabilities with softmax function $P = \text{softmax}(S_n)$;
- **Step 4:** Get the weighted value matrix with $Z = V \cdot P$.

The process can be unified into a single function:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{Q \cdot K^T}{\sqrt{d_k}}\right) \cdot V. \quad (1)$$

The intuition behind Eq. 1 is simple. Step 1 calculates scores between two different vectors, and the score is to determine the degree of attention that we put on other words when encode the word at the current position. Step 2 normalize the scores to have more stable gradients for better training, and step 3 transfer the scores into probabilities. Finally, each value vector is multiplied by the sum-upped probability, and vectors with larger probabilities will be focused on more by the following layers.

The encoder-decoder attention layer in the decoder module is almost the same as self-attention layer in the encoder-module, except that the key matrix K and value matrix V

are derived from the encoder module and the query matrix Q is derived from the previous layer.

Note that the above process is irrelevant to the position of each word, thus the self-attention layer lacks the ability of capturing the positional information of the words in a sentence. To address this, a positional encoding with dimension d_{model} is added to the original input embedding to get the final input vector of the word. Specifically, the position is encoded with the following equation:

$$PE(pos, 2i) = \sin\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right); \quad (2)$$

$$PE(pos, 2i + 1) = \cos\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right), \quad (3)$$

in which pos denotes the position of the word in a sentence, and i represents the current dimension of the positional encoding.




Figure 3. (Left) The process of self-attention. (Right) Multi-head attention. The image is from [149].

2.2. Multi-Head Attention

The self-attention layer is further improved by adding a mechanism called multi-head attention in order to boost the performance of the vanilla self-attention layer. Note that for a given reference word, we often want to focus on several other words when going through the sentence. Thus, a single-head self-attention layer limits the ability of focusing on a specific position (or several specific positions) while does not influence the attention on other positions that is equally important at the same time. This is achieved by giving attention layers different representation subspace. Specifically, different query, key and value matrices are used for different heads, and they can project the input vectors into different representation subspace after training due to the random initialization.

In detail, given an input vector and the number of heads h , the input vector is firstly transformed into three different groups of vectors, *i.e.*, the query group, the key group and the value group. There are h vectors in each group with

dimension $d_{q'} = d_{k'} = d_{v'} = d_{model}/h = 64$. Then, vectors derived from different inputs are packed together into three different groups of matrices $\{Q_i\}_{i=1}^h$, $\{K_i\}_{i=1}^h$ and $\{V_i\}_{i=1}^h$. Then, the process of multi-head attention is shown as follow:

$$MultiHead(Q', K', V') = Concat(head_1, \dots, head_h)W^o, \\ \text{where } head_i = Attention(Q_i, K_i, V_i), \quad (4)$$

where Q' is the concatenation of $\{Q_i\}_{i=1}^h$ (so is K' and V') and $W^o \in \mathbb{R}^{d_{model} \times d_{model}}$ is the linear projection matrix.

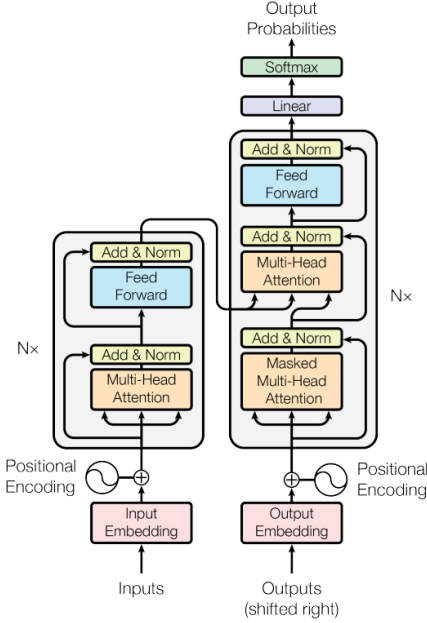


Figure 4. Detailed structure of transformer (the image is from [149]).

2.3. Other Parts in Transformer

Residual in the encoder and decoder. As shown in Fig. 4, a residual connection is added in each sub-layer in the encoder and decoder in order to strengthen the flow of information and get a better performance. A layer-normalization [5] is followed afterward. The output of operations mentioned above can be described as:

$$LayerNorm(X + Attention(X)). \quad (5)$$

Note that X is used as the input of self-attention layer here, since the query, key and value matrices Q , K and V are all derived from the same input matrix X .

Feed-forward neural network. A feed-forward NN is applied after the self-attention layers in each encoder and decoder. Specifically, the feed-forward NN is consist of two linear transformation layers and a nonlinear activation function within them, which can be denoted as the following function:

$$FFNN(X) = W_2\sigma(W_1X), \quad (6)$$

where W_1 and W_2 are the two parameter matrices of the two linear transformation layers, and σ represents the nonlinear activation function, such as GELU [57]. The dimensionality of the hidden layer is $d_h = 2048$.

Final layer in decoder. The final layer in decoder aims to turn the stack of vectors back into a word. This is achieved by a linear layer followed by a softmax layer. The linear layer project the vector into a logits vector with d_{word} dimensions in which d_{word} is the number of words in the vocabulary. Then, a softmax layer is used to transform the logits vector into probabilities.

Most of the transformers used in the computer vision tasks utilize the encoder module of the original transformer. In short, it can be treated as a new feature selector that is different from the convolutional neural networks (CNNs) and the recurrent neural networks (RNNs). Compared to CNN which only focuses on local characteristics, transformer is able to capture long distance characteristics which means that the global information can be easily derived by transformer. Compared to RNN whose hidden state must be computed in sequence, transformer is much efficient since the output of the self-attention layer and the fully-connected layers can be computed in parallel and accelerated easily. Thus, it is meaningful to further study the application of transformer in the area of not only NLP but also computer vision.

3. Revisiting Transformers for NLP

Before the advent of Transformer, recurrent neural networks (e.g., GRU [31] and LSTM [60]) with added attention empower most of the state-of-the-art language models. However, in RNNs, the information flow needs to be processed sequentially from the previous hidden states to the next one, which precludes the acceleration and parallelization during training, and thus hinders the potential of RNNs to process longer sequences or build larger models. In 2017, Vaswani *et al.* [149] proposes Transformer, a novel encoder-decoder architecture solely built on multi-head self-attention mechanisms and feed-forward neural networks, aiming to solve seq-to-seq natural language tasks (e.g., machine translation) with acquiring global dependencies at ease. The success of Transformer demonstrates that leveraging attention mechanisms alone can achieve comparable performances with attentive RNNs. Moreover, the architecture of Transformer favors massively parallel computing, enabling training on larger datasets and thus leading to the burst of large pre-trained models (PTM) for natural language processing.

BERT [34] and its variants (e.g., SpanBERT [74], RoBERTa [97]) are a series of PTMs built on the multi-layer Transformer encoder architecture. Two tasks are conducted on BookCorpus [188] and English Wikipedia datasets at the pre-training stage of BERT: 1) Masked language modeling

(MLM) via first randomly masking out some tokens in the input and then train the model to predict; 2) Next sentence prediction using paired sentences as input and predicting whether the second sentence is the original one in the document. After pre-training, BERT can be fine-tuned by adding one output layer alone on a wide range of downstream tasks. More specifically, When performing sequence level tasks (e.g., sentiment analysis), BERT uses the representation of the first token for classification; while for token-level tasks (e.g., Name Entity Recognition), all tokens are fed into the softmax layer for classification. At the time of release, BERT achieves the state-of-the-art results on 11 natural language processing task, setting up a milestone in pre-trained language models. Generative Pre-Trained Transformer series (e.g., GPT [119], GPT-2 [120]) are another type of pre-trained models based on the Transformer decoder architecture, which uses masked self-attention mechanisms. The major difference between GPT series and BERT lies in the way of pre-training. Unlike BERT, GPT series are one-directional language models pre-trained by Left-to-Right (LTR) language modeling. Besides, sentence separator ([SEP]) and classifier token ([CLS]) are only involved in the fine-tuning stage of GPT but BERT learns those embeddings during pre-training. Because the one-directional pre-training strategy of GPT, it shows superiority in many natural language generation tasks. More recently, a gigantic transformer-based model, GPT-3, with incredibly 175 billion parameters has been introduced [12]. By pre-training on 45TB compressed plaintext data, GPT-3 claims the ability to directly process different types of downstream natural language tasks without fine-tuning, achieving strong performances on many NLP datasets, including both natural language understanding and generation. Besides the aforementioned transformer-based PTMs, many other models have been proposed since the introduction of Transformer. For this is not the major topic in our survey, we simply list a few representative models in Table 2 for interested readers.

Apart from the PTMs trained on large corpora for general natural language processing tasks, transformer-based models have been applied in many other NLP related domains or multi-modal tasks. **BioNLP Domain.** Transformer-based models have outperformed many traditional biomedical methods. BioBERT [82] uses Transformer architecture for biomedical text mining tasks; SciBERT [9] is developed by training Transformer on 114M scientific articles covering biomedical and computer science field, aiming to execute NLP tasks related to scientific domain more precisely; Huang *et al.* [65] proposes ClinicalBERT utilizing Transformer to develop and evaluate continuous representations of clinical notes and as a side effect, the attention map of ClinicalBERT can be used to explain predictions and thus discover high-quality connections between different medical contents. **Multi-Modal Tasks.** Ow-

Table 2. List of representative language models built on Transformer.

Models	Architecture	Params	Fine-tuning
GPT [119]	Transformer Dec.	117M	Yes
GPT-2 [120]	Transformer Dec.	117M~1542M	No
GPT-3 [12]	Transformer Dec.	125M~175B	No
BERT [34]	Transformer Enc.	110M~340M	Yes
RoBERTa [97]	Transformer Enc.	355M	Yes
XLNet [164]	Two-Stream Transformer Enc.	≈ BERT	Yes
ELECTRA [32]	Transformer Enc.	335M	Yes
UniLM [35]	Transformer Enc.	340M	Yes
BART [83]	Transformer	110% of BERT	Yes
T5 [121]	Transformer	220M~11B	Yes
ERNIE (THU) [180]	Transformer Enc.	114M	Yes
KnowBERT [112]	Transformer Enc.	253M~523M	Yes

¹ Transformer is the standard encoder-decoder architecture. Transformer Enc. and Dec. represent the encoder and decoder part of standard Transformer. Decoder uses mask self-attention to prevent attending to the future tokens.

² The data of the Table is from [118].

ing to the success of Transformer across text-based NLP tasks, many researches are committed to exploiting the potential of Transformer to process multi-modal tasks (e.g., video-text, image-text and audio-text). VideoBERT [138] uses a CNN-based module pre-processing the video to get the representation tokens, based on which a Transformer encoder is trained to learn the video-text representations for downstream tasks, such as video caption. VisualBERT [85] and VL-BERT [137] propose single-stream unified Transformer to capture visual elements and image-text relationship for downstream tasks like visual question answering (VQA) and visual commonsense reasoning (VCR). Moreover, several studies such as SpeechBERT [29] explore the possibility of encoding audio and text pairs with a Transformer encoder to process auto-text tasks like Speech Question Answering (SQA).

The rapid development of transformer-based models on varieties of natural language processing as well as NLP-related tasks demonstrates its structural superiority and versatility. This empowers Transformer to become a universal module in many other AI fields beyond natural language processing. The following part of this survey will focus on the applications of Transformer in a wide range of computer vision tasks emerged in the past two years.

4. Visual Transformer

In this section, we provide a comprehensive review of the transformer-based models in computer vision, including the applications in image classification, high-level vision, low-level vision and video processing. We also briefly summarize the applications of self-attention mechanism and model compression methods for efficient transformer.

4.1. Image Classification

Inspired by great success of transformer on natural language processing, some researchers have tried to examine whether similar models can learn useful representations for images. As a higher dimensional, noiser, and more redundant modality than text, images are believed to be difficult for generative modeling.

Wu et al. [159] use ResNet as a convenient baseline and use visual transformers to replace the last stage of convolutions. Specifically, they apply convolutional layers to extract low-level features and fed to the visual transformer. For the visual transformer, they apply a *tokenizer* to group pixels into a small number of *visual tokens*, each representing a semantic concept in the image. Then the transformers are utilized to model relationships between tokens and the *visual tokens* are directly used for image classification. iGPT [21], ViT [36] and DeiT [145] are the works purely using transformer for image classification.

4.1.1 iGPT

It has been a long time since the original wave of generative pre-training methods for images, Chen *et al.* [21] re-examine this class of methods and combine with the recent progress of self-supervised methods. The approach consists of a pre-training stage followed by a fine-tuning stage. In pre-training, auto-regressive and BERT objectives are explored. Besides, sequence transformer architecture is applied to predict pixels instead of language tokens in NLP. Pre-training can be viewed as a favorable initialization or as a regularizer when used in combination with early stopping. During fine-tuning, they adds a small classification head to the model, which is used to optimize a classification objective and adapts all weights.

Given an unlabeled dataset X consisting of high dimensional data $x = (x_1, \dots, x_n)$. They train the model by minimizing the negative log-likelihood of the data:

$$L_{AR} = \mathbb{E}_{x \sim X} [-\log p(x)] \quad (7)$$

where $p(x)$ is the density of the data of images, which can be modeled as:

$$p(x) = \prod_{i=1}^n p(x_{\pi_i} | x_{\pi_1}, \dots, x_{\pi_{i-1}}, \theta) \quad (8)$$

where the identity permutation $\pi_i = i$ is adopted for $1 \leq i \leq n$, also known as raster order. They also consider the BERT objective, which samples a sub-sequence $M \subset [1, n]$ such that each index i independently has probability 0.15 of appearing in M . The M is called the BERT mask, and the model is trained by minimizing the negative log-likelihood of the "masked" elements x_M conditioned

on the "unmasked" ones $x_{[1,n] \setminus M}$:

$$L_{BERT} = \mathbb{E}_{x \sim X} \mathbb{E}_M \sum_{i \in M} [-\log p(x_i | x_{[1,n] \setminus M})] \quad (9)$$

In pre-training, they pick one of L_{AR} or L_{BERT} and minimize the loss over the pre-training dataset.

They use the GPT-2 [120] formulation of the transformer decoder block. In particular, layer norms precede both the attention and Multi-Layer Perceptron (MLP) operations, and all operations lie strictly on residual paths. The only mixing across sequence elements occurs in the attention operation, and to ensure proper conditioning when training the AR objective, they apply the standard upper triangular mask to the $n \times n$ matrix of attention logits. When using the BERT objective, no attention logit masking is required: after applying content embeddings to the input sequence, they zero out the positions.

Following the final transformer layer, they apply a layer norm and learn a projection from the output to logits parameterizing the conditional distributions at each sequence element. When training BERT, they simply ignore the logits at unmasked positions.

During fine-tuning, they average pool the output of the final layer norm n^L across the sequence dimension to extract a d -dimensional vector of features per example:

$$f^L = \langle n_i^L \rangle_i \quad (10)$$

They learn a projection from f^L to class logits, which is utilized to minimize a cross entropy loss L_{CLF} . In practice, they find empirically that the joint objective $L_{GEN} + L_{CLF}$ works even better, where $L_{GEN} \in \{L_{AR}, L_{BERT}\}$.

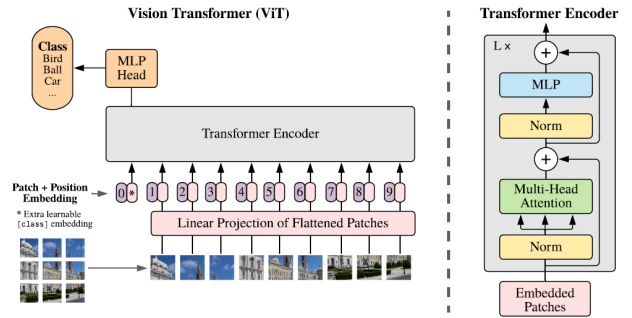


Figure 5. The framework of the Vision Transformer (The image is from [36]).

4.1.2 ViT

Recently, Dosovitskiy et al. [36] propose a pure transformer, namely Vision Transformer (ViT), which performs well on image classification tasks when applied directly to sequences of image patches. They follow as closely as possible the design of the original Transformer. Figure 5 shows the framework of Vision Transformer.

To handle 2D images, the image $x \in \mathbb{R}^{H \times W \times C}$ is reshaped into a sequence of flattened 2D patches $x_p \in \mathbb{R}^{N \times (P^2 \cdot C)}$. (H, W) is the resolution of the original image and (P, P) is the resolution of each image patch. $N = HW/P^2$ is then the effective sequence length for the transformer. Since the transformer uses constant widths through all of its layers, a trainable linear projection maps each vectorized patch to the model dimension D , the output of which they refer to as patch embeddings.

Similar to BERT’s `[class]` token, a learnable embedding is employed to the sequence of embedding patches, whose state at the output of the transformer encoder serves as the image representation. During both pre-training and fine-tuning, the classification head is attached to the same size. In addition, 1D position embeddings are added to the patch embeddings to retain positional information. They have explored different 2D-aware variants of position embeddings, which does not obtain significant gains over standard 1D position embeddings. The joint embeddings serves as input to the encoder. It is worth noting that the Vision Transformer only employs the encoder of the standard transformer and a MLP head is followed by the output of the transformer encoder.

Typically, the ViT is firstly pre-trained on large datasets, and fine-tune to smaller downstream tasks. For this, the pre-trained prediction head is removed and a zero-initialized $D \times K$ feedforward layer is attached, where K is the number of downstream classes. It is often beneficial to fine-tune at higher resolution than pre-training. When feeding images of higher resolution, the patch size is kept the same, which results in a larger effective sequence length. The Vision Transformer can handle arbitrary sequence lengths, however, the pre-trained position embeddings may no longer be meaningful. The authors therefore perform 2D interpolation of the pre-trained position embeddings according to their location in the original image. Note that this resolution adjustment and patch extraction are the only points at which an inductive bias about the 2D structure of the images is manually injected into the Vision Transformer.

Such models yield modest results when trained on mid-sized datasets such as ImageNet, achieving accuracies of a few percentage points below ResNets of comparable size. Transformers lack some inductive biases inherent to CNNs, such as translation equivariance and locality, and therefore do not generalize well when trained on insufficient amounts of data. However, the picture changes if the models are trained on large datasets (14M-300M images). The authors find that large scale training trumps inductive bias. Transformers attain excellent results when pre-trained at sufficient scale and transferred to tasks with fewer datapoints. The Vision Transformer, pretrained on the JFT-300M dataset, approaches or beats state of the art on multiple image recognition benchmarks, reaching accuracy of

88.36% on ImageNet, 99.50% on CIFAR-10, 94.55% on CIFAR-100, and 77.16% on the VTAB suite of 19 tasks. The detail results of iGPT and ViT are shown in table 3.

Touvron et al. [145] propose a competitive convolution-free transformer by training on Imagenet only. DeiT-B has the same architecture as ViT-Base which has 86M parameters. With a strong data augmentation, the reference vision transformer (DeiT-B) achieves top-1 accuracy of 83.1% (single-crop evaluation) on ImageNet with no external data. In addition, the authors observe that using a convnet teacher gives better performance than using a transformer. Specifically, DeiT-B can achieve 84.40% top-1 accuracy with the help of distillation token.

In conclusion, iGPT recalls the generative pre-training method and combine it with self-supervised methods, the results is not very satisfied. ViT has achieves much better results especially when it utilizes a larger dataset (JFT-300). DeiT performs better with a more cautious training strategy and distillation. However, the structure of ViT is basically the same as the transformer in NLP, how to explicit the correlations in the intra-patch and inter-patch is still a challenging problem. Besides, the patches with the same size are treated equal in ViT. As we all know, the complexity of each patch is different and this characteristic has not been fully employed now.

4.2. High-level Vision

Recently there have been increasing interests for adopting transformer for high-level computer vision tasks, such as object detection [17, 187, 8, 168], lane detection [96], segmentation [155, 152] and pose estimation [66, 67, 91]. In this section we make a review for these methods.

4.2.1 Generic Object Detection

Traditional object detectors are mainly built upon CNNs and recently using transformer for object detection has attracted increasing interests due to the good capability of transformer.

Some object detection methods have tried to utilize the self-attention mechanism of transformer to enhance the specific modules for modern detectors, like feature fusion module [175] and prediction head [27], which will be further discussed in Section 4.6.

Transformer-based object detection methods can be coarsely categorized into two groups, *i.e.*, transformer-based set prediction methods [16, 187, 140, 184] and transformer-based backbone methods [8, 107], as shown in Fig. 6.

Transformer-based set prediction for detection. As an pioneer for transformer-based detector, the detection transformer (DETR) proposed by Carion [17] redesigns the framework of object detection, which is a simple and fully

Table 3. Experimental results comparison between CNN and transformers. [†]Result using knowledge distillation.

Method	Network	Pre-training policy	Pre-training data	Dataset	Params (M)	Acc (Top-1)
BiT-L [76]	CNN	Supervised	JFT-300M	CIFAR-10	60	99.37
				CIFAR-100	60	93.51
				ImageNet	60	87.54
iGPT [21]	Transformer	Self-supervised	ImageNet	CIFAR-10	1362	99.00
				CIFAR-100	1362	91.70
ViT [36]	Transformer	Supervised	ImageNet	CIFAR-10	86	98.13
				CIFAR-100	86	87.13
				ImageNet	86	77.91
			JFT-300M	CIFAR-10	86	99.00
				CIFAR-100	86	91.87
				ImageNet	86	84.15
				CIFAR-10	632	99.50
				CIFAR-100	632	94.55
				ImageNet	632	88.36
DeiT-B [145]	Transformer	Supervised	ImageNet	ImageNet	86	83.10
					86	84.40 [†]

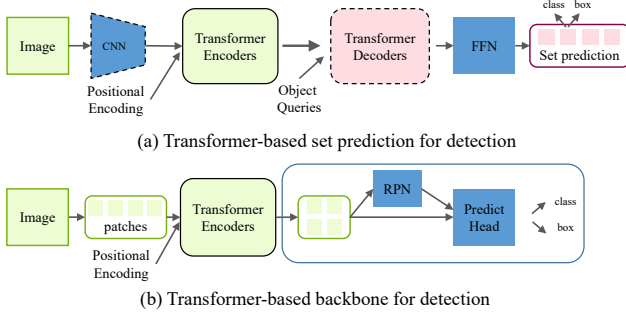


Figure 6. General framework of transformer-based object detection. (a) The object detection task is formulated as set prediction task via transformer, *e.g.*, DETR [16]. (b) The vision transformer is utilized to replace the CNN backbone for traditional object detectors, *e.g.*, ViT-FRCNN [8].

end-to-end object detector. DETR treats the object detection task as an intuitive set prediction problem and gets rid of traditional hand-crafted components like anchor generation and non-maximum suppression (NMS) post-processing. As shown in Fig. 7, DETR starts with an CNN backbone to extract features from the input image. To supplement the image features with position information, the fixed positional encodings are added to the flatten features before fed into the encode-decoder transformer. The transformer decoder consumes the embeddings from encoder along with N learnt positional encodings (object queries), and produces N output embeddings, where N is a predefined parameter and is typically larger than the number of objects in an image. The final predictions are computed with simple feed-forward networks (FFN), which include the bounding box coordinates and class labels to indicate the specific class of the object or no object. Unlike original

transformer that produce predictions sequentially, DETR decodes N objects in parallel at the same time. DETR employs an bipartite matching algorithm to assign the predicted and ground-truth objects. As shown in Eq. (11), the Hungarian loss is exploited to compute the loss function for all matched pairs of objects.

$$\mathcal{L}_{\text{Hungarian}}(y, \hat{y}) = \sum_{i=1}^N \left[-\log \hat{p}_{\hat{\sigma}(i)}(c_i) + \mathbb{1}_{\{c_i \neq \emptyset\}} \mathcal{L}_{\text{box}}(b_i, \hat{b}_{\hat{\sigma}(i)}) \right], \quad (11)$$

where y and \hat{y} are the ground truth and prediction of objects respectively, $\hat{\sigma}$ is the optimal assignment, c_i and $\hat{p}_{\hat{\sigma}(i)}(c_i)$ are the target class label and predicted label, b_i and $\hat{b}_{\hat{\sigma}(i)}$ are ground truth and predicted bounding box, respectively. DETR shows impressive performance on object detection, with comparable accuracy and speed with the popular and well-established Faster R-CNN baseline on COCO benchmark.




Figure 7. The overall architecture of DETR (figure from [17]).

DETR is a new design for object detection framework based on transformer and sheds light on the community to develop fully end-to-end detectors. However, the vanilla DETR also comes with several challenges, *e.g.*, longer training schedule and poor performance for small objects. Deformable DETR proposed by Zhu *et al.* [187] is a popular method for address the above issues and greatly improves the detection performance. Instead of looking

over all spatial locations on image feature maps by original multi-head attention in transformer, the deformable attention module is proposed to attend to a small set of key positions around a reference point. In this way, the computational complexity is greatly reduced and also benefit to fast convergence. More importantly, the deformable attention module can be easily applied for fusing multi-scale features. Deformable DETR achieves better performance than DETR with $10\times$ less training cost and $1.6\times$ faster inference speed. Some additional improvements are also applied for Deformable DETR, including an effective iterative bounding box refinement method and a two-stage scheme, which result in further performance gain.

Aiming at the high computation complexity problem of DETR, Zheng *et al.* [184] propose a Adaptive Clustering Transformer (ACT) to reduce the computation cost of the pre-trained DETR without any training process. ACT adaptively clusters the query features using a locality sensitivity hashing method and broadcast the attention output to the queries represented by the selected prototypes. By replacing the self-attention module of the pre-trained DETR model with the proposed ACT without any re-training, the computation cost can be considerably reduced with suffering little accuracy degradation. Additionally, the performance drop can be further reduced by utilizing a multi-task knowledge distillation (MTKD) method, which exploit the original transformer to distill the ACT module with a few epochs of fine-tuning.

Sun *et al.* [140] investigate the slow convergence problem of DETR model and reveal that the cross-attention module in the transformer decoder is the main reason behind. To this end, an encoder-only version of DETR is proposed and achieves considerably improvement for the detection accuracy and training convergence. Moreover, a new bipartite matching scheme is designed for more stable training and faster convergence. Two transformer-based set prediction models are proposed to improve encoder-only DETR with feature pyramids, namely TSP-FCOS and TSP-RCNN, which achieve better performance than the original DETR model.

Inspired by the pre-training transformer scheme in natural language processing, Dai *et al.* [33] propose a method for unsupervisedly pre-train DETR (UP-DETR) for object detection. Specifically, a novel unsupervised pretext task named random query patch detection is proposed to pre-train the DETR model. With this scheme, UP-DETR improve the detection accuracy on a relatively small dataset, *i.e.*, PASCAL VOC, by a large margin. On the COCO benchmark with sufficient training data, UP-DETR still outperform DETR, which demonstrate the effectiveness of unsupervised pre-training scheme.

Transformer-based backbone for detection. Unlike DETR that redesign object detection as a set prediction

tasks via transformer, Beal *et al.* [8] propose to utilize transformer as a backbone for common detection framework like Faster R-CNN [124]. The input image is divided into several patches and fed into a vision transformer, whose output embedding features are reorganized according to the spatial information and go through a detection head for the final results. The proposed ViT-FRCNN could naturally benefit from massive pre-training transformer backbone.

Detection results for different transformer-based object detectors mentioned above on COCO 2012 val set are listed in Table 4. Transformer-based methods have shown strong performance compared to CNN-based detectors, in terms of accuracy and running speed.

4.2.2 Other Detection Tasks

Pedestrian Detection. Considering that the distribution of objects is very dense in occluded or crowd scenes, additional analysis and adaptation are often required when applying common detection networks to pedestrian detection task. Lin *et al.* [92] reveal that sparse uniform queries and weak attention field in decoder harm the performance when directly applying DETR [16] or Deformable DETR [187] for pedestrian detection. To this end, they propose Pedestrian End-to-end Detector (PED) to alleviate the identified drawbacks. In PED, a new decoder named Dense Queries and Rectified Attention field (DQRF) is utilized to support dense queries and alleviate the noisy or narrow attention field of the queries. Moreover, they propose V-Match to fully leverage visible annotations for further improvements.

Lane Detection. Based on PolyLaneNet [142], Liu *et al.* [96] propose to improve performance of curve lane detection by learning the global context with transformer network. Similar to PolyLaneNet, the proposed method (LSTR) regards the lane detection as a task of fitting lanes with polynomials and uses neural networks to predict the parameters of polynomials. To capture long and thin structures for lanes and the global context, LSTR introduces transformer network into the architecture to process low-level features extracted by convolutional neural networks. Besides, LSTR uses Hungarian Loss to optimize network parameters. As is shown in [96], LSTR achieves 2.82% higher accuracy and $3.65\times$ FPS than PolyLaneNet with only $0.2\times$ parameters. The combination of transformer network, convolutional neural network and Hungarian Loss realizes a tiny, fast and precise lane detection framework.

4.2.3 Segmentation

DETR [17] can be naturally extended for panoptic segmentation task by appending a mask head on the decoder and achieve competitive results. Wang *et al.* [152] propose Max-DeepLab to directly predict panoptic segmentation results with a mask transformer, without surrogate sub-tasks

Table 4. Comparison of different transformer-based object detectors using ResNet-50 on COCO 2017 val set. Running speed (FPS) is evaluated on NVIDIA Tesla V100 GPU as reported in [187]. [†]Estimated speed according to the reported number in the paper. ^{*}ViT backbone is pretrained on ImageNet-21k. * ViT backbone is pretrained on an private dataset with 1.3B images.

Method	Epochs	AP	AP ₅₀	AP ₇₅	AP _s	AP _M	AP _L	#Params	FLOPs	FPS
<i>CNN based</i>										
FCOS [144]	36	41.0	59.8	44.1	26.2	44.6	52.2	-	177G	23 [†]
Faster R-CNN + FPN [124]	109	42.0	62.1	45.5	26.6	45.4	53.4	42M	180G	26
<i>Transformer based</i>										
DETR [16]	500	42.0	62.4	44.2	20.5	45.8	61.1	41M	86G	28
DETR-DC5 [16]	500	43.3	63.1	45.9	22.5	47.3	61.1	41M	187G	12
Deformable DETR [187]	50	46.2	65.2	50.0	28.8	49.2	61.7	40M	173G	19
TSP-FCOS [140]	36	43.1	62.3	47.0	26.6	46.8	55.9	-	189G	20 [†]
TSP-RCNN [140]	96	45.0	64.5	49.6	29.7	47.7	58.0	-	188G	15 [†]
ACT+MKKD (L=32) [184]	-	43.1	-	-	61.4	47.1	22.2	-	169G	14 [†]
ACT+MKKD (L=16) [184]	-	40.6	-	-	59.7	44.3	18.5	-	156G	16 [†]
ViT-B/16-FRCNN [*] [8]	21	36.6	56.3	39.3	17.4	40.0	55.5	-	-	-
ViT-B/16-FRCNN [*] [8]	21	37.8	57.4	40.1	17.8	41.4	57.3	-	-	-
UP-DETR [33]	150	40.5	60.8	42.6	19.0	44.4	60.0	41M	-	-
UP-DETR [33]	300	42.8	63.0	45.3	20.8	47.1	61.7	41M	-	-

like box detection. Similar with DETR, Max-DeepLab streamlines the panoptic segmentation tasks in an end-to-end fashion and directly predicts a set of non-overlapping masks and corresponding labels. A panoptic quality (PQ) style loss is utilized for training the model. Additionally, unlike prior methods that stack transformer on top of a CNN backbone, Max-DeepLab employs a dual-path framework for better combining CNN with transformer.

Wang *et al.* [155] propose a transformer-based video instance segmentation (VisTR) model, which takes a sequence of images as input and produces corresponding instance prediction results. An instance sequence matching strategy is proposed to assign the predictions with ground truth. To obtain the mask sequence for each instance, VisTR utilizes the instance sequence segmentation module to accumulate the mask features from multiple frames and segment the mask sequence with 3D CNN.

There is also an attempt to use transformer for cell instance segmentation [113], which is based on the DETR panoptic segmentation model. The proposed Cell-DETR additionally adds skip connections to bridge feature from backbone CNN and the CNN decoder in segmentation head to obtain better fused features. Cell-DETR shows state-of-the-art performance for cell instance segmentation from microscopy imagery.

Recently there are several works to explore transformer architecture for point cloud learning [38, 50, 181]. Guo *et al.* [50] propose to replace the original self-attention with the more suitable offset-attention module for point cloud learning, which includes implicit Laplace operator and normalization refinement. Zhao *et al.* [181] design a novel transformer architecture (Point Transformer) for processing

point clouds. The proposed self-attention layer is invariant to the permutation of the point set and thus is suitable for the point set processing tasks. Point Transformer shows strong performance for semantic segmentation task from 3D point cloud.

4.2.4 Pose Estimation

Human and hand pose estimation aim to predict the joint coordinates or mesh vertices from input RGB/D images and have attracted tremendous interest from the research community. Articulated pose estimation is basically a structured prediction task and some methods [66, 67, 91] explore to utilize transformer to model the global structure information of human pose.

Huang *et al.* [66] propose a transformer based network for 3D hand pose estimation from point sets. The encoder first utilizes a PointNet [116] to extract point-wise features from input point cloud and further adopts standard multi-head self-attention module to produce embeddings. To expose more global pose-related information to the decoder, a feature extractor like PointNet++ [117] is adopted to extract hand joint-wise features, which are fed into the decoder as positional encodings.

Similarly, Huang *et al.* [67] propose hand-object transformer network (HOT-Net) for 3D hand-object pose estimation. Unlike the above method that utilizing transformer to directly predict 3D hand pose from input point clouds, HOT-Net instead uses a ResNet to generate initial 2D hand-object pose and feeds it into a transformer to predict the 3D hand-object pose. Therefore, a spectral graph convolution network is instead exploited to extract input embeddings for

encoder.

Lin *et al.* [91] propose a mesh transformer (METRO) to predict 3D human pose and mesh from a single RGB image. METRO extracts image features via a CNN and perform position encoding by concatenating a template human mesh to the image feature. A progressive dimensionality reduction transformer encoder is proposed to gradually reduce the embedding dimensions and finally produce 3D coordinates of human joint and mesh vertex. To encourage the learning of non-local relationships between human joints, METRO randomly mask some input queries during training.

4.2.5 Discussions

As discussed above, transformers have shown strong performance on several high-level tasks, including detection, segmentation and pose estimation. The most important issues to handle for adopting transformer on high-level tasks include the input embedding, position encoding and the prediction loss. Some methods also proposed to improve the self-attention module from different perspectives, *e.g.*, deformable attention [187], adaptive clustering [184] and point transformer [181]. Nevertheless, transformers for high-level vision tasks have just been preliminarily explored and more research topics are worthy for digging. For example, are the feature extraction modules like CNN or PointNet before transformer necessary for the final performance? How to make vision transformer beneficial from large scale pre-training data like BERT or GPT-3 do in NLP community. Is it possible to pre-train a single transformer model and fine-tune for different downstream tasks with a few epochs of fine-tuning? We hope that more further research effort could be dedicated for exploring more powerful transformers for high-level vision.

4.3. Low-level Vision

Besides high-level vision tasks, few works apply transformers on low-level vision fields, such as image super-resolution, generation, *etc.* Compared with classification, segmentation and detection whose outputs are labels or boxes, the low-level tasks often take images as outputs (*e.g.*, high-resolution or denoised images), which is more challenging.

Parmar *et al.* [110] take the first step at generalizing the transformer model to formulate image translation and generation task and propose Image Transformer. The Image Transformer consists of two parts: an encoder for extracting image representation and a decoder to generate pixels. For each pixel with value 0 – 255, a $256 \times d$ dimensional embedding is learned for encoding each value into a d dimensional vector, which is taken as the input of the encoder. The architecture of encoder and decoder are the same as that in [149]. The detailed structure of each layer in decoder is

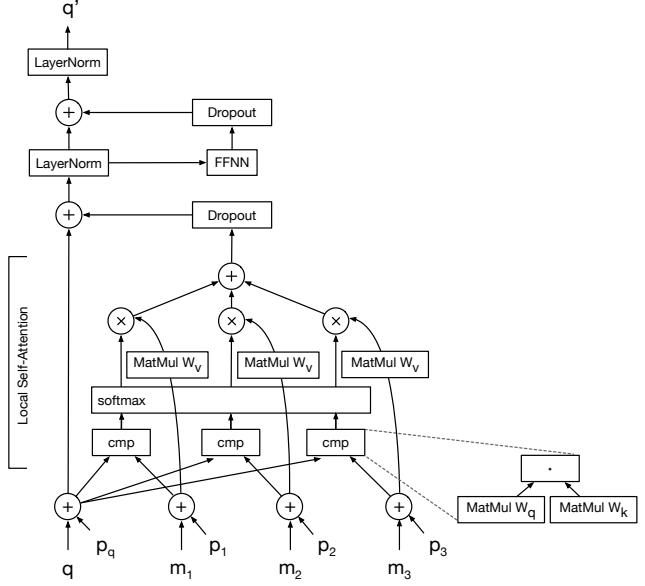


Figure 8. Diagram of one layer in the decoder of Image Transformer (figure from [110]). Each output pixel q' is generated by previously generated pixels m_1, m_2, \dots and the corresponding input pixel q .

shown in Figure 8. Each output pixel q' is generated by calculating self-attention between the input pixel q and previously generated pixels m_1, m_2, \dots with position embedding p_1, p_2, \dots . For image-conditioned generation, such as super-resolution and inpainting, an encoder-decoder architecture is used, where the input of encoder is the low-resolution images or the corrupted images. For unconditional and class-conditional generation (*i.e.*, noise to image), only decoder is used for inputting noise vectors. Since the input for the decoder is the former generated pixels which will bring large computation cost when producing high-resolution images, a local self-attention scheme is proposed, which only use the closest generated pixels as inputs for the decoder. As a result, the Image Transformer can achieve competitive performance with the CNN-based models on image generation and translation tasks, which shows the effectiveness of transformer-based models on low-level vision task.

Compared with using each pixel as inputs of transformer models, recent works use patches (set of pixels) as inputs. Yang *et al.* [163] propose Texture Transformer Network for Image Super-Resolution (TTSR). They use the transformer architecture in the references-based image super-resolution problem, which aims to transfer relevant textures from reference images to the low-resolution images. Taking the low-resolution image and reference image as the query Q and key K , the relevance $r_{i,j}$ is calculated between each patch q_i in Q and k_i in K by

$$r_{i,j} = \frac{q_i}{\|q_i\|} \cdot \frac{k_i}{\|k_i\|} \quad (12)$$

Then a hard-attention module is proposed to select high-resolution features V according to the reference image to match the low-resolution image by using the relevance. The hard-attention map is calculated by

$$h_i = \arg \max_j r_{i,j} \quad (13)$$

Then the most relevant reference patch from is $t_i = v_{h_i}$, where t_i in T is the transferred features. After that, a soft-attention module is used to transfer V to the low-resolution feature F . The soft-attention can be calculated by

$$s_i = \max_j r_{i,j}. \quad (14)$$

Therefore, the equation to transfer the high-resolution texture images to the low-resolution images can be formulated as:

$$F_{out} = F + \text{Conv}(\text{Concat}(F, T)) \odot S, \quad (15)$$

where F_{out} and F denotes the output and input features of the low-resolution image, S is the soft-attention and T is the transferred features from the high-resolution texture image. By introducing the transformer based architecture, TTSR can successfully transfer the texture information from the high resolution referenced images to the low-resolution image for super-resolution task.

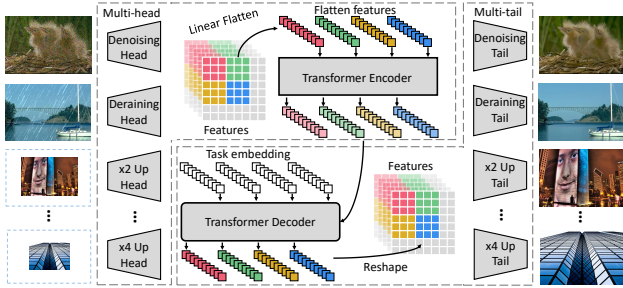


Figure 9. Diagram of the architecture of IPT (figure from [20]). They use multi-head and multi-tail structure to handling different image processing tasks.

The above methods use transformer models on single task, while Chen *et al.* [20] propose Image Processing Transformer (IPT) to fully utilize the advantages of transformers by using large scale pre-training and achieve the state-of-the-art performance in several image processing tasks including super-resolution, denoising and deraining. As shown in Figure 9, IPT consists of multi-heads, an encoder, a decoder and multi-tails. The multi-head and multi-tail structure and task embeddings is introduced for different image processing tasks. The features are divided into patches to put into the encoder-decoder architecture and then the outputs are reshape to features with the same size. As transformer models show advantages on large scale pre-training, IPT use the ImageNet dataset for pre-training. Specifically, Images from ImageNet dataset are

downgraded to generated corrupted images by manually adding noising, rainstreaks or downsampling. The degraded images are then taken as the inputs of IPT and the clean images are taken as the optimization goal of the outputs. A self-supervised method is also introduced to enhance the generalization ability of IPT model. The trained model is then fine-tuned on each task using the corresponding head, tail and task embedding. IPT largely improve the performance on the image processing tasks (*e.g.*, 2dB in image denoising task), which demonstrate the huge potential of transformer-based model on low-level vision fields.

Besides single image generation, Wang *et al.* [?] propose the SceneFormer to utilize transformer in 3D indoor scene generation. By treating a scene as a sequence of objects, the transformer decoder can be naturally applied for predict series of objects and their location, category and size. As a result, the SceneFormer achieve better performance in user study than the conventional CNN-based methods.

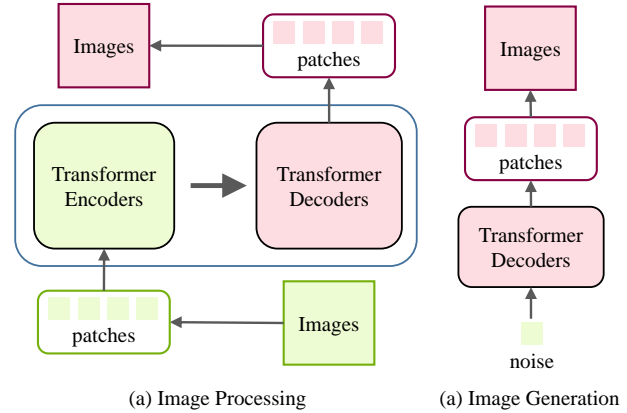


Figure 10. A generic framework for transformer in image processing and generation.

In conclusion, different with classification and detection tasks, the outputs of image generation and processing are images. Figure 10 shows the illustration of using transformers in low-level vision. By taking images as a sequence of pixels or patches, the transformer encoder take the sequence as input sentence and the transformer decoder can successfully produce desired images. In the future work, designing suitable architecture for different image processing tasks would be a meaningful direction.

4.4. Video Processing

The transformer performs surprisingly well on sequence-based tasks, especially on NLP tasks. In computer vision, spatial and temporal dimension information is favored in video tasks. Therefore, the transformer is applied to a number of video tasks, such as frame synthesis [98], action recognition [48], and video retrieval [95].

4.4.1 High-level Video Processing

Human Action Recognition. Video human action task refers to identify and localize human actions in videos. The contextual stuff plays a critical role in recognizing human actions. Rohit *et al.* proposes the action transformer [48] to model the underlying relationship between the human of interest and the surrounding stuff. In specific, the I3D is used as the backbone to extract high-level feature maps. The extracted features from intermediate feature maps by ROI Pooling is viewed as the query (Q). The key (K), values (V) are calculated from the intermediate features. The self-attention mechanism is conducted on three components and outputs the classification and regressions predictions. Lohit *et al.* [100] proposes an interpretable differentiable module, named temporal transformer network to reduce the intra-class variance and increases the inter-class variance. Fayyaz and Gall propose a temporal transformer to perform the action recognition task under weakly supervised settings.

Face Alignment. The video-based face alignment task aims to localize the facial landmarks. The temporal dependency and spatial information are important to the final performance. However, the former methods failed to capture both the temporal information across consecutive frames and the complementary spatial information on a still frame. Therefore, Liu *et al.* [95] uses the two-stream transformer networks to separately learn the temporal and spatial features. Two streams are jointly optimized in an end-to-end manner and the features are weighted to get the final predictions.

Video Retrieval. The key to content-based video retrieval is to find the similarity between videos. To overcome the shortcomings by leveraging only the image-level of video-level features, Shao *et al.* [132] suggests to use the transformer to model the long-range semantic dependency. Moreover, the supervised contrastive learning strategy is introduced to perform hard negative mining. The results on benchmark datasets demonstrate the performance and speed advantages. Gabeur *et al.* [46] presents a multi-modal transformer to learn different cross-modal cues so as to represent the video.

Activity Recognition. The activity recognition refers to identify the activity of a human or a group. The former methods to solve this problem are based on the individual actors' location. Gavriluk *et al.* proposes an actor-transformer [47] architecture to learn the representation. The actor-transformer takes the static and dynamic representations generated by the 2D and 3D networks as input. The output of the transformer is the predicted activity.

Video Object Detection. To detect the objects from a video, global and local information are required. Chen *et al.* introduces the memory enhanced global-local aggregation (MEGA) [22] to capture more content. The representative features enhance the overall performance and address

the *ineffective* and *insufficient* problems. Yin *et al.* [166] proposes a spatiotemporal transformer to aggregate the spatial and temporal information. Together with another spatial feature encoding component, these two components perform well on 3D video object detection task.

Multi-task Learning. The untrimmed video usually contains many frames irrelevant to the target tasks. Therefore, it is crucial to mine the relevant information and remove the redundant information. To cope with the multi-task learning on untrimmed videos, Seong *et al.* adopts the video multi-task transformer network [130] to extract the information. For the CoVieW dataset, the tasks are scene recognition, action recognition and importance score prediction. Two pre-trained networks on ImageNet and Places365 extract the scene features and object features. The Multi-task transformers are stacked to fuse the features with the help of the class conversion matrix (CCM).

4.4.2 Low-level Video Processing

Frame/Video Synthesis. The frame synthesis task refers to synthesize the frames between two continuing frames or after a frame sequence. The video synthesis task aims to synthesize a video. Liu *et al.* proposes the ConvTransformer [98] which comprises five components: feature embedding, position encoding, encoder, query decoder, and the synthesis feed-forward network. Compared with the LSTM based works, the ConvTransformer achieves superior results with more parallelizable architecture. Schatz *et al.* [129] uses a recurrent transformer network to synthesize human actions from novel views.

Video Inpainting. The video inpainting task aims to complete the missing regions in a frame. This challenging task requires to merge information along the spatial and temporal dimensions. Zeng *et al.* proposes a spatial-temporal transformer network [173] for this task. All the input frames are taken as input and they are filled in parallel. The spatial-temporal adversarial loss is used to optimize the transformer network.

4.5. Multimodality

Visual Captioning/Summarization. The goal of visual caption task is to generate text for the untrimmed images or videos. The event detection and description modules are two main parts. Zhou *et al.* [186] proposes an end-to-end optimized transformer to solve the dense video captioning task. The encoder transforms the video into representations. The proposal decoder generates event proposals from the encoding. The caption decoder masks the coding with the proposal and outputs the description. Bilkhu *et al.* [11] uses the C3D and I3D network to extract features and use a transformer to generate the predictions. The algorithm performs well on both single and dense summarization tasks.

Li *et al.* [84] exploits the transformer based on the EnTangled Attention (ETA) module to tackle the image captioning task. Sun *et al.* [34] proposes a visual-linguistic framework to learn the representations without supervision. Zhou *et al.* [185] unifies vision-language pretraining for image captioning and VQA (visual question answering). The model can be applied to a number of tasks, including video captioning, action classification and *etc.*

4.6. Self-attention for Computer Vision

In the above sections, we have reviewed the methods using transformer architecture for visual tasks. Self-attention is the key part of transformer. In this section, we delve deep into the self-attention based methods for challenging tasks in computer vision, *e.g.*, semantic segmentation, instance segmentation, object detection, keypoint detection and depth estimation. We start with formulating the algorithm of self-attention in Section 4.6.1 and summarize the existing applications using self-attention for computer vision in Section 4.6.2.

4.6.1 General Formulation of Self-attention

Self-attention module [149] for machine translation computes the responses at one position in a sequence by attending to all positions and weighted summing them up accordingly in an embedding space, which can be viewed as a form of non-local filtering operations [154, 13] that is applicable in computer vision. We follow the convention [154] to formulate the self-attention module. Given an input signal (*e.g.*, image, sequence, video and feature) $X \in \mathbb{R}^{n \times c}$, where $n = h \times w$ indicates the number of pixels in feature, and c is the number of channels, the output signal is generated as:

$$y_i = \frac{1}{C(x_i)} \sum_{\forall j} f(x_i, x_j) g(x_j), \quad (16)$$

where $x_i \in \mathbb{R}^{1 \times c}$ and $y_i \in \mathbb{R}^{1 \times c}$ indicate the i^{th} position (*e.g.*, space, time and spacetime) of the input signal X and output signal Y , respectively. Subscript j is the index that enumerates all positions. A pairwise function $f(\cdot)$ computes a representing relationship such as affinity between i and all j . The function $g(\cdot)$ computes a representation of the input signal at the position j . The response is normalized by a factor $C(x_i)$.

Note that there are many choices for the pairwise function $f(\cdot)$, for example, a simple extension of the Gaussian function can be used to compute the similarity in an embedding space, thus the function $f(\cdot)$ can be formulated as:

$$f(x_i, x_j) = e^{\theta(x_i)\phi(x_j)^T} \quad (17)$$

where $\theta(\cdot)$ and $\phi(\cdot)$ can be any embedding layers. If we consider the $\theta(\cdot), \phi(\cdot), g(\cdot)$ in the form of linear embedding: $\theta(X) = XW_\theta, \phi(X) = XW_\phi, g(X) = XW_g$ where $W_\theta \in \mathbb{R}^{c \times d_k}, W_\phi \in \mathbb{R}^{c \times d_k}, W_g \in \mathbb{R}^{c \times d_v}$, and set the normalization factor as $C(x_i) = \sum_{\forall j} f(x_i, x_j)$, the E.q. 16 can be rewritten as:

$$y_i = \frac{e^{x_i w_{\theta,i} w_{\phi,j}^T x_j^T}}{\sum_j e^{x_i w_{\theta,i} w_{\phi,j}^T x_j^T}} x_j w_{g,j}, \quad (18)$$

where $w_{\theta,i} \in \mathbb{R}^{c \times 1}$ is the i^{th} row of the weight matrix W_θ . For a given index i , $\frac{1}{C(x_i)} f(x_i, x_j)$ becomes the softmax output along the dimension j , thus the formulation can be further rewritten as:

$$Y = \text{softmax}(XW_\theta W_\phi^T X) g(X), \quad (19)$$

where $Y \in \mathbb{R}^{n \times c}$ is the output signal of the same size as X . Compared to the query, key and value representations $Q = XW_q, K = XW_k, V = XW_v$ from translation module, once $W_q = W_\theta, W_k = W_\phi, W_v = W_g$, E.q. 19 can be formulated as:

$$Y = \text{softmax}(QK^T)V = \text{Attention}(Q, K, V), \quad (20)$$

where the self-attention module [149] proposed for machine translation to some extent is exactly the same as above non-local filtering operations proposed for computer vision.

Generally, the final output signal of self-attention module for computer vision will be wrapped as:

$$Z = YW_Z + X \quad (21)$$

where Y is generated through E.q. 19. If W_Z is initialized as zero, this self-attention module can be inserted to any existing model without breaking its initial behavior.

4.6.2 Applications on Visual Tasks

Self-attention module is considered as a building block of the convolutional neural network architectures, which have low scaling properties concerning the large receptive fields. The self-attention building block is always used on top of the networks to capture long-range interactions and enhance high-level semantic features for computer vision tasks. In what follows, we review the proposed self-attention based methods for image based tasks, such as image classification, semantic segmentation and object detection.

Image Classification. Trainable attention for classification consists of two main streams: hard attention [4, 103, 162] regarding the use of an image region and soft attention [151, 70, 52, 122] generating non-rigid feature maps. Ba *et al.* [4] first proposes the visual attention term for the image classification task, and uses attention to select relevant regions and locations within the input image, which

can also reduce the computational complexity of proposed model w.r.t. the size of input image. AG-CNN [49] proposes to crop a sub-region from global image by the attention heat map for medical image classification. Instead of using hard attention and recalibrating the crop of feature maps, SENet [64] proposes soft self-attention to reweight the channel-wise responses of the convolutional features. Jetley *et al.* [70] uses attention maps generated by corresponding estimators to reweight intermediate features in deep neural networks. Han *et al.* [52] utilize the attribute-aware attention to enhance the representation of CNNs.

Semantic Segmentation. PSANet [182], OCNNet [167], DANet [45] and CFNet [177] are the first works that introduce the self-attention module to semantic segmentation task, which consider and augment the relation and similarity [176, 88, 55, 106, 156, 87] between the contextual pixels. DANet [45] simultaneously leverages self-attention module on spatial and channel dimensions. A^2 Net [23] proposes to group the pixels into a set of regions, and then augment the pixel representations by aggregating the region representations with the generated attention weights. DGCNet [179] proposes a dual graph convolutional network to model coordinate space similarity and feature space similarity in a single framework. To alleviate the huge amount of parameters brought by calculating pixel similarities in self-attention module, several works [169, 69, 68, 89, 79] are proposed to improve the efficiency of the self-attention module for semantic segmentation. For example, CGNL [169] applies the Taylor series of the RBF kernel function to approximate the pixel similarities. CCNet [69] approximates the original self-attention scheme via two consecutive criss-cross attention. ISSA [68] proposes to factorize the dense affinity matrix as the product of two sparse affinity matrices. There are other related works using attention based graph reasoning module [90, 24, 89] to enhance both the local and global representations.

Object Detection. Ramachandran *et al.* [122] proposes an attention-based layer and then swaps the conventional convolution layers to build a fully attentional detector that outperforms the typical RetinaNet [93] on COCO benchmark [94]. GCNet [15] finds that the global contexts modeled by non-local operations are almost the same for different query positions within an image, and proposes to unify the simplified formulation and SENet [64] into a general framework for global context modeling [86, 62, 41, 111]. Vo *et al.* [150] designs a bidirectional operation to gather and distribute the information from a query position to all possible positions. Zhang *et al.* [175] suggests that previous methods failed to interact the cross-scale features, and proposes the self-attention module based Feature Pyramid Transformer to fully exploit interactions across both space and scales.

Conventional detection methods usually exploit a single

visual representation (*e.g.*, bounding box and corner point) for predicting the final results. Hu *et al.* [63] proposes a self-attention based relation module to process a set of objects simultaneously through interaction between their appearance feature. Cheng *et al.* [28] propose RelationNet++ with the bridging visual representations (BVR) module to combine different heterogeneous representations into a single one similar as that in self-attention module. Specifically, the master representation is treated as the query input and the auxiliary representations are regarded as the key input. Thus the enhanced feature can bridge the information from auxiliary representations and benefit final detection results.

Other Vision Tasks. Zhang *et al.* [178] proposes a resolution-wise attention module to learn enhanced feature maps when training multi-resolution networks to obtain accurate human keypoint locations for pose estimation task. Besides, Chang *et al.* [18] improves the accuracy of human keypoint detection model with the help of an attention mechanism based feature fusion block.

Johnston *et al.* [73] directly leverages self-attention module to explore a more general contextual information for improving the self-supervised monocular trained depth estimation. Chen *et al.* [25] also proposes an attention-based aggregation network to capture context information that differs in diverse scenes for depth estimation. Aich *et al.* [1] proposes bidirectional attention modules that utilize the forward and backward attention operations for better results of monocular depth estimation.

4.7. Efficient Transformer

Though Transformer models has achieve success in various tasks, high memory and computing resource are still required, which block the implementation on resource-limited devices (*e.g.*, mobile phones). In this section, we review the researches about compressing & accelerating transformer models for efficient implementation, including network pruning, low-rank decomposition, knowledge distillation, network quantization, compacting architecture design. Table 5 lists a few representative works for compressing transformer-based models.

4.7.1 Pruning & Decomposition

In transformer based pre-trained models (*e.g.*, BERT), multiple attention operations are paralleled to independently model the relationship between different tokens [149, 34], while no all the heads are necessary for a specific task. Michel *et al.* [101] empirically observe that a large percentage of attention heads can be removed at test time without significantly impacting performance. The required numbers of heads vary across different layers and even one head is enough for some layers. Considering the redundancy on attention heads, importance scores are defined to esti-

Table 5. List of representative compressed transformer-based models. The data of the Table is from [118].

Models	Compress Type	#Layer	Params	Speed Up
BERT _{BASE} [34]	Baseline	12	110M	$\times 1$
ALBERT [80]	Decomposition	12	12M	$\times 5.6$
BERT-of-Theseus [161]	Architecture design	6	66M	$\times 1.94$
Q-BERT [133]	Quantization	12	-	-
Q8BERT [171]		12	-	-
TinyBERT [72]	Distillation	4	14.5M	$\times 9.4$
DistilBERT [128]		6	6.6m	$\times 1.63$
BERT-PKD [139]		3~6	45.7~67M	$\times 3.73 \sim 1.64$
MobileBERT [141]		24	25.3M	$\times 4.0$
PD [146]		6	67.5M	$\times 2.0$

mate the influence of each head on the final output in [101], and unimportant heads can be removed for efficient deployment. Dalvi *et al.* [114] further analyzes the redundancy in pretrained transformer models from two perspectives, *i.e.*, general redundancy and task-specific redundancy. Following the lottery ticket hypothesis *et al.* [43], Prasanna *et al.* [114] analyzes the lotteries in BERT and show that good sub-networks also exist in Transformer-based models. Both the FFN layers and attention heads are reduced in [114] to achieve high compression rates.

Besides the width of transformer models, the depth, *i.e.*, the number of layers can also be reduced to accelerate the inference process [39]. Different from that different attention heads in transformer models can be computed in parallel, different layers have to be calculated sequentially due to that the input of the next layer depends on the output of previous layers. Fan *et al.* [39] proposes a layer-wisely dropping strategy to regularize the training of models, and then the whole layers are removed together at the test phase. Considering that the available resources in different devices may vary, Hou *et al.* [61] propose to adaptively reduce the width and depth of the pre-defined transformer models, and obtains multiple models with different sizes simultaneously. Important attention heads and neurons are shared across different sub-networks via a rewiring mechanism.

Beyond the pruning methods that directly discarding part modules in transformer models, matrix decomposition aims to approximate the large matrices with multiple small matrices based on the low-rank assumption. For example, Wang *et al.* [157] decompose the standard matrix multiplication in the transformer models and achieve more efficient inference.

4.7.2 Knowledge Distillation

Knowledge distillation aims to train student networks by transferring knowledge from giant teacher networks [58, 14, 3]. Compared to the teacher networks, student networks usually have thinner and shallower architectures, which are




Figure 11. Span-based dynamic convolution (figure from [71]). LConv denotes the light-weight depth-wise convolution.

easier to be deployed on resource-limited resources. Both the output and intermediate features of neural networks can also be used to transfer effective information from teacher to students. Focused on Transform models, Mukherjee *et al.* [104] use the pre-trained BERT [34] as the teacher to guide the training of small models, with the help of large amounts of unlabeled data. Wang *et al.* [153] train the student networks to mimic outputs of self-attention layers in the pre-trained teacher models. The dot-product between values is introduced as a new form of knowledge for guiding students. A teacher assistant [102] is also introduced in [153], which reduces the gap between the large pre-trained Transformer models and compact student networks to make the mimicking takes more easily. Considering various types of layers in the transformer model (*i.e.*, self-attention layer, embedding layer, prediction layers), Jiao *et al.* [72] design different objective functions to transfer knowledge from teachers to students. For examples, the outputs of the embedding layers of the students models are to imitate the those of teachers via MSE losses. A learnable linear transformation is also imposed to map different features into a same space. For the output of prediction layers, the KL-divergence is adopted to measure the difference between different models.

4.7.3 Quantization

Quantization aim to reduce the number of bits to represent network weight or intermediate features [148, 165]. Quantization methods for general neural networks have been well discussed and achieve performance comparable to the original networks [109, 44, 7]. Recently, how to specially quantize transformer models has been attracted much attention [10, 40]. Shridhar *et al.* [134] suggest to embed the input into binary high-dimensional vectors, and then the binary input representation are used to train the binary neural networks. Cheong *et al.* [26] represent the weights in the transformer models by low-bit (*e.g.*, 4-bit) representa-




Figure 12. Different methods for compressing transformers.

tion. Zhao *et al.* [183] empirically investigate various quantization methods and show that k-means quantization has a huge development potential. Aimed at the machine translation task, Prato *et al.* [115] proposes a fully quantized transformer, which is the first 8-bit quality model without any loss in translation quality as claimed in the paper.

4.7.4 Compact architecture design

Beyond compressing pre-defined transformer models to small ones, some works try to design compact models directly [160, 71]. Jiang *et al.* [71] simplify the calculation of self-attention by proposing a new module named as span-based dynamic convolution, which combining the full-connected layers and convolutional layers, as shown in Figure 11. The local dependence between representation from different tokens is calculated by convolution operations, which is much efficient than the dense full-connected layers in the standard transformers. Depth-wise convolution is also used to further reduce the computing cost. Interesting hamburger layers are proposed in [2], which use matrix decomposition to substitute the original self-attention layers. Matrix decomposition can be calculated more efficiently than the standard self-attention operations while reflect the dependence between different tokens well. The architectures of efficient transformers can also auto-designed with neural architecture search (NAS) [51, 135], which automatically search the combination ways of different components.

The self-attention operation in Transformer models calculate the dot product between representations from different input tokens in a given sequence (patches in image recognition task [37]), whose complexity is $O(N)$, where N is the length of the sequence. Recently, massive methods focus on reducing the complexity to $O(N)$ to make transformer models scalable to long sequences. For example, Katharopoulos *et al.* [75] approximate the self-attention as a linear dot-product of kernel feature maps and reveal the relationship between tokens via recurrent neural networks. Zaheer *et al.* [172] treat each token as a vertex in a graph and the inner product calculation between two tokens is denoted an edge. Inspired graph theories [136, 30], various sparse graph are combined together to approximate the

dense graph in the transformer model, which also achieve $O(N)$ complexity. From a theoretical perspective, Yun *et al.* [170] prove that a sparse transformer with $O(N)$ complexity is enough to reflect any kind of relationship between tokens and can make universal approximation, which provides theoretical guarantees for further research about transformer with $O(N)$ complexity.

Discussion. The above methods discover redundancy of transformer models from different perspectives (Figure 12). Pruning and decomposition based methods usually require pre-defined models with redundancy, *i.e.*, Pruning focuses on reducing the number of components (*e.g.*, layers, heads) in transformer models while decomposition represents an original matrix with multiple small matrices. The compact models also can be directly designed via expertise or automatic search. The obtained compact models can be further represented with low-bits via quantization methods for efficient deployment on resource-limited devices.

5. Conclusions and Discussions

Transformer is becoming a hot topic in computer vision area due to its competitive performance and tremendous potential compared to convolutional neural networks. To discover and utilize the power of transformer, as summarized in the survey, a number of solutions have been proposed in recent years. These methods show excellent performance on a wide range of visual tasks, including basic image classification, high-level vision, low-level vision and video processing. Nevertheless, the potential of transformer for computer vision is yet not fully explored and several challenges are remaining to be solved. In the following, we discuss the challenges of visual transformer and provide some insights on the future prospects.

5.1. Challenges

Although researchers have proposed many transformer-based models to tackle computer vision tasks, these works are the initiatory solutions and have much room for improvement. For example, the transformer architecture in ViT [36] follows the standard transformer for NLP [149]. The improved version specially for CV remains to be explored. Moreover, the applications of transformer on more tasks in addition to those mentioned above are also required.

The generalization and robustness of transformers for computer vision are also challenging. Compared to CNNs, pure transformers lack some inductive biases and heavily rely on the large-scale training on massive data [36]. In consequence, the quality of data much influence the generalization and robustness of transformers. Although ViT show great performance on the downstream image classification tasks like CIFAR [77] and VTAB [174], directly applying ViT backbone on object detection has not achieve better results than CNNs [8]. There is still a long way to

well generalize pretrained transformer on more generalized visual tasks.

Although there have been a series of works on explaining transformers in NLP [131, 158], it is still hard to clearly explain why transformer works well on visual tasks. The inductive biases including translation equivariance and locality are supposed to be the key factors of the CNN’s success, but transformer lacks the inductive bias. The current literature usually analyzes the effect in intuitive way [36, 19]. For example, Dosovitskiy *et al.* [36] claim that large scale training can trump inductive bias. Position embeddings are added into the image patches to retain positional information which is important in vision. Inspired by the heavy parameter usage in transformers, over-parameterization [99, 105] may be a potential point to interpretability of visual transformers.

Last but not the least, developing efficient transformer models for CV is also an open problem. Transformer models are usually huge and computation-expensive, *e.g.*, the base ViT model [36] requires 18B FLOPs to process an image. As a contrast, the lightweight CNN model GhostNet [53, 54] can achieve similar performance with only about 600M FLOPs. Although several methods have been proposed to compress transformer, their complexities are still large. And these methods which are originally designed for NLP may not be suitable for CV. Thus, efficient transformer models are urgent to deploy visual transformer on resource-limited devices.

5.2. Future Prospects

In order to push forward the development of visual transformers, we provide several potential directions for future study.

One research direction is the effectiveness and the efficiency of transformers in computer vision. The goal is to develop high-effectiveness and high-efficiency visual transformers, that is, transformers with high performance and low resource cost. The performance determines whether the model can be applied on real-world applications while the resource cost influences the deployment on the devices. The effectiveness is usually correlated with the efficiency, so how to achieve a better trade-off between them is a meaningful topic for future study.

Besides, most of the existing visual transformer models are designed to handle single task. Many NLP models such as GPT-3 [12] have shown the ability of transformer to deal with multiple tasks in one model. IPT [20] in CV area is also able to process multiple low-level vision tasks, such as super-resolution, image denoising and deraining. We believe that more tasks can be involved in only one model. Unifying all the visual tasks and even other tasks in one transformer (*i.e.*, a grand unified model) is an exciting topic.

There have been various types of neural networks, such

as CNN, RNN and transformer. In computer vision area, CNNs used to be the mainstream choice [56, 143], but now transformer is becoming popular. CNNs are able to capture inductive biases such as translation equivariance and locality, while ViT utilizes large scale training to trump inductive bias [36]. From the current evidence [36], CNNs perform well on small datasets while transformers is better on large datasets. In the future, CNN or transformer, that is the question.

By training with large-scale data, transformers can achieve state-of-the-art performance on both NLP [12, 34] and CV benchmarks [36]. It is potential that what the neural networks need is big data rather than the inductive bias which is assumed by human. What is more, can it obtains satisfactory results with a very simple computational paradigm (*e.g.*, with only fully connected layers) and massive data training?

References

- [1] Shubhra Aich, Jean Marie Uwabeza Vianney, Md Amirul Islam, Mannat Kaur, and Bingbing Liu. Bidirectional attention network for monocular depth estimation. *arXiv preprint arXiv:2009.00743*, 2020. 15
- [2] Anonymous. Is attention better than matrix decomposition? In *Submitted to International Conference on Learning Representations*, 2021. under review. 17
- [3] Jimmy Ba and Rich Caruana. Do deep nets really need to be deep? *Advances in neural information processing systems*, 27:2654–2662, 2014. 16
- [4] Jimmy Ba, Volodymyr Mnih, and Koray Kavukcuoglu. Multiple object recognition with visual attention. In *International Conference on Learning Representations*, 2014. 14
- [5] Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E Hinton. Layer normalization. *arXiv preprint arXiv:1607.06450*, 2016. 4
- [6] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. Neural machine translation by jointly learning to align and translate. *arXiv preprint arXiv:1409.0473*, 2014. 1
- [7] Yu Bai, Yu-Xiang Wang, and Edo Liberty. Proxquant: Quantized neural networks via proximal operators. *arXiv preprint arXiv:1810.00861*, 2018. 16
- [8] Josh Beal, Eric Kim, Eric Tzeng, Dong Huk Park, Andrew Zhai, and Dmitry Kislyuk. Toward transformer-based object detection. *arXiv preprint arXiv:2012.09958*, 2020. 7, 8, 9, 10, 17
- [9] Iz Beltagy, Kyle Lo, and Arman Cohan. Scibert: A pretrained language model for scientific text. *arXiv preprint arXiv:1903.10676*, 2019. 5
- [10] Aishwarya Bhandare, Vamsi Sripathi, Deepthi Karkada, Vivek Menon, Sun Choi, Kushal Datta, and Vikram Sale-tore. Efficient 8-bit quantization of transformer neural machine language translation model. *arXiv preprint arXiv:1906.00532*, 2019. 16
- [11] Manjot Bilkhu, Siyang Wang, and Tushar Dobhal. Attention is all you need for videos: Self-attention based video

- summarization using universal transformers. *arXiv preprint arXiv:1906.02792*, 2019. 13
- [12] Tom B Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *arXiv preprint arXiv:2005.14165*, 2020. 1, 5, 18
- [13] Antoni Buades, Bartomeu Coll, and J-M Morel. A non-local algorithm for image denoising. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 60–65, 2005. 14
- [14] Cristian Bucilua, Rich Caruana, and Alexandru Niculescu-Mizil. Model compression. In *Proceedings of the 12th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 535–541, 2006. 16
- [15] Yue Cao, Jiarui Xu, Stephen Lin, Fangyun Wei, and Han Hu. Gcnet: Non-local networks meet squeeze-excitation networks and beyond. In *Proceedings of the IEEE International Conference on Computer Vision Workshops*, 2019. 15
- [16] Nicolas Carion, Francisco Massa, Gabriel Synnaeve, Nicolas Usunier, Alexander Kirillov, and Sergey Zagoruyko. End-to-end object detection with transformers. *arXiv preprint arXiv:2005.12872*, 2020. 2, 7, 8, 9, 10
- [17] Nicolas Carion, Francisco Massa, Gabriel Synnaeve, Nicolas Usunier, Alexander Kirillov, and Sergey Zagoruyko. End-to-end object detection with transformers. In *ECCV*, 2020. 7, 8, 9
- [18] Yuan Chang, Zixuan Huang, and Qiwei Shen. The same size dilated attention network for keypoint detection. In *International Conference on Artificial Neural Networks*, pages 471–483, 2019. 15
- [19] Hila Chefer, Shir Gur, and Lior Wolf. Transformer interpretability beyond attention visualization. *arXiv preprint arXiv:2012.09838*, 2020. 18
- [20] Hanting Chen, Yunhe Wang, Tianyu Guo, Chang Xu, Yiping Deng, Zhenhua Liu, Siwei Ma, Chunjing Xu, Chao Xu, and Wen Gao. Pre-trained image processing transformer. *arXiv preprint arXiv:2012.00364*, 2020. 2, 12, 18
- [21] Mark Chen, Alec Radford, Rewon Child, Jeffrey Wu, Heewoo Jun, David Luan, and Ilya Sutskever. Generative pre-training from pixels. In *International Conference on Machine Learning*, pages 1691–1703. PMLR, 2020. 1, 2, 6, 8
- [22] Yihong Chen, Yue Cao, Han Hu, and Liwei Wang. Memory enhanced global-local aggregation for video object detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 10337–10346, 2020. 13
- [23] Yunpeng Chen, Yannis Kalantidis, Jianshu Li, Shuicheng Yan, and Jiashi Feng. A²-nets: Double attention networks. *Advances in neural information processing systems*, pages 352–361, 2018. 15
- [24] Yunpeng Chen, Marcus Rohrbach, Zhicheng Yan, Yan Shuicheng, Jiashi Feng, and Yannis Kalantidis. Graph-based global reasoning networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 433–442, 2019. 15
- [25] Yuru Chen, Haitao Zhao, and Zhengwei Hu. Attention-based context aggregation network for monocular depth estimation. *arXiv preprint arXiv:1901.10137*, 2019. 15
- [26] Robin Cheong and Robel Daniel. transformers. zip: Compressing transformers with pruning and quantization. Technical report, tech. rep., Stanford University, Stanford, California, 2019. 16
- [27] Cheng Chi, Fangyun Wei, and Han Hu. Relationnet++: Bridging visual representations for object detection via transformer decoder. In *NeurIPS*, 2020. 7
- [28] Cheng Chi, Fangyun Wei, and Han Hu. Relationnet++: Bridging visual representations for object detection via transformer decoder. *Advances in Neural Information Processing Systems*, 2020. 15
- [29] Yung-Sung Chuang, Chi-Liang Liu, and Hung-Yi Lee. Speechbert: Cross-modal pre-trained language model for end-to-end spoken question answering. *arXiv preprint arXiv:1910.11559*, 2019. 5
- [30] Fan Chung and Linyuan Lu. The average distances in random graphs with given expected degrees. *Proceedings of the National Academy of Sciences*, 99(25):15879–15882, 2002. 17
- [31] Junyoung Chung, Caglar Gulcehre, KyungHyun Cho, and Yoshua Bengio. Empirical evaluation of gated recurrent neural networks on sequence modeling. *arXiv preprint arXiv:1412.3555*, 2014. 4
- [32] Kevin Clark, Minh-Thang Luong, Quoc V Le, and Christopher D Manning. Electra: Pre-training text encoders as discriminators rather than generators. *arXiv preprint arXiv:2003.10555*, 2020. 5
- [33] Zhigang Dai, Bolun Cai, Yugeng Lin, and Junying Chen. UP-DETR: unsupervised pre-training for object detection with transformers. *arXiv preprint arXiv:2011.09094*, 2020. 2, 9, 10
- [34] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *NAACL-HLT (1)*, 2019. 1, 4, 5, 14, 15, 16, 18
- [35] Li Dong, Nan Yang, Wenhui Wang, Furu Wei, Xiaodong Liu, Yu Wang, Jianfeng Gao, Ming Zhou, and Hsiao-Wuen Hon. Unified language model pre-training for natural language understanding and generation. In *Advances in Neural Information Processing Systems*, pages 13063–13075, 2019. 5
- [36] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, et al. An image is worth 16x16 words: Transformers for image recognition at scale. *arXiv preprint arXiv:2010.11929*, 2020. 2, 6, 8, 17, 18
- [37] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, et al. An image is worth 16x16 words: Transformers for image recognition at scale. *arXiv preprint arXiv:2010.11929*, 2020. 17

- [38] Nico Engel, Vasileios Belagiannis, and Klaus Dietmayer. Point transformer. *arXiv preprint arXiv:2011.00931*, 2020. **10**
- [39] Angela Fan, Edouard Grave, and Armand Joulin. Reducing transformer depth on demand with structured dropout. *arXiv preprint arXiv:1909.11556*, 2019. **16**
- [40] Chaofei Fan. Quantized transformer. Technical report, Technical report, Stanford University, Stanford, California, 2019. URL <https://...> **16**
- [41] Qi Fan, Wei Zhuo, Chi-Keung Tang, and Yu-Wing Tai. Few-shot object detection with attention-rpn and multi-relation detector. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 4013–4022, 2020. **15**
- [42] Robert B Fisher. Cvonline: The evolving, distributed, non-proprietary, on-line compendium of computer vision. Retrieved January 28, 2006 from <http://homepages.inf.ed.ac.uk/rbf/CVonline>, 2008. **2**
- [43] Jonathan Frankle and Michael Carbin. The lottery ticket hypothesis: Finding sparse, trainable neural networks. *arXiv preprint arXiv:1803.03635*, 2018. **16**
- [44] Joshua Fromm, Meghan Cowan, Matthai Philipose, Luis Ceze, and Shwetak Patel. Riptide: Fast end-to-end binarized neural networks. *Proceedings of Machine Learning and Systems*, 2:379–389, 2020. **16**
- [45] Jun Fu, Jing Liu, Haijie Tian, Yong Li, Yongjun Bao, Zhiwei Fang, and Hanqing Lu. Dual attention network for scene segmentation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 3146–3154, 2019. **15**
- [46] Valentin Gabeur, Chen Sun, Karteek Alahari, and Cordelia Schmid. Multi-modal transformer for video retrieval. In *European Conference on Computer Vision (ECCV)*, pages 214–229, 2020. **13**
- [47] Kirill Gavrilyuk, Ryan Sanford, Mehrsan Javan, and Cees GM Snoek. Actor-transformers for group activity recognition. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 839–848, 2020. **13**
- [48] Rohit Girdhar, Joao Carreira, Carl Doersch, and Andrew Zisserman. Video action transformer network. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 244–253, 2019. **12, 13**
- [49] Qingji Guan, Yaping Huang, Zhun Zhong, Zhedong Zheng, Liang Zheng, and Yi Yang. Diagnose like a radiologist: Attention guided convolutional neural network for thorax disease classification. In *arXiv preprint arXiv:1801.09927*, 2018. **15**
- [50] Meng-Hao Guo, Jun-Xiong Cai, Zheng-Ning Liu, Tai-Jiang Mu, Ralph R Martin, and Shi-Min Hu. Pct: Point cloud transformer. *arXiv preprint arXiv:2012.09688*, 2020. **10**
- [51] Yong Guo, Yin Zheng, Mingkui Tan, Qi Chen, Jian Chen, Peilin Zhao, and Junzhou Huang. Nat: Neural architecture transformer for accurate and compact architectures. In *Advances in Neural Information Processing Systems*, pages 737–748, 2019. **17**
- [52] Kai Han, Jianyuan Guo, Chao Zhang, and Mingjian Zhu. Attribute-aware attention model for fine-grained representation learning. In *Proceedings of the 26th ACM international conference on Multimedia*, pages 2040–2048, 2018. **14, 15**
- [53] Kai Han, Yunhe Wang, Qi Tian, Jianyuan Guo, Chunjing Xu, and Chang Xu. Ghostnet: More features from cheap operations. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 1580–1589, 2020. **18**
- [54] Kai Han, Yunhe Wang, Qiulin Zhang, Wei Zhang, Chunjing Xu, and Tong Zhang. Model rubik’s cube: Twisting resolution, depth and width for tinynets. *Advances in Neural Information Processing Systems*, 33, 2020. **18**
- [55] Junjun He, Zhongying Deng, Lei Zhou, Yali Wang, and Yu Qiao. Adaptive pyramid context network for semantic segmentation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 7519–7528, 2019. **15**
- [56] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. pages 770–778, 2016. **1, 18**
- [57] Dan Hendrycks and Kevin Gimpel. Gaussian error linear units (gelus). *arXiv preprint arXiv:1606.08415*, 2016. **4**
- [58] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015. **16**
- [59] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997. **1**
- [60] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997. **4**
- [61] Lu Hou, Zhiqi Huang, Lifeng Shang, Xin Jiang, Xiao Chen, and Qun Liu. Dynabert: Dynamic bert with adaptive width and depth. *Advances in Neural Information Processing Systems*, 33, 2020. **16**
- [62] Ting-I Hsieh, Yi-Chen Lo, Hwann-Tzong Chen, and Tyng-Luh Liu. One-shot object detection with co-attention and co-excitation. In *Advances in Neural Information Processing Systems*, pages 2725–2734, 2019. **15**
- [63] Han Hu, Jiayuan Gu, Zheng Zhang, Jifeng Dai, and Yichen Wei. Relation networks for object detection. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 3588–3597, 2018. **15**
- [64] Jie Hu, Li Shen, and Gang Sun. Squeeze-and-excitation networks. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 7132–7141, 2018. **15**
- [65] Kexin Huang, Jaan Altosaar, and Rajesh Ranganath. Clinicalbert: Modeling clinical notes and predicting hospital readmission. *arXiv preprint arXiv:1904.05342*, 2019. **5**
- [66] Lin Huang, Jianchao Tan, Ji Liu, and Junsong Yuan. Hand-transformer: Non-autoregressive structured modeling for 3d hand pose estimation. In *European Conference on Computer Vision*, pages 17–33, 2020. **7, 10**
- [67] Lin Huang, Jianchao Tan, Jingjing Meng, Ji Liu, and Junsong Yuan. Hot-net: Non-autoregressive transformer for 3d hand-object pose estimation. In *Proceedings of the*

- 28th ACM International Conference on Multimedia, pages 3136–3145, 2020. 7, 10
- [68] Lang Huang, Yuhui Yuan, Jianyuan Guo, Chao Zhang, Xilin Chen, and Jingdong Wang. Interlaced sparse self-attention for semantic segmentation. *arXiv preprint arXiv:1907.12273*, 2019. 15
- [69] Zilong Huang, Xinggang Wang, Lichao Huang, Chang Huang, Yunchao Wei, and Wenyu Liu. Ccnet: Criss-cross attention for semantic segmentation. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 603–612, 2019. 15
- [70] Saumya Jetley, Nicholas A Lord, Namhoon Lee, and Philip HS Torr. Learn to pay attention. In *International Conference on Learning Representations*, 2018. 14, 15
- [71] Zi-Hang Jiang, Weihao Yu, Daquan Zhou, Yunpeng Chen, Jiashi Feng, and Shuicheng Yan. Convbert: Improving bert with span-based dynamic convolution. *Advances in Neural Information Processing Systems*, 33, 2020. 2, 16, 17
- [72] Xiaoqi Jiao, Yichun Yin, Lifeng Shang, Xin Jiang, Xiao Chen, Linlin Li, Fang Wang, and Qun Liu. TinyBERT: Distilling BERT for natural language understanding. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 4163–4174, Nov. 2020. 2, 16
- [73] Adrian Johnston and Gustavo Carneiro. Self-supervised monocular trained depth estimation using self-attention and discrete disparity volume. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 4756–4765, 2020. 15
- [74] Mandar Joshi, Danqi Chen, Yinhan Liu, Daniel S Weld, Luke Zettlemoyer, and Omer Levy. Spanbert: Improving pre-training by representing and predicting spans. *Transactions of the Association for Computational Linguistics*, 8:64–77, 2020. 4
- [75] Angelos Katharopoulos, Apoorv Vyas, Nikolaos Pappas, and François Fleuret. Transformers are rnns: Fast autoregressive transformers with linear attention. In *International Conference on Machine Learning*, pages 5156–5165. PMLR, 2020. 17
- [76] Alexander Kolesnikov, Lucas Beyer, Xiaohua Zhai, Joan Puigcerver, Jessica Yung, Sylvain Gelly, and Neil Houlsby. Big transfer (bit): General visual representation learning. In *ECCV*, 2020. 8
- [77] Alex Krizhevsky and Geoffrey Hinton. Learning multiple layers of features from tiny images. Technical report, Cite-seer, 2009. 17
- [78] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with deep convolutional neural networks. In *NeurIPS*, pages 1097–1105, 2012. 1
- [79] Saumya Kumar, Ye Lyu, Francesco Nex, and Michael Ying Yang. Cabinet: Efficient context aggregation network for low-latency semantic segmentation. *arXiv preprint arXiv:2011.00993*, 2020. 15
- [80] Zhenzhong Lan, Mingda Chen, Sebastian Goodman, Kevin Gimpel, Piyush Sharma, and Radu Soricut. Albert: A lite bert for self-supervised learning of language representations. *arXiv preprint arXiv:1909.11942*, 2019. 16
- [81] Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998. 1
- [82] Jinhyuk Lee, Wonjin Yoon, Sungdong Kim, Donghyeon Kim, Sunkyu Kim, Chan Ho So, and Jaewoo Kang. Biobert: a pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics*, 36(4):1234–1240, 2020. 5
- [83] Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Ves Stoyanov, and Luke Zettlemoyer. Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension. *arXiv preprint arXiv:1910.13461*, 2019. 5
- [84] Guang Li, Linchao Zhu, Ping Liu, and Yi Yang. Entangled transformer for image captioning. In *2019 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 8928–8937, 2019. 14
- [85] Liunian Harold Li, Mark Yatskar, Da Yin, Cho-Jui Hsieh, and Kai-Wei Chang. Visualbert: A simple and performant baseline for vision and language. *arXiv preprint arXiv:1908.03557*, 2019. 5
- [86] Wei Li, Kai Liu, Lizhe Zhang, and Fei Cheng. Object detection based on an adaptive attention mechanism. *Scientific Reports*, pages 1–13, 2020. 15
- [87] Xiangtai Li, Li Zhang, Ansheng You, Maoke Yang, Kuiyuan Yang, and Yunhai Tong. Global aggregation then local distribution in fully convolutional networks. *arXiv preprint arXiv:1909.07229*, 2019. 15
- [88] Xia Li, Zhisheng Zhong, Jianlong Wu, Yibo Yang, Zhouchen Lin, and Hong Liu. Expectation-maximization attention networks for semantic segmentation. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 9167–9176, 2019. 15
- [89] Yin Li and Abhinav Gupta. Beyond grids: Learning graph representations for visual recognition. *Advances in Neural Information Processing Systems*, pages 9225–9235, 2018. 15
- [90] Xiaodan Liang, Zhiting Hu, Hao Zhang, Liang Lin, and Eric P Xing. Symbolic graph reasoning meets convolutions. *Advances in Neural Information Processing Systems*, pages 1853–1863, 2018. 15
- [91] Kevin Lin, Lijuan Wang, and Zicheng Liu. End-to-end human pose and mesh reconstruction with transformers. *arXiv preprint arXiv:2012.09760*, 2020. 7, 10, 11
- [92] Matthieu Lin, Chuming Li, Xingyuan Bu, Ming Sun, Chen Lin, Junjie Yan, Wanli Ouyang, and Zhidong Deng. Detr for pedestrian detection. *arXiv preprint arXiv:2012.06785*, 2020. 9
- [93] Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection. In *ICCV*, 2017. 15
- [94] Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *European conference on computer vision*, pages 740–755, 2014. 15

- [95] Hao Liu, Jiwen Lu, Jianjiang Feng, and Jie Zhou. Two-stream transformer networks for video-based face alignment. *IEEE transactions on pattern analysis and machine intelligence*, 40(11):2546–2554, 2017. [12](#), [13](#)
- [96] Ruijin Liu, Zejian Yuan, Tie Liu, and Zhiliang Xiong. End-to-end lane shape prediction with transformers. In *WACV*, 2021. [7](#), [9](#)
- [97] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692*, 2019. [4](#), [5](#)
- [98] Zhouyong Liu, Shun Luo, Wubin Li, Jingben Lu, Yufan Wu, Chunguo Li, and Luxi Yang. Convtransformer: A convolutional transformer network for video frame synthesis. *arXiv preprint arXiv:2011.10185*, 2020. [12](#), [13](#)
- [99] Roi Livni, Shai Shalev-Shwartz, and Ohad Shamir. On the computational efficiency of training neural networks. *Advances in neural information processing systems*, 27:855–863, 2014. [18](#)
- [100] Suhas Lohit, Qiao Wang, and Pavan Turaga. Temporal transformer networks: Joint learning of invariant and discriminative time warping. In *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 12426–12435, 2019. [13](#)
- [101] Paul Michel, Omer Levy, and Graham Neubig. Are sixteen heads really better than one? In *Advances in Neural Information Processing Systems*, pages 14014–14024, 2019. [2](#), [15](#), [16](#)
- [102] Seyed Iman Mirzadeh, Mehrdad Farajtabar, Ang Li, Nir Levine, Akihiro Matsukawa, and Hassan Ghasemzadeh. Improved knowledge distillation via teacher assistant. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 5191–5198, 2020. [16](#)
- [103] Volodymyr Mnih, Nicolas Heess, Alex Graves, et al. Recurrent models of visual attention. *Advances in neural information processing systems*, pages 2204–2212, 2014. [14](#)
- [104] Subhabrata Mukherjee and Ahmed Hassan Awadallah. Xtremedistil: Multi-stage distillation for massive multilingual models. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 2221–2234, 2020. [16](#)
- [105] Behnam Neyshabur, Zhiyuan Li, Srinadh Bhojanapalli, Yann LeCun, and Nathan Srebro. Towards understanding the role of over-parametrization in generalization of neural networks. *arXiv preprint arXiv:1805.12076*, 2018. [18](#)
- [106] Ozan Oktay, Jo Schlemper, Loic Le Folgoc, Matthew Lee, Mattias Heinrich, Kazunari Misawa, Kensaku Mori, Steven McDonagh, Nils Y Hammerla, Bernhard Kainz, et al. Attention u-net: Learning where to look for the pancreas. *arXiv preprint arXiv:1804.03999*, 2018. [15](#)
- [107] Xuran Pan, Zhuofan Xia, Shiji Song, Li Erran Li, and Gao Huang. 3d object detection with pointformer. *arXiv preprint arXiv:2012.11409*, 2020. [7](#)
- [108] Ankur Parikh, Oscar Täckström, Dipanjan Das, and Jakob Uszkoreit. A decomposable attention model for natural language inference. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 2249–2255, 2016. [1](#)
- [109] Eunhyeok Park and Sungjoo Yoo. Profit: A novel training method for sub-4-bit mobilenet models. In *European Conference on Computer Vision*, pages 430–446. Springer, 2020. [16](#)
- [110] Niki Parmar, Ashish Vaswani, Jakob Uszkoreit, Łukasz Kaiser, Noam Shazeer, Alexander Ku, and Dustin Tran. Image transformer. *ICML*, 2018. [2](#), [11](#)
- [111] Hughes Perreault, Guillaume-Alexandre Bilodeau, Nicolas Saunier, and Maguelonne Héritier. Spotnet: Self-attention multi-task network for object detection. In *2020 17th Conference on Computer and Robot Vision (CRV)*, pages 230–237, 2020. [15](#)
- [112] Matthew E Peters, Mark Neumann, Robert L Logan IV, Roy Schwartz, Vidur Joshi, Sameer Singh, and Noah A Smith. Knowledge enhanced contextual word representations. *arXiv preprint arXiv:1909.04164*, 2019. [5](#)
- [113] Tim Prangemeier, Christoph Reich, and Heinz Koeppl. Attention-based transformers for instance segmentation of cells in microstructures. *arXiv preprint arXiv:2011.09763*, 2020. [10](#)
- [114] Sai Prasanna, Anna Rogers, and Anna Rumshisky. When bert plays the lottery, all tickets are winning. *arXiv preprint arXiv:2005.00561*, 2020. [16](#)
- [115] Gabriele Prato, Ella Charlaix, and Mehdi Rezagholizadeh. Fully quantized transformer for machine translation. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: Findings*, pages 1–14, 2020. [2](#), [17](#)
- [116] Charles R Qi, Hao Su, Kaichun Mo, and Leonidas J Guibas. Pointnet: Deep learning on point sets for 3d classification and segmentation. In *CVPR*, pages 652–660, 2017. [10](#)
- [117] Charles Ruizhongtai Qi, Li Yi, Hao Su, and Leonidas J Guibas. Pointnet++: Deep hierarchical feature learning on point sets in a metric space. *Advances in neural information processing systems*, 30:5099–5108, 2017. [10](#)
- [118] Xipeng Qiu, Tianxiang Sun, Yige Xu, Yunfan Shao, Ning Dai, and Xuanjing Huang. Pre-trained models for natural language processing: A survey. *arXiv preprint arXiv:2003.08271*, 2020. [5](#), [16](#)
- [119] Alec Radford, Karthik Narasimhan, Tim Salimans, and Ilya Sutskever. Improving language understanding by generative pre-training, 2018. [5](#)
- [120] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9, 2019. [5](#), [6](#)
- [121] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *arXiv preprint arXiv:1910.10683*, 2019. [5](#)
- [122] Prajit Ramachandran, Niki Parmar, Ashish Vaswani, Irwan Bello, Anselm Levskaya, and Jonathon Shlens. Stand-alone self-attention in vision models. *arXiv preprint arXiv:1906.05909*, 2019. [14](#), [15](#)

- [123] Shaoqing Ren, Kaiming He, Ross Girshick, and Jian Sun. Faster r-cnn: Towards real-time object detection with region proposal networks. In *Advances in neural information processing systems*, pages 91–99, 2015. 1
- [124] Shaoqing Ren, Kaiming He, Ross Girshick, and Jian Sun. Faster R-CNN: Towards real-time object detection with region proposal networks. In *Advances in Neural Information Processing Systems (NIPS)*, 2015. 9, 10
- [125] Frank Rosenblatt. *The perceptron, a perceiving and recognizing automaton Project Para*. Cornell Aeronautical Laboratory, 1957. 1
- [126] FRANK ROSENBLATT. Principles of neurodynamics. perceptrons and the theory of brain mechanisms. Technical report, CORNELL AERONAUTICAL LAB INC BUF-FALO NY, 1961. 1
- [127] David E Rumelhart, Geoffrey E Hinton, and Ronald J Williams. Learning internal representations by error propagation. Technical report, California Univ San Diego La Jolla Inst for Cognitive Science, 1985. 1
- [128] Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019. 16
- [129] Kara Marie Schatz, Erik Quintanilla, Shruti Vyas, and Yogesh Singh Rawat. A recurrent transformer network for novel view action synthesis. In *ECCV (27)*, pages 410–426, 2020. 13
- [130] Hongje Seong, Junhyuk Hyun, and Euntai Kim. Video multitask transformer network. In *Proceedings of the IEEE International Conference on Computer Vision Workshops*, pages 0–0, 2019. 13
- [131] Sofia Serrano and Noah A Smith. Is attention interpretable? In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 2931–2951, 2019. 18
- [132] Jie Shao, Xin Wen, Bingchen Zhao, and Xiangyang Xue. Temporal context aggregation for video retrieval with contrastive learning. 13
- [133] Sheng Shen, Zhen Dong, Jiayu Ye, Linjian Ma, Zhewei Yao, Amir Gholami, Michael W Mahoney, and Kurt Keutzer. Q-bert: Hessian based ultra low precision quantization of bert. In *AAAI*, pages 8815–8821, 2020. 16
- [134] Kumar Shridhar, Harshil Jain, Akshat Agarwal, and Denis Kleyko. End to end binarized neural networks for text classification. In *Proceedings of SustaiNLP: Workshop on Simple and Efficient Natural Language Processing*, pages 29–34, 2020. 16
- [135] David R So, Chen Liang, and Quoc V Le. The evolved transformer. *arXiv preprint arXiv:1901.11117*, 2019. 17
- [136] Daniel A Spielman and Shang-Hua Teng. Spectral sparsification of graphs. *SIAM Journal on Computing*, 40(4):981–1025, 2011. 17
- [137] Weijie Su, Xizhou Zhu, Yue Cao, Bin Li, Lewei Lu, Furu Wei, and Jifeng Dai. Vi-bert: Pre-training of generic visual-linguistic representations. *arXiv preprint arXiv:1908.08530*, 2019. 5
- [138] Chen Sun, Austin Myers, Carl Vondrick, Kevin Murphy, and Cordelia Schmid. Videobert: A joint model for video and language representation learning. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 7464–7473, 2019. 5
- [139] Siqi Sun, Yu Cheng, Zhe Gan, and Jingjing Liu. Patient knowledge distillation for bert model compression. *arXiv preprint arXiv:1908.09355*, 2019. 16
- [140] Zhiqing Sun, Shengcao Cao, Yiming Yang, and Kris Kitani. Rethinking transformer-based set prediction for object detection. *arXiv preprint arXiv:2011.10881*, 2020. 2, 7, 9, 10
- [141] Zhiqing Sun, Hongkun Yu, Xiaodan Song, Renjie Liu, Yiming Yang, and Denny Zhou. Mobilebert: a compact task-agnostic bert for resource-limited devices. *arXiv preprint arXiv:2004.02984*, 2020. 16
- [142] Lucas Tabelini, Rodrigo Berriel, Thiago M Paixão, Claudine Badue, Alberto F De Souza, and Thiago Oliveira-Santos. Polylanenet: Lane estimation via deep polynomial regression. *arXiv preprint arXiv:2004.10924*, 2020. 9
- [143] Mingxing Tan and Quoc Le. Efficientnet: Rethinking model scaling for convolutional neural networks. In *ICML*, 2019. 18
- [144] Zhi Tian, Chunhua Shen, Hao Chen, and Tong He. Fcos: Fully convolutional one-stage object detection. In *ICCV*, pages 9627–9636, 2019. 10
- [145] Hugo Touvron, Matthieu Cord, Matthijs Douze, Francisco Massa, Alexandre Sablayrolles, and Hervé Jégou. Training data-efficient image transformers & distillation through attention. *arXiv preprint arXiv:2012.12877*, 2020. 6, 7, 8
- [146] Iulia Turc, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Well-read students learn better: The impact of student initialization on knowledge distillation. *arXiv preprint arXiv:1908.08962*, 2019. 16
- [147] Shimon Ullman et al. *High-level vision: Object recognition and visual cognition*, volume 2. MIT press Cambridge, MA, 1996. 2
- [148] Vincent Vanhoucke, Andrew Senior, and Mark Z Mao. Improving the speed of neural networks on cpus. 2011. 16
- [149] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30:5998–6008, 2017. 1, 2, 3, 4, 11, 14, 15, 17
- [150] Xuan-Thuy Vo, Lihua Wen, Tien-Dat Tran, and Kang-Hyun Jo. Bidirectional non-local networks for object detection. In *International Conference on Computational Collective Intelligence*, pages 491–501, 2020. 15
- [151] Fei Wang, Mengqing Jiang, Chen Qian, Shuo Yang, Cheng Li, Honggang Zhang, Xiaogang Wang, and Xiaoou Tang. Residual attention network for image classification. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 3156–3164, 2017. 14
- [152] Huiyu Wang, Yukun Zhu, Hartwig Adam, Alan L. Yuille, and Liang-Chieh Chen. Max-deeplab: End-to-end panoptic segmentation with mask transformers. *arXiv preprint arXiv:2012.00759*, 2020. 2, 7, 9

- [153] Wenhui Wang, Furu Wei, Li Dong, Hangbo Bao, Nan Yang, and Ming Zhou. Minilm: Deep self-attention distillation for task-agnostic compression of pre-trained transformers. *arXiv preprint arXiv:2002.10957*, 2020. 16
- [154] Xiaolong Wang, Ross Girshick, Abhinav Gupta, and Kaiming He. Non-local neural networks. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 7794–7803, 2018. 14
- [155] Yuqing Wang, Zhaoliang Xu, Xinlong Wang, Chunhua Shen, Baoshan Cheng, Hao Shen, and Huaxia Xia. End-to-end video instance segmentation with transformers. *arXiv preprint arXiv:2011.14503*, 2020. 2, 7, 10
- [156] Yude Wang, Jie Zhang, Meina Kan, Shiguang Shan, and Xilin Chen. Self-supervised equivariant attention mechanism for weakly supervised semantic segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 12275–12284, 2020. 15
- [157] Ziheng Wang, Jeremy Wohlwend, and Tao Lei. Structured pruning of large language models. *arXiv preprint arXiv:1910.04732*, 2019. 16
- [158] Sarah Wiegrefe and Yuval Pinter. Attention is not not explanation. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 11–20, 2019. 18
- [159] Bichen Wu, Chenfeng Xu, Xiaoliang Dai, Alvin Wan, Peizhao Zhang, Masayoshi Tomizuka, Kurt Keutzer, and Peter Vajda. Visual transformers: Token-based image representation and processing for computer vision. *arXiv preprint arXiv:2006.03677*, 2020. 6
- [160] Zhanghao Wu, Zhijian Liu, Ji Lin, Yujun Lin, and Song Han. Lite transformer with long-short range attention. *arXiv preprint arXiv:2004.11886*, 2020. 17
- [161] Canwen Xu, Wangchunshu Zhou, Tao Ge, Furu Wei, and Ming Zhou. Bert-of-theseus: Compressing bert by progressive module replacing. *arXiv preprint arXiv:2002.02925*, 2020. 16
- [162] Kelvin Xu, Jimmy Ba, Ryan Kiros, Kyunghyun Cho, Aaron Courville, Ruslan Salakhutdinov, Rich Zemel, and Yoshua Bengio. Show, attend and tell: Neural image caption generation with visual attention. In *International conference on machine learning*, pages 2048–2057, 2015. 14
- [163] Fuzhi Yang, Huan Yang, Jianlong Fu, Hongtao Lu, and Baining Guo. Learning texture transformer network for image super-resolution. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 5791–5800, 2020. 2, 11
- [164] Zhilin Yang, Zihang Dai, Yiming Yang, Jaime Carbonell, Russ R Salakhutdinov, and Quoc V Le. Xlnet: Generalized autoregressive pretraining for language understanding. In *Advances in neural information processing systems*, pages 5753–5763, 2019. 5
- [165] Zhaohui Yang, Yunhe Wang, Kai Han, Chunjing Xu, Chao Xu, Dacheng Tao, and Chang Xu. Searching for low-bit weights in quantized neural networks. In *NeurIPS*, 2020. 16
- [166] Junbo Yin, Jianbing Shen, Chenye Guan, Dingfu Zhou, and Ruigang Yang. Lidar-based online 3d video object detection with graph-based message passing and spatiotemporal transformer attention. In *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 11495–11504, 2020. 13
- [167] Yuhui Yuan and Jingdong Wang. Ocnet: Object context network for scene parsing. *arXiv preprint arXiv:1809.00916*, 2018. 15
- [168] Zhenxun Yuan, Xiao Song, Lei Bai, Wengang Zhou, Zhe Wang, and Wanli Ouyang. Temporal-channel transformer for 3d lidar-based video object detection in autonomous driving. *arXiv preprint arXiv:2011.13628*, 2020. 7
- [169] Kaiyu Yue, Ming Sun, Yuchen Yuan, Feng Zhou, Errui Ding, and Fuxin Xu. Compact generalized non-local network. In *Advances in Neural Information Processing Systems*, pages 6510–6519, 2018. 15
- [170] Chulhee Yun, Yin-Wen Chang, Srinadh Bhojanapalli, Ankit Singh Rawat, Sashank J Reddi, and Sanjiv Kumar. $o(n)$ connections are expressive enough: Universal approximability of sparse transformers. *arXiv preprint arXiv:2006.04862*, 2020. 17
- [171] Ofir Zafrit, Guy Boudoukh, Peter Izsak, and Moshe Wasserblat. Q8bert: Quantized 8bit bert. *arXiv preprint arXiv:1910.06188*, 2019. 16
- [172] Manzil Zaheer, Guru Guruganesh, Avinava Dubey, Joshua Ainslie, Chris Alberti, Santiago Ontanon, Philip Pham, Anirudh Ravula, Qifan Wang, Li Yang, et al. Big bird: Transformers for longer sequences. *arXiv preprint arXiv:2007.14062*, 2020. 17
- [173] Yanhong Zeng, Jianlong Fu, and Hongyang Chao. Learning joint spatial-temporal transformations for video inpainting. In *European Conference on Computer Vision*, pages 528–543. Springer, 2020. 2, 13
- [174] Xiaohua Zhai, Joan Puigcerver, Alexander Kolesnikov, Pierre Ruyssen, Carlos Riquelme, Mario Lucic, Josip Djolonga, Andre Susano Pinto, Maxim Neumann, Alexey Dosovitskiy, et al. A large-scale study of representation learning with the visual task adaptation benchmark. *arXiv preprint arXiv:1910.04867*, 2019. 17
- [175] Dong Zhang, Hanwang Zhang, Jinhui Tang, Meng Wang, Xiansheng Hua, and Qianru Sun. Feature pyramid transformer. In *ECCV*, 2020. 7, 15
- [176] Fan Zhang, Yanqin Chen, Zhihang Li, Zhibin Hong, Jingtuo Liu, Feifei Ma, Junyu Han, and Errui Ding. Acfnnet: Attentional class feature network for semantic segmentation. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 6798–6807, 2019. 15
- [177] Hang Zhang, Han Zhang, Chenguang Wang, and Junyuan Xie. Co-occurrent features in semantic segmentation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 548–557, 2019. 15
- [178] Kun Zhang, Peng He, Ping Yao, Ge Chen, Rui Wu, Min Du, Huimin Li, Li Fu, and Tianyao Zheng. Learning enhanced resolution-wise features for human pose estimation. In *2020 IEEE International Conference on Image Processing (ICIP)*, pages 2256–2260, 2020. 15

- [179] Li Zhang, Xiangtai Li, Anurag Arnab, Kuiyuan Yang, Yunhai Tong, and Philip HS Torr. Dual graph convolutional network for semantic segmentation. *arXiv preprint arXiv:1909.06121*, 2019. 15
- [180] Zhengyan Zhang, Xu Han, Zhiyuan Liu, Xin Jiang, Maosong Sun, and Qun Liu. Ernie: Enhanced language representation with informative entities. *arXiv preprint arXiv:1905.07129*, 2019. 5
- [181] Hengshuang Zhao, Li Jiang, Jiaya Jia, Philip Torr, and Vladlen Koltun. Point transformer. *arXiv preprint arXiv:2012.09164*, 2020. 10, 11
- [182] Hengshuang Zhao, Yi Zhang, Shu Liu, Jianping Shi, Chen Change Loy, Dahua Lin, and Jiaya Jia. Psanet: Point-wise spatial attention network for scene parsing. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 267–283, 2018. 15
- [183] Zihan Zhao, Yuncong Liu, Lu Chen, Qi Liu, Rao Ma, and Kai Yu. An investigation on different underlying quantization schemes for pre-trained language models. In *CCF International Conference on Natural Language Processing and Chinese Computing*, pages 359–371. Springer, 2020. 17
- [184] Minghang Zheng, Peng Gao, Xiaogang Wang, Hongsheng Li, and Hao Dong. End-to-end object detection with adaptive clustering transformer. *arXiv preprint arXiv:2011.09315*, 2020. 2, 7, 9, 10, 11
- [185] Luowei Zhou, Hamid Palangi, Lei Zhang, Houdong Hu, Jason J Corso, and Jianfeng Gao. Unified vision-language pre-training for image captioning and vqa. In *AAAI*, pages 13041–13049, 2020. 14
- [186] Luowei Zhou, Yingbo Zhou, Jason J Corso, Richard Socher, and Caiming Xiong. End-to-end dense video captioning with masked transformer. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 8739–8748, 2018. 2, 13
- [187] Xizhou Zhu, Weijie Su, Lewei Lu, Bin Li, Xiaogang Wang, and Jifeng Dai. Deformable detr: Deformable transformers for end-to-end object detection. *arXiv preprint arXiv:2010.04159*, 2020. 2, 7, 8, 9, 10, 11
- [188] Yukun Zhu, Ryan Kiros, Rich Zemel, Ruslan Salakhutdinov, Raquel Urtasun, Antonio Torralba, and Sanja Fidler. Aligning books and movies: Towards story-like visual explanations by watching movies and reading books. In *Proceedings of the IEEE international conference on computer vision*, pages 19–27, 2015. 4