

# CS5489 - Machine Learning

## Lecture 3a - Linear Classifiers

Dr. Antoni B. Chan

Dept. of Computer Science, City University of Hong Kong

### Outline

1. Discriminative linear classifiers
2. Logistic regression
3. Support vector machines (SVM)

## Classification with Generative Model

- Steps to build a classifier
  1. Collect training data (features  $\mathbf{x}$  and class labels  $y$ )
  2. Learn class-conditional distribution (CCD),  $p(\mathbf{x}|y)$ .
  3. Use Bayes' rule to calculate class probability,  $p(y|\mathbf{x})$ .
- **Note:** the data is used to learn the CCD -- the classifier is secondary.
  - Density estimation is an "ill-posed" problem -- which density to use? how much data is needed?
- Advice from Vladimir Vapnik (inventor of SVM):

When solving a problem, try to avoid solving a more general problem as an intermediate step.

- **Discriminative solution**
  - Solve for the classifier  $p(y|\mathbf{x})$  directly!
- Terminology
  - **"Discriminative"** - learn to directly discriminate the classes apart using the features.
  - **"Generative"** - learn model of how the features are generated from different classes.

## Revisit the Naive Bayes Gaussian Classifier

- CCDs: assume the same variance for all Gaussians:
  - $p(\mathbf{x}|y = 1) = \prod_{i=1}^D \mathcal{N}(x_i|\mu_i, \sigma^2)$
  - $p(\mathbf{x}|y = 2) = \prod_{i=1}^D \mathcal{N}(x_i|\nu_i, \sigma^2)$
- prior:
  - $p(y = 1) = \pi_1, p(y = 2) = \pi_2$ .
- look at the log-ratio of CCDs,

$$\begin{aligned}
\log \frac{p(\mathbf{x}|y=1)}{p(\mathbf{x}|y=2)} &= \log \frac{\prod_{i=1}^D \mathcal{N}(x_i|\mu_i, \sigma^2)}{\prod_{i=1}^D \mathcal{N}(x_i|\nu_i, \sigma^2)} \\
&= \sum_{i=1}^D \log \mathcal{N}(x_i|\mu_i, \sigma^2) - \log \mathcal{N}(x_i|\nu_i, \sigma^2) \\
&= \sum_{i=1}^D -\frac{1}{2\sigma^2}(x_i - \mu_i)^2 + \frac{1}{2\sigma^2}(x_i - \nu_i)^2 \\
&= \frac{1}{2\sigma^2} \sum_{i=1}^D (2x_i\mu_i - \mu_i^2 - 2x_i\nu_i + \nu_i^2) \\
&= \frac{1}{2\sigma^2} \sum_{i=1}^D 2(\mu_i - \nu_i)x_i - \mu_i^2 + \nu_i^2
\end{aligned}$$

- Thus

$$\log \frac{p(\mathbf{x}|y=1)}{p(\mathbf{x}|y=2)} = \frac{1}{\sigma^2} \sum_{i=1}^D (\mu_i - \nu_i)x_i + \frac{1}{2\sigma^2} \sum_{i=1}^D (\nu_i^2 - \mu_i^2)$$

- **Bayes decision rule:** Compute the posterior probability of each class  $p(y = j|\mathbf{x})$ 
  - select class 1 when:

$$\begin{aligned}
&\log p(y=1|\mathbf{x}) > \log p(y=2|\mathbf{x}) \\
&\log p(\mathbf{x}|y=1) + \log p(y=1) > \log p(\mathbf{x}|y=2) + \log p(y=2) \\
&\log \frac{p(\mathbf{x}|y=1)}{p(\mathbf{x}|y=2)} + \log \frac{p(y=1)}{p(y=2)} > 0
\end{aligned}$$


- substituting for the CCDs and priors, the BDR is:
  - select class  $y = 1$  when:

$$\sum_{i=1}^D \frac{1}{\sigma^2} (\mu_i - \nu_i)x_i + \frac{1}{2\sigma^2} \sum_{i=1}^D (\nu_i^2 - \mu_i^2) + \log \frac{\pi_1}{\pi_2} > 0$$

- Example

In [8]: pfig

Out[8]:



- BDR in this case is a **linear** function

- select class  $y = 1$  when:

$$\sum_{i=1}^D \underbrace{\frac{1}{\sigma^2}(\mu_i - \nu_i)}_{w_i} x_i + \underbrace{\frac{1}{2\sigma^2} \sum_{i=1}^D (\nu_i^2 - \mu_i^2) + \log \frac{\pi_1}{\pi_2}}_b > 0$$

- $w_i$  is a per-feature weight
- $b$  is a bias term

- the BDR in this case is a **linear** classifier:

- select class  $y = 1$  when
  - $\sum_{i=1}^D w_i x_i + b > 0$
  - equivalently,  $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b > 0$

- Here we obtain the weights  $\mathbf{w}$  by learning the CCDs

- assuming Naive Bayes Gaussians with shared variance.
- this is a generative model since we learn how the data is generated for each class (CCDs).

- How to learn the linear classifier in a discriminative way?

- directly learn the posterior  $p(y|\mathbf{x})$ .
- we will look at a generic linear classifier.

## Linear Classifier

- **Setup**

- Observation (feature vectors)  $\mathbf{x} \in \mathbb{R}^d$
- Class  $y \in \{-1, +1\}$

- **Goal:** given a feature vector  $\mathbf{x}$ , predict its class  $y$ .

- Calculate a *linear function* of the feature vector  $\mathbf{x}$ .
  - $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = \sum_{j=1}^d w_j x_j + b$ 
    - $\mathbf{w} \in \mathbb{R}^d$  are the weights of the linear function.
    - multiply each feature value with a weight, and then add together.
- Predict from the value:
  - if  $f(\mathbf{x}) > 0$  then predict Class  $y = 1$
  - if  $f(\mathbf{x}) < 0$  then predict Class  $y = -1$
  - Equivalently,  $y = \text{sign}(f(\mathbf{x}))$


## Geometric Interpretation

- The linear classifier separates the features space into 2 *half-spaces*
  - corresponding to feature values belonging to Class +1 and Class -1
  - the class boundary is normal to  $\mathbf{w}$ .
    - also called the *separating hyperplane*.
- Example:

$$\mathbf{w} = \begin{bmatrix} 2 \\ 1 \end{bmatrix}, b = 0$$

```
In [11]: linclass
```

```
Out[11]:
```



## Separating Hyperplane

- In a  $d$ -dimensional feature space, the parameters are  $\mathbf{w} \in \mathbb{R}^d$ .
- The equation  $\mathbf{w}^T \mathbf{x} + b = 0$  defines a  $(d - 1)$ -dim. linear surface:
  - for  $d = 2$ ,  $\mathbf{w}$  defines a 1-D line.
  - for  $d = 3$ ,  $\mathbf{w}$  defines a 2-D plane.
  - ...
  - in general, we call it a hyperplane.

## Learning the classifier


- How to set the classifier parameters  $(\mathbf{w}, b)$ ?
  - Learn them from training data!
- Classifiers differ in the objectives used to learn the parameters  $(\mathbf{w}, b)$ .
  - We will look at two examples:
    - *logistic regression*
    - *support vector machine (SVM)*

## Logistic regression

- Use a probabilistic approach
  - Map the linear function  $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$  to probability values between 0 and 1 using a *sigmoid* function.
  - $\sigma(z) = \frac{1}{1+e^{-z}}$


```
In [13]: sigmoidplot
```

```
Out[13]:
```



- Given a feature vector  $x$ , the probability of a class is:
  - $p(y = +1|\mathbf{x}) = \sigma(f(\mathbf{x}))$
  - $p(y = -1|\mathbf{x}) = 1 - \sigma(f(\mathbf{x}))$
- Note: here we are directly modeling the class posterior probability!
  - not the class-conditional  $p(\mathbf{x}|y)$

In [15]: `lrexample`



## Learning the parameters

- Given training data  $\{\mathbf{x}_i, y_i\}_{i=1}^N$ , learn the function parameters  $(\mathbf{w}, b)$  using maximum likelihood estimation.
- maximize the likelihood of the data  $\{\mathbf{x}_i, y_i\}$  according to the posterior:

$$(\mathbf{w}^*, b^*) = \operatorname{argmax}_{\mathbf{w}, b} \sum_{i=1}^N \log p(y_i|\mathbf{x}_i)$$

- posterior is a Bernoulli distribution (given  $\mathbf{x}$ ):


$$p(y|\mathbf{x}) = \begin{cases} \sigma(f(\mathbf{x})), & y = 1 \\ 1 - \sigma(f(\mathbf{x})), & y = -1 \end{cases}$$

- Note the following property:

$$1 - \sigma(z) = \sigma(-z)$$

In [17]: `sigmoidplot`

Out [17]:



- Thus,

$$p(y|\mathbf{x}) = \begin{cases} \sigma(f(\mathbf{x})), & y = 1 \\ \sigma(-f(\mathbf{x})), & y = -1 \end{cases}$$

- Simplifying the 2 cases into one equation,

$$p(y|\mathbf{x}) = \sigma(yf(\mathbf{x}))$$

- Taking the log,

$$\begin{aligned} \log p(y|\mathbf{x}) &= \log \sigma(yf(\mathbf{x})) \\ &= \log \frac{1}{1 + e^{-yf(\mathbf{x})}} \\ &= -\log(1 + e^{-yf(\mathbf{x})}) \end{aligned}$$

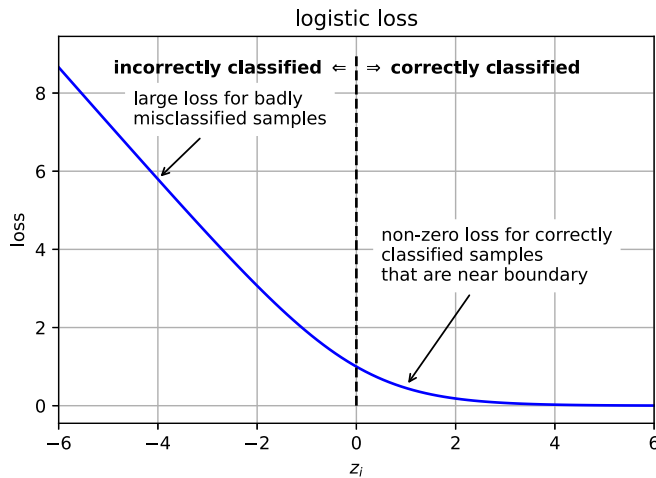
- Substituting into the MLE formulation:

$$\begin{aligned} (\mathbf{w}^*, b^*) &= \operatorname{argmax}_{\mathbf{w}, b} \sum_{i=1}^N \log p(y_i|\mathbf{x}_i) \\ &= \operatorname{argmin}_{\mathbf{w}, b} \sum_{i=1}^N \log(1 + e^{-y_i f(\mathbf{x}_i)}) \end{aligned}$$

- the term on the right is a *data-fit term*
  - wants to make the parameters  $(\mathbf{w}, b)$  to well fit the data.
  - Define  $z_i = y_i f(\mathbf{x}_i)$ 
    - Interesting observation:
      - $z_i > 0$  when sample  $\mathbf{x}_i$  is classified correctly
      - $z_i < 0$  when sample  $\mathbf{x}_i$  is classified incorrectly
      - $z_i = 0$  when sample is on classifier boundary
  - logistic loss function:  $L(z_i) = \log(1 + \exp(-z_i))$

In [19]: `lossfig`

Out[19]:




## Regularization

- to prevent *overfitting*, add a prior distribution on  $\mathbf{w}$ .
  - prefer solutions that are likely under the prior.

$$(\mathbf{w}^*, b^*) = \underset{\mathbf{w}, b}{\operatorname{argmax}} \log p(\mathbf{w}) + \sum_{i=1}^N \log p(y_i | \mathbf{x}_i)$$

- assume Gaussian distribution on  $\mathbf{w}$  with variance  $C/2$ 
  - $p(\mathbf{w}) = \mathcal{N}(\mathbf{w} | 0, \frac{C}{2} \mathbf{I})$ 
    - small values of  $C$  keep  $\mathbf{w}$  close to 0.
    - large values of  $C$  allow larger values of  $\mathbf{w}$ .
  - $\log p(\mathbf{w}) = -\frac{1}{C} \mathbf{w}^T \mathbf{w} + \text{constant}$



- Substituting,


$$(\mathbf{w}^*, b^*) = \underset{\mathbf{w}, b}{\operatorname{argmin}} \frac{1}{C} \mathbf{w}^T \mathbf{w} + \sum_{i=1}^N \log(1 + \exp(-y_i f(\mathbf{x}_i)))$$

- the first term is the *regularization term*
  - Note:  $\mathbf{w}^T \mathbf{w} = \sum_{j=1}^d w_j^2$
  - penalty term that keeps entries in  $\mathbf{w}$  from getting too large.
  - $C$  is the regularization *hyperparameter*
    - larger  $C$  values apply less penalty on large  $\mathbf{w}$  → allow large values in  $\mathbf{w}$ .
    - smaller  $C$  values apply more penalty on large  $\mathbf{w}$  → discourage large values in  $\mathbf{w}$ .
- the second term is the *data fit term* - same as before.

## Optimization

- **no closed-form solution**

- use an iterative optimization algorithm to find the optimal solution
- e.g., *gradient descent* - step downhill in each iteration.
  - $\mathbf{w} \leftarrow \mathbf{w} - \eta \frac{dE}{d\mathbf{w}}$
  - where  $E$  is the objective function
  - $\eta$  is the *learning rate* (how far to step in each iteration).



## Example: Iris Data


```
In [20]: # load iris data each row is (petal length, sepal width, class)
irisdata = loadtxt('iris2.csv', delimiter=',', skiprows=1)

X = irisdata[:,0:2] # the first two columns are features (petal length, sepal width)
Y = irisdata[:,2]   # the third column is the class label (versicolor=1, virginica=2)
                    # --> automatically mapped to (-1, +1) when training classifier

print(X.shape)
```

(100, 2)

```
In [22]: # show the data
plt.figure()
plt.scatter(X[:,0], X[:,1], c=Y, cmap=mycmap, edgecolors='k')
irisaxis(axbox)
```



```
In [23]: # randomly split data into 50% train and 50% test set
trainX, testX, trainY, testY = \
    model_selection.train_test_split(X, Y,
    train_size=0.5, test_size=0.5, random_state=4487)

print(trainX.shape)
print(testX.shape)
```

(50, 2)

(50, 2)




```
In [24]: # learn logistic regression classifier
# (C is a regularization hyperparameter)
logreg = linear_model.LogisticRegression(C=100)
logreg.fit(trainX, trainY)

print("w =", logreg.coef_)
print("b =", logreg.intercept_)
```

```
w = [[9.51275841 0.89596567]]
b = [-48.68254369]
```

- Equation:
  - $f(\mathbf{x}) = (4.87 * \text{petal\_length}) - (0.62 * \text{sepal\_width}) - 21.68$
- Interpretation:
  - large petal length makes  $f(\mathbf{x})$  positive, so large petal length is associated with class +1.

```
In [26]: # show the posterior and training data
plt.figure(figsize=(8,6))
plot_posterior(logreg, axbox, mycmap)
plt.scatter(trainX[:,0], trainX[:,1], c=trainY, cmap=mycmap, edgecolors='k')
plt.title('posterior p(y=+1|x) with training data');
```




```
In [27]: # predict from the model
predY = logreg.predict(testX)

# calculate accuracy
acc = metrics.accuracy_score(testY, predY)
print("test accuracy =", acc)
```

```
test accuracy = 0.92
```

```
In [29]: postfig
```

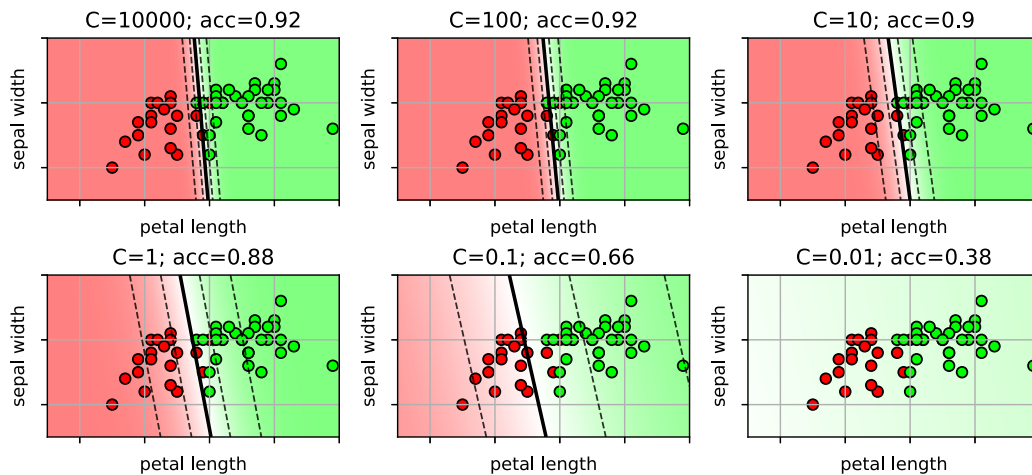


## Selecting the regularization hyperparameter

- the regularization hyperparameter  $C$  has a large effect on the decision boundary and the accuracy.
  - larger  $C$  makes the classifier more confident (posterior probabilities saturate to 0 and 1)
    - more likely to overfit
  - smaller  $C$  makes the classifier less confident (wider range of posterior probabilities).
    - less likely to overfit
- How to set the value of  $C$ ?


In [31]: `lrC`

Out[31]:



## Cross-validation

- Use *cross-validation* on the training set to select the best value of  $C$ .
  - Run many experiments on the training set to see which parameters work on different versions of the data.
    - Split the data into batches of training and validation data.
    - Try a range of  $C$  values on each split.
    - Pick the value that works best over all splits.



### Procedure

- select a range of  $C$  values to try
- Repeat  $K$  times
  - Split the training set into training data and validation data
  - Learn a classifier for each value of  $C$
  - Record the accuracy on the validation data for each  $C$
- Select the value of  $C$  that has the highest average accuracy over all  $K$  folds.
- Retrain the classifier using all data and the selected  $C$ .

- scikit-learn already has built-in `cross_validation` module (more later).

- for logistic regression, use *LogisticRegressionCV* class

```
In [32]: # learn logistic regression classifier using CV
# Cs is an array of possible C values
# cv is the number of folds
# n_jobs=-1 means run in parallel with all cores
logreg = linear_model.LogisticRegressionCV(Cs=logspace(-4,4,20), cv=5, n_jobs=-1)
logreg.fit(trainX, trainY)

print("w=", logreg.coef_)
print("b=", logreg.intercept_)

# predict from the model
predY = logreg.predict(testX)


# calculate accuracy
acc = metrics.accuracy_score(testY, predY)
print("test accuracy=", acc)

w= [[4.61911023 0.72397804]]
b= [-24.24716682]
test accuracy= 0.9
```

## Which C was selected?

```
In [33]: print("C =", logreg.C_)
# calculate the average score for each C
avgscores = mean(logreg.scores_[2],0) # 2 is the class label
plt.semilogx(logreg.Cs_, avgscores, 'ko-')
plt.xlabel('C'); plt.ylabel('average CV accuracy'); plt.grid(True);
```

C = [4.2813324]



## Multi-class classification

- So far, we have only learned a classifier for 2 classes (+1, -1)
  - called a **binary classifier**
- For more than 2 classes, split the problem up into several binary classifier problems.
  - **1-vs-rest**
    - *Training*: for each class, train a classifier for that class versus the other classes.
      - For example, if there are 3 classes, then train 3 binary classifiers: 1 vs {2,3}; 2 vs {1,3}; 3 vs {1,2}
    - *Prediction*: calculate probability for each binary classifier. Select the class with highest probability.

## Example on 3-class Iris data

```
In [34]: # load iris data each row is (petal length, sepal width, class)
irisdata = loadtxt('iris3.csv', delimiter=',', skiprows=1)

X = irisdata[:,0:2] # the first two columns are features (petal length, sepal width)
Y = irisdata[:,2]   # the third column is the class label (setosa=0, versicolor=1, virginica=2)

print(X.shape)
```

```
(150, 2)
```

```
In [35]: # randomly split data into 50% train and 50% test set
trainX, testX, trainY, testY = \
    model_selection.train_test_split(X, Y,
    train_size=0.5, test_size=0.5, random_state=4487)

print(trainX.shape)
print(testX.shape)
```


```
(75, 2)
```

```
(75, 2)
```

```
In [36]: # look at training data

axbox3 = [0.8, 7, 1.5, 4.5]
# make a colormap for viewing 3 classes
mycmap3 = matplotlib.colors.LinearSegmentedColormap.from_list('mycmap', ['#FF0000', '#00FF00', '#0000FF'])

plt.scatter(trainX[:,0], trainX[:,1], c=trainY, cmap=mycmap3, edgecolors='k')
plt.axis(axbox3); plt.grid(True);
plt.xlabel('petal length'); plt.ylabel('sepal width');
```



```
In [37]: # learn logistic regression classifier (one-vs-all)
mlogreg = linear_model.LogisticRegression(C=10, multi_class='ovr')
mlogreg.fit(trainX, trainY)

# now contains 3 hyperplanes and 3 bias terms (one for each class)
print("w=", mlogreg.coef_)
print("b=", mlogreg.intercept_)

# predict from the model
predY = mlogreg.predict(testX)

# calculate accuracy
acc = metrics.accuracy_score(testY, predY)
print("test accuracy=", acc)
```

```
w= [[-3.54884501  1.11513222]
     [-0.03185278 -2.23119433]
     [ 5.41926127 -1.69411622]]
b= [ 6.26247277  6.1229662 -21.79941354]
test accuracy= 0.9733333333333334
```

- the individual 1-vs-rest binary classifiers


```
In [39]: print("w=", mlogreg.coef_)
```

```
print("b=", mlogreg.intercept_)
```

```
mlrfig
```

```
w= [[-3.54884501  1.11513222]
     [-0.03185278 -2.23119433]
     [ 5.41926127 -1.69411622]]
b= [ 6.26247277  6.1229662 -21.79941354]
```

Out[39]:




- the final classifier, combining all 1 vs rest classifiers

In [41]:

```
lr3class
```

Out[41]:



## Multiclass logistic regression

- Another way to get a multi-class classifier is to define a multi-class objective.
  - One weight vector  $\mathbf{w}_c$  for each class  $c$ .
  - linear function for each class,  $f_c(\mathbf{x}) = \mathbf{w}_c^T \mathbf{x}$ .
- Define probabilities with **softmax** function
  - analogous to sigmoid function for binary logistic regression.

$$p(y = c|\mathbf{x}) = \frac{e^{f_c(\mathbf{x})}}{e^{f_1(\mathbf{x})} + \dots + e^{f_K(\mathbf{x})}}$$


- The class with largest response of  $f_c(\mathbf{x})$  will have the highest probability.
- Example with  $K = 3$ .

$$p(y = 1|\mathbf{x}) = \frac{e^{f_1(\mathbf{x})}}{e^{f_1(\mathbf{x})} + e^{f_2(\mathbf{x})} + e^{f_3(\mathbf{x})}}$$

In [43]:

```
sfmfig
```

Out[43]:



## Parameter estimation

- Estimate the  $\{\mathbf{w}_j\}$  parameters using MLE.
- Let  $(\mathbf{x}, \mathbf{y})$  be a data sample pair:
  - $\mathbf{x}$  feature vector.
  - $\mathbf{y} = [y_1, \dots, y_K]$  is a one-hot vector, where  $y_c = 1$  when class  $c$ , and 0 otherwise.
- Data likelihood of  $(\mathbf{x}, \mathbf{y})$ .

$$\text{likelihood:} \quad p(\mathbf{y}|\mathbf{x}) = \prod_{j=1}^K p(y = j|\mathbf{x})^{y_j}$$

$$\text{log-likelihood:} \quad \log p(\mathbf{y}|\mathbf{x}) = \sum_{j=1}^K y_j \log p(y = j|\mathbf{x})$$

$$\text{negative log-likelihood:} \quad -\log p(\mathbf{y}|\mathbf{x}) = -\sum_{j=1}^K y_j \log p(y = j|\mathbf{x})$$

- equivalent to the *cross-entropy loss*
- Given dataset  $\{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^N$ 
  - maximize the data log-likelihood:

$$\max_{\{\mathbf{w}_j\}} \sum_{i=1}^N \log p(\mathbf{y}_i|\mathbf{x}_i) = \max_{\{\mathbf{w}_j\}} \sum_{i=1}^N \sum_{j=1}^K y_{ij} \log p(y = j|\mathbf{x}_i)$$

- i.e., minimize the cross-entropy loss

```
In [44]: # learn logistic regression classifier
mlogreg = linear_model.LogisticRegression(C=10,
                                           multi_class='multinomial')
# use multi-class and corresponding solver
mlogreg.fit(trainX, trainY)

# now contains 3 hyperplanes and 3 bias terms (one for each class)
print("w=", mlogreg.coef_)
print("b=", mlogreg.intercept_)

# predict from the model
predY = mlogreg.predict(testX)

# calculate accuracy
acc = metrics.accuracy_score(testY, predY)
print("test accuracy=", acc)
```

```
w = [[-4.13092437  1.30718735]
```


```

[-0.71717021  0.23609022]
[ 4.84809458 -1.54327757]]
b= [ 11.46078594  5.40723484 -16.86802078]
test accuracy= 0.9733333333333334

```

```
In [46]: lr3classm
```

```
Out[46]:
```



- individual weight vectors work together to partition the space

```
In [48]:
```

```

print("w=", mlogreg.coef_)
print("b=", mlogreg.intercept_)

```


```
lr3lvr
```

```

w= [[-4.13092437  1.30718735]
     [-0.71717021  0.23609022]
     [ 4.84809458 -1.54327757]]
b= [ 11.46078594  5.40723484 -16.86802078]

```

```
Out[48]:
```



## Logistic Regression Summary

- **Classifier:**

- linear function  $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$
- Given a feature vector  $\mathbf{x}$ , the probability of a class is:
  - $p(y = +1|\mathbf{x}) = \sigma(f(\mathbf{x}))$
  - $p(y = -1|\mathbf{x}) = 1 - \sigma(f(\mathbf{x}))$
  - *sigmoid* function:  $\sigma(z) = \frac{1}{1+e^{-z}}$
- logistic loss function:  $L(z) = \log(1 + \exp(-z))$

- **Training:**

- Maximize the likelihood of the training data.
- Use regularization to prevent overfitting.
  - Use cross-validation to pick the regularization hyperparameter  $C$ .

- **Classification:**

- Given a new sample  $\mathbf{x}^*$ :
  - pick class with highest probability  $p(y|\mathbf{x}^*)$ :

$$y^* = \begin{cases} +1, & p(y = +1|\mathbf{x}^*) > p(y = -1|\mathbf{x}^*) \\ -1, & \text{otherwise} \end{cases}$$

- alternatively, just use  $f(\mathbf{x}^*)$

$$y^* = \begin{cases} +1, & f(\mathbf{x}^*) > 0 \\ -1, & \text{otherwise} \end{cases} = \text{sign}(f(\mathbf{x}_*))$$

- **Extend to multi-class:**

- $K$  linear functions, one for each class.
- compute probability using softmax function
- MLE equivalent to cross-entropy loss

In [ ]: