

# CS5489 - Machine Learning

## Lecture 3b - Support Vector Machines

Dr. Antoni B. Chan

Dept. of Computer Science, City University of Hong Kong

### Outline

1. Discriminative classifiers
2. Logistic regression
3. **Support vector machines**


### Support vector machines

- With logistic regression we used a maximum-likelihood framework to learn the separating hyperplane.
- Let's consider a purely geometric approach...

### Linearly-Separable Data

- For now, assume the training data is *linearly separable*
  - the two classes in the training data can be separated by a line (hyperplane)

In [3]: `lsdatafig`

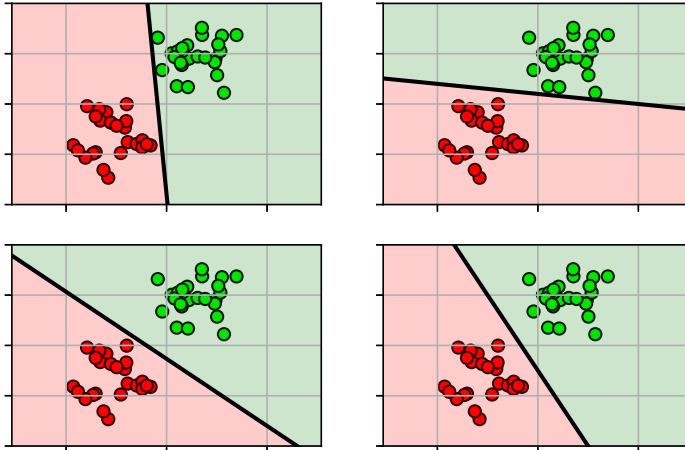


# Which is the best separating line?

- there are many possible solutions...

In [6]: `seplinefig`

Out[6]:




## Maximum margin

- Define the space between the separating line and the closest point as the *margin*.
  - think of this space as the "amount of wiggle room" for accommodating errors in estimating  $\mathbf{w}$ .

In [9]: `margfig`


Out[9]:



- **Idea:** the best separating line is the one that *maximizes the margin*.
  - i.e., puts the most distance between the closest points and the decision boundary.

In [11]: `margfigs`


Out[11]:



- *the solution...*

- by symmetry, there should be at least one margin point on each side of the boundary
- the points on the margins are called the **support vectors**
  - the points support (define) the margin

In [13]: maxmfig




## Why is maximizing the margin good?

- the true  $\mathbf{w}$  is uncertain
  - maximizing the margin allows the most uncertainty (wiggle room) for  $\mathbf{w}$ , while keeping all the points correctly classified.

In [15]: maxmfigw

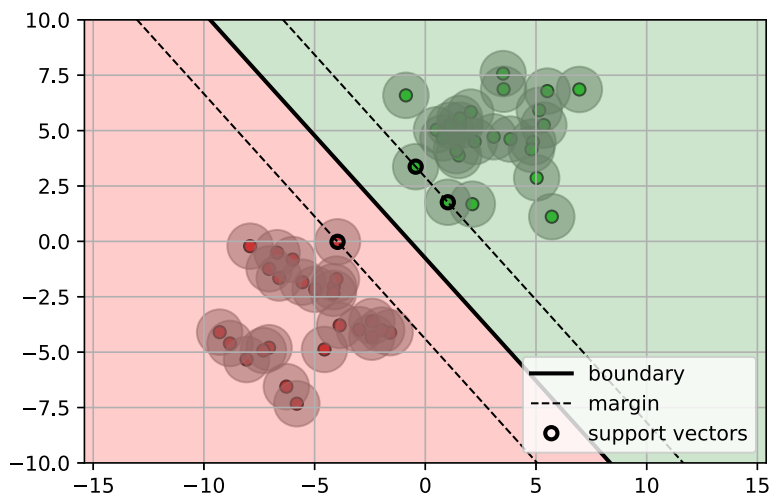
Out[15]:



- the data points are uncertain
  - maximizing the margin allows the most wiggle of the points, while keeping all the points correctly classified.

In [17]: `maxmfigg`

Out[17]:



## SVM Training

- Given a training set  $\{\mathbf{x}_i, y_i\}_{i=1}^N$ .
- First define the margin distance:
  - Distance from a point  $\mathbf{x}_i$  to hyperplane  $\mathbf{w}$ :

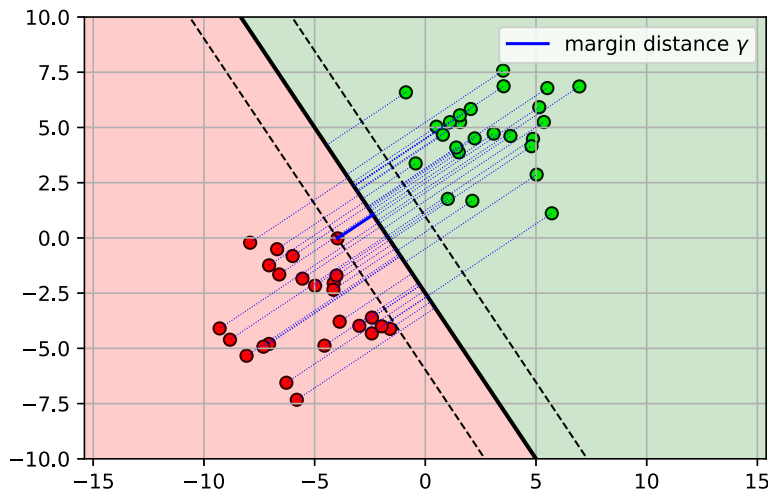
$$d_i = \frac{|f(\mathbf{x}_i)|}{\|\mathbf{w}\|}$$

- Margin distance is the minimum distance among all points:

$$\gamma = \min_i \frac{|f(\mathbf{x}_i)|}{\|\mathbf{w}\|}$$

In [20]: `dfig`

Out[20]:



- The hyperplane  $\mathbf{w}$  appears in both numerator and denominator!
  - $\gamma = \min_i \frac{|f(\mathbf{x}_i)|}{\|\mathbf{w}\|} = \min_i \frac{|\mathbf{w}^T \mathbf{x}_i + b|}{\|\mathbf{w}\|}$
- Changing the length of  $\mathbf{w}$  won't affect the margin distance.
  - $\hat{\mathbf{w}} = a\mathbf{w}, \hat{b} = ab$
  - $\frac{|\hat{\mathbf{w}}^T \mathbf{x}_i + \hat{b}|}{\|\hat{\mathbf{w}}\|} = \frac{|a\mathbf{w}^T \mathbf{x}_i + ab|}{\|a\mathbf{w}\|} = \frac{|\mathbf{w}^T \mathbf{x}_i + b|}{\|\mathbf{w}\|}$
- Margin distance is determined by the direction of  $\mathbf{w}$ , not the length.

## Normalization

- Since the length of  $\mathbf{w}$  doesn't matter, we can assume a normalization for  $\mathbf{w}$ .
- Two possibilities:
  1. set denominator to 1:  $\|\mathbf{w}\| = 1$ 
    - $\|\mathbf{w}\|$  is a unit-norm vector
  2. set numerator to 1:  $\min_i |f(\mathbf{x}_i)| = 1$ 
    - the point  $\mathbf{x}_i$  on the margin has  $f(\mathbf{x}_i) = 1$ .
- Which is better?

## SVM Optimization Problem

- Choose the 2nd option.
  - constraint:  $\min_i |f(\mathbf{x}_i)| = 1$
  - margin:  $\gamma = \frac{1}{\|\mathbf{w}\|}$
- Maximize the margin:

$$(\hat{\mathbf{w}}, b) = \operatorname{argmax}_{\mathbf{w}, b} \frac{1}{\|\mathbf{w}\|} \quad \text{s.t.} \quad \min_i |f(\mathbf{x}_i)| = 1$$

- Invert the objective to turn into a minimization problem

$$(\hat{\mathbf{w}}, b) = \underset{\mathbf{w}, b}{\operatorname{argmin}} \frac{1}{2} \|\mathbf{w}\|^2 \quad \text{s.t.} \quad \min_i |f(\mathbf{x}_i)| = 1$$

- Closer look:
  - original constraint:  $\min_i |f(\mathbf{x}_i)| = 1$ 
    - the minimum over all  $i$  is 1.
  - equivalent constraint:  $|f(\mathbf{x}_i)| \geq 1, \forall i$ 
    - since we are minimizing  $\|\mathbf{w}\|$ ,  $\mathbf{w}$  will shrink at the optimum so that at least one  $\mathbf{x}_i$  will have  $|f(\mathbf{x}_i)| = |\mathbf{w}^T \mathbf{x}_i + b| = 1$ ,
- Note: if the points are correctly classified...
  - for points in class  $y_i = 1$ , then  $f(\mathbf{x}_i) > 0$ .
  - for points in class  $y_i = -1$ , then  $f(\mathbf{x}_i) < 0$ .
- Thus, the constraint:  $|f(\mathbf{x}_i)| \geq 1, \forall i$ 
  - can be rewritten:  $y_i f(\mathbf{x}_i) \geq 1, \forall i$

## SVM Training Objective

- given a training set  $\{\mathbf{x}_i, y_i\}_{i=1}^N$ , optimize:

$$\underset{\mathbf{w}, b}{\operatorname{argmin}} \frac{1}{2} \mathbf{w}^T \mathbf{w} \\ \text{s.t. } y_i f(\mathbf{x}_i) \geq 1, \quad \forall i$$

- the objective minimizes the inverse of the margin distance, i.e., maximizes the margin.
- the inequality constraints ensure that all points are either on or outside of the margin.
  - a point on the margin has  $f(\mathbf{x}_i) = 1$ .


## SVM Prediction

- given a new data point  $\mathbf{x}_*$ , use sign of linear function to predict class
  - $y_* = \operatorname{sign} f(\mathbf{x}_*) = \operatorname{sign}(\mathbf{w}^T \mathbf{x}_* + b)$

## Example

- Regions defined by  $f(\mathbf{x})$ :
  - $f(\mathbf{x}) = 0$  -- decision boundary
  - $f(\mathbf{x}) = \pm 1$  -- positive and negative margins
  - $f(\mathbf{x}) > 1$  -- points correctly classified as class 1
  - $f(\mathbf{x}) < -1$  -- points correctly classified as class -1


In [22]: `egsvmfig`



## What about non-separable data?

- use the same linear classifier
  - allow some training samples to **violate** the margin
    - i.e., are inside the margin (or even mis-classified)
- Define "slack" variable  $\xi_i \geq 0$ 
  - $\xi_i = 0$  means sample is outside of margin area (no slack)
  - $\xi_i > 0$  means sample is inside of margin area (slack)

In [24]: `slackfig`



- constraint now includes slack variable
  - $y_i f(\mathbf{x}_i) \geq 1 - \xi_i, \forall i$
- Two possibilities:
  - $\xi_i = 0$ , then sample  $\mathbf{x}_i$  has normal margin constraint:  $y_i f(\mathbf{x}_i) \geq 1$

- $\xi_i > 0$ , then sample  $\mathbf{x}_i$  is inside margin:  $y_i f(\mathbf{x}_i) = 1 - \xi_i$ 
  - Note that:  $y_i f(\mathbf{x}_i) < 1$ , inside margin.

## Soft-SVM optimization problem

- Penalize each training sample that violates the margin by summing over  $\xi_i$ .
  - penalty controlled by hyperparameter  $C$ .
    - smaller value means allow more violations (less penalty)
    - larger value means don't allow violations (more penalty)


$$\begin{aligned} \operatorname{argmin}_{\mathbf{w}, b} \quad & \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^N \xi_i \\ \text{s. t. } & y_i f(\mathbf{x}_i) \geq 1 - \xi_i, \quad \forall i \\ & \xi_i \geq 0 \end{aligned}$$

- Example with different  $C$ .

In [26]:

Cmargfigs

Out[26]:



## Loss function

- After some massaging, the objective function is:

$$\operatorname{argmin}_{\mathbf{w}, b} \quad \frac{1}{C} \mathbf{w}^T \mathbf{w} + \sum_{i=1}^N \max(0, 1 - y_i f(\mathbf{x}_i))$$


- hinge loss function:  $L(z_i) = \max(0, 1 - z_i)$ 
  - Note:  $\max(a, b)$  returns whichever value ( $a$  or  $b$ ) is largest.

In [28]:

lossfig

Out[28]:





## Example: Iris Data

```
In [29]: # load iris data each row is (petal length, sepal width, class)
irisdata = loadtxt('iris2.csv', delimiter=',', skiprows=1)

X = irisdata[:,0:2] # the first two columns are features (petal length, sepal width)
Y = irisdata[:,2]   # the third column is the class label (versicolor=1, virginica=2)

print(X.shape)
```

```
(100, 2)
```

```
In [30]: # randomly split data into 50% train and 50% test set
trainX, testX, trainY, testY = \
    model_selection.train_test_split(X, Y,
    train_size=0.5, test_size=0.5, random_state=4487)

print(trainX.shape)
print(testX.shape)
```

```
(50, 2)
(50, 2)
```

```
In [31]: # fit the SVM using all the data and the best C
clf = svm.SVC(kernel='linear', C=2)
clf.fit(trainX, trainY)

# get line parameters
w = clf.coef_[0]
b = clf.intercept_[0]
print(w)
print(b)
```


```
[2.87200943 0.10399865]
-13.95923965217775
```

```
In [32]: # indices of data points that are support vectors (on or inside the margin)
clf.support_
```

```
Out[32]: array([ 0,  7, 12, 31, 36, 41, 13, 20, 22, 25, 33, 46], dtype=int32)
```

```
In [34]: svmfig
```

Out [34]:



- SVM doesn't have its own dedicated cross-validation function
- Use the `GridSearchCV` to run cross-validation for a list of parameters
  - calculate average accuracy for each parameter
  - select parameter with average highest accuracy, retrain model with all data
  - Speed up: each parameter can be trained/tested separately, specify number of parallel jobs using `n_jobs`

In [35]:

```
# setup the list of parameters to try
paramgrid = {'C': logspace(-3,3,13)}
print(paramgrid)

# setup the cross-validation object
# pass the SVM object, parameter grid, and number of CV folds
# set number of parallel jobs to -1 (use all cores)
svmcv = model_selection.GridSearchCV(svm.SVC(kernel='linear'), paramgrid, cv=5,
                                     n_jobs=-1, verbose=True)

# run cross-validation (train for each split)
svmcv.fit(trainX, trainY);
```

```
{'C': array([1.00000000e-03, 3.16227766e-03, 1.00000000e-02, 3.16227766e-02,
            1.00000000e-01, 3.16227766e-01, 1.00000000e+00, 3.16227766e+00,
            1.00000000e+01, 3.16227766e+01, 1.00000000e+02, 3.16227766e+02,
            1.00000000e+03])}
```

Fitting 5 folds for each of 13 candidates, totalling 65 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 12 concurrent workers.
[Parallel(n_jobs=-1)]: Done 26 tasks      | elapsed:    1.7s
[Parallel(n_jobs=-1)]: Done 65 out of 65 | elapsed:    1.8s finished
```

In [36]:

```
# show the test error for each parameter set
for m,p in zip(svmcv.cv_results_['mean_test_score'], svmcv.cv_results_['params']):
    print("mean={:.4f} {}".format(m,p))
```

```
mean=0.6200 {'C': 0.001}
mean=0.6200 {'C': 0.0031622776601683794}
mean=0.6200 {'C': 0.01}
mean=0.6600 {'C': 0.03162277660168379}
mean=0.9400 {'C': 0.1}
mean=0.9600 {'C': 0.31622776601683794}
```

```

mean=0.9400 {'C': 1.0}
mean=0.9400 {'C': 3.1622776601683795}
mean=0.9400 {'C': 10.0}
mean=0.9200 {'C': 31.622776601683793}
mean=0.9000 {'C': 100.0}
mean=0.9000 {'C': 316.22776601683796}
mean=0.9000 {'C': 1000.0}

```

In [37]:

```

# make a plot
allC      = [p['C'] for p in svmcv.cv_results_['params']]
allscores = svmcv.cv_results_['mean_test_score']

plt.figure()
plt.semilogx(allC, allscores, 'kx-')
plt.xlabel('C'); plt.ylabel('accuracy')
plt.grid(True)


# view best results and best retrained estimator
print(svmcv.best_params_)
print(svmcv.best_score_)
print(svmcv.best_estimator_)

```

```

{'C': 0.31622776601683794}
0.96
SVC(C=0.31622776601683794, kernel='linear')

```




In [38]:

```

plt.figure()
plot_svm(svmcv.best_estimator_, axbox, mycmap)
plt.scatter(trainX[:,0], trainX[:,1], c=trainY, cmap=mycmap, edgecolors='k')
plt.xlabel('petal length'); plt.ylabel('sepal width')
plt.title('class boundary with training data');

```




```
In [39]: # Directly use svmcv to make predictions
predY = svmcv.predict(testX)

acc = metrics.accuracy_score(testY, predY)
print("test accuracy = " + str(acc))
```

test accuracy = 0.88

```
In [41]: tsvmfig
```



## Multi-class SVM

- In sklearn, `svm.SVC` implements "1-vs-1" multi-class classification.
  - Train binary classifiers on all pairs of classes.
    - 3-class Example: 1 vs 2, 1 vs 3, 2 vs 3
  - To label a sample, pick the class with the most votes among the binary classifiers.
- Problem:
  - 1v1 classification is very slow when there are a large number of classes.
    - if there are  $C$  classes, need to train  $C(C - 1)/2$  binary classifiers!

# 1-vs-all SVM

- Use the `multiclass.OneVsRestClassifier` to build a 1-vs-all classifier from any binary classifier.
  - pass it the binary classifier as the base classifier.
- For `GridSearchCV`, the binary SVM is embedded inside the 1-vs-all wrapper class.
  - use `'estimator__C'` as the parameter label for `C` in the SVM.
  - notation `A__B` means the cross-validated parameter `B` is nested in parameter `A`.

Simple classifier

```
svm.SVC
kernel: 'linear'
C: 10
...
```

Nested classifier

```
OneVsRestClassifier
n_jobs: -1
estimator:
  svm.SVC
  kernel: 'linear'
  C: 10
  ...
```

```
In [43]: msvm = multiclass.OneVsRestClassifier(svm.SVC(kernel='linear'))

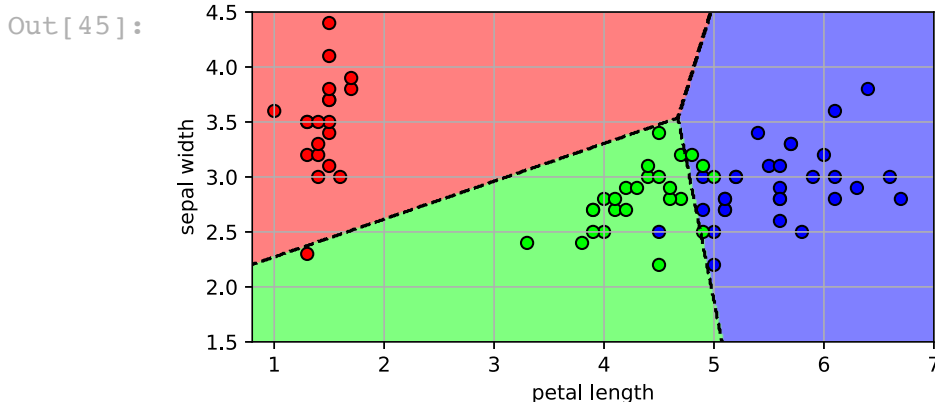
# setup the parameters and run CV
paramgrid = {'estimator__C': logspace(-3,3,13)}
msvmcv = model_selection.GridSearchCV(msvm, paramgrid, cv=5, n_jobs=-1, verbose=True)
msvmcv.fit(trainX, trainY)
print(msvmcv.best_params_)
```

Fitting 5 folds for each of 13 candidates, totalling 65 fits  
{'estimator\_\_C': 10.0}

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 12 concurrent worke
rs.
[Parallel(n_jobs=-1)]: Done 28 tasks      | elapsed:    0.1s
[Parallel(n_jobs=-1)]: Done 42 out of 65 | elapsed:    0.1s remaining:
0.1s
[Parallel(n_jobs=-1)]: Done 65 out of 65 | elapsed:    0.2s finished
```

## 3-class decision boundaries

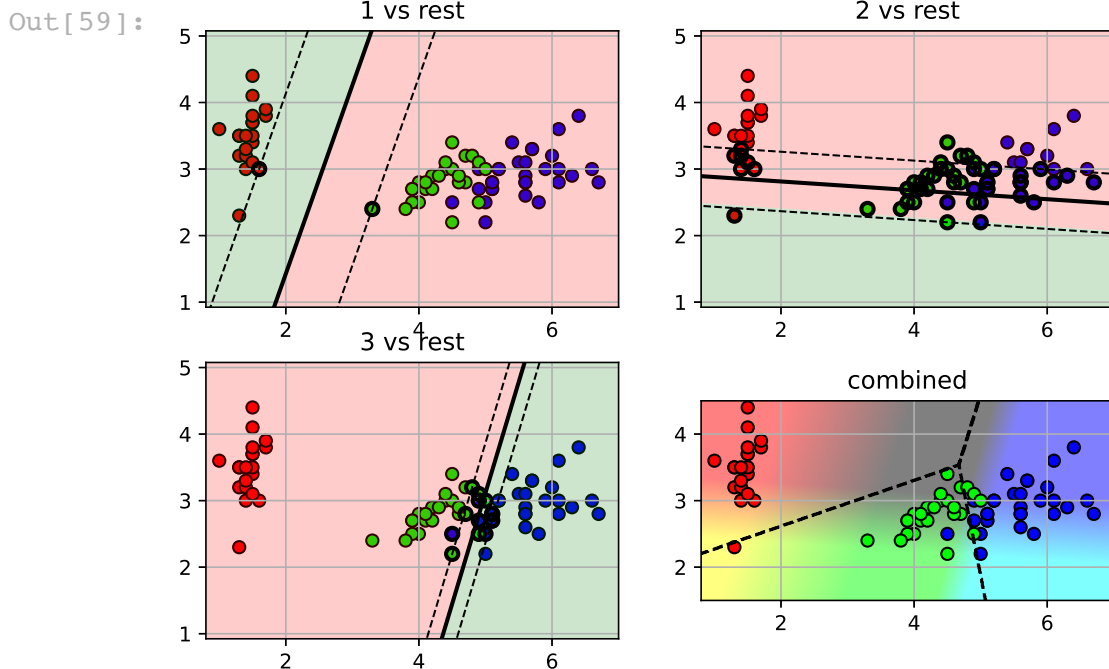
```
In [45]: svm3fig
```



# Decision boundaries for each binary classifier

```
In [59]: for bclf in msvmcv.best_estimator_.estimators_:
          print(bclf.coef_)
          bfig
```

```
[[ -1.04615354  0.36923066]]
[[ -0.14902716 -2.23820701]]
[[  4.44425028 -1.33309667]]
```



## SVM Summary

- **Classifier:**
  - linear function  $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$
  - given new sample  $\mathbf{x}_*$ , predict  $y_* = \text{sign}(f(\mathbf{x}_*))$ .
- **Training:**
  - Maximize the margin of the training data.
    - i.e., maximize the separation between the points and the decision boundary.
  - Allow some training samples to violate the margin.
    - Use cross-validation to pick the hyperparameter  $C$ .

## Summary

- **Linear classifiers:**
  - separate the data using a linear surface (hyperplane).
  - $y = \text{sign}(\mathbf{w}^T \mathbf{x} + b)$
- **Two formulations:**


- logistic regression - maximize the probability of the data
- support vector machine - maximize the margin of the hyperplane

- **Loss functions**

- SVM - ensure a margin of 1 between boundary and closest point
- LR - push the classification boundary as far as possible from all points

In [49]: `lossfig`

Out [49]:



- **Advantages:**

- SVM works well on high-dimensional features ( $d$  large), and has low generalization error.
- LR has well-calibrated probabilities.

- **Disadvantages:**

- decision surface can only be linear!
  - Next lecture we will see how to deal with non-linear decision surfaces.