

CS5489 - Machine Learning

Lecture 8a - Neural Networks

Dr. Antoni B. Chan


Dept. of Computer Science, City University of Hong Kong

Outline


- History
- Perceptron
- Multi-class logistic regression
- Multi-layer perceptron (MLP)

Original idea

- *Perceptron*
 - Warren McCulloch and Walter Pitts (1943), Rosenblatt (1957)
 - Simulate a neuron in the brain
 - 1) take binary inputs (input from nearby neurons)
 - 2) multiply by weights (synapses, dendrites)
 - 3) sum and threshold to get binary output (output axon)
 - Train weights from data.



- Multiple outputs handled by using multiple perceptrons




• **Problem:**

- linear classifier, can't solve harder problems

Multi-layer Perceptron

- Add *hidden* layers between input and output neurons
 - each layer extracts some features from the previous layers
 - can represent complex non-linear functions
 - train weights using *backpropagation* algorithm. (1970-80s)
 - (now called a *neural network*)



- **Problem:**
 - difficult to train.
 - sensitive to initialization.
 - computationally expensive (at the time).


Decline in the 1990s

- Because of those problems, NN became less popular in the 1990s
 - Support vector machines (SVM) had good accuracy
 - easy to use - only one global optimum.
 - learning is not sensitive to initialization.
 - theory about generalization guarantees.
 - Not a lot of data, so kernel methods were still okay.

Deep learning

- There was a resurgence in NN in the 2000s, due to a number of factors:
 - improvements in network architecture
 - developed nodes that are easier to train
 - better training algorithms
 - better ways to prevent overfitting
 - better initialization methods
 - faster computers
 - massively parallel GPUs
 - more labeled data
 - from Internet
 - crowd-sourcing for labeling data (Amazon Turk)

- We can train NN with more and more layers \Rightarrow Deep Learning




Outline

- History
- **Perceptron**
- Multi-class logistic regression
- Multi-layer perceptron (MLP)

Perceptron

- Model a single neuron
 - input $\mathbf{x} \in \mathbb{R}^d$ is a d -dim vector
 - apply a weight to the inputs
 - sum and threshold to get the output
- Formally,
 - $y = f(\sum_{j=0}^d w_j x_j) = f(\mathbf{w}^T \mathbf{x})$
 - \mathbf{w} is the weight vector.
 - $f(a)$ is the activation function
 - $f(a) = \begin{cases} 1, & a > 0 \\ 0, & \text{otherwise} \end{cases}$



Perceptron training criteria

- Train the perceptron on data $D = \{(\mathbf{x}_i, y_i)\}_{i=1}^N$
- Only look at the points that are misclassified.
 - Loss is based on how badly misclassified

$$E(\mathbf{w}) = \sum_{i=1}^N \begin{cases} -y_i \mathbf{w}^T \mathbf{x}_i, & \mathbf{x}_i \text{ is misclassified} \\ 0, & \text{otherwise} \end{cases}$$

- Minimize the loss: $\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} E(\mathbf{w})$


Perceptron Loss Function

- Define $z_i = y_i \mathbf{w}^T \mathbf{x}_i$,
- The loss function is

$$L(z_i) = \max(0, -z_i)$$


In [5]: `lossfig`

Out[5]:



Training algorithm

- We can learn the model by applying gradient descent.
 - Move \mathbf{w} in the direction to decrease the loss $E(\mathbf{w})$.
 - Gradient descent update: $\mathbf{w} \leftarrow \mathbf{w} - \eta \frac{d}{d\mathbf{w}} E(\mathbf{w})$
 - η is the learning rate for gradient descent



- Computers were slow back then...

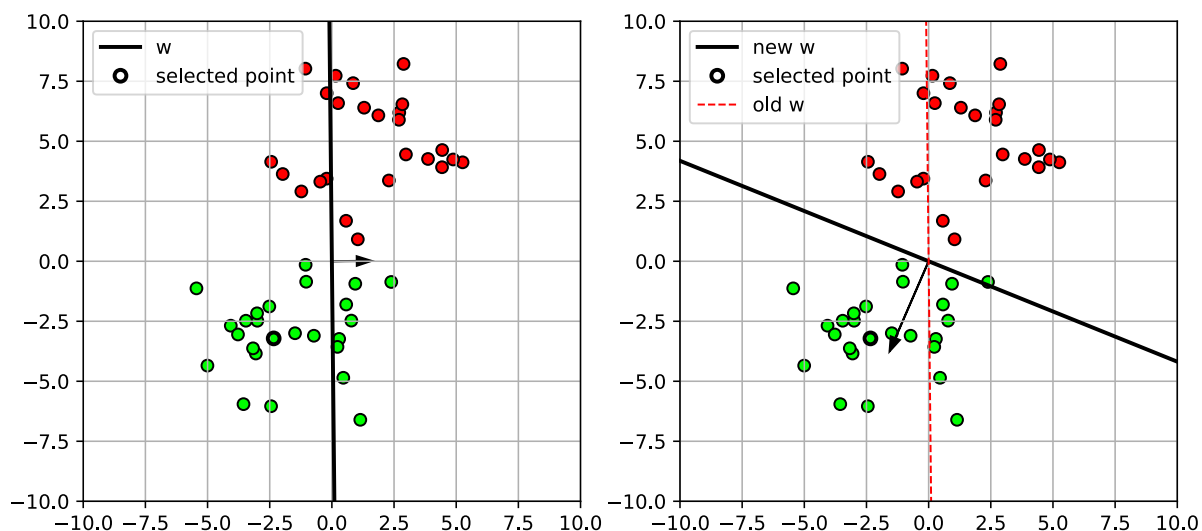
- Solution: only look at one data point at a time and use gradient descent.
 - loss of a misclassified point \mathbf{x}_i : $E_i(\mathbf{w}) = -y_i \mathbf{w}^T \mathbf{x}_i$
 - Gradient: $\frac{d}{d\mathbf{w}} E_i(\mathbf{w}) = -y_i \mathbf{x}_i$
- **Perceptron Algorithm**
 - For each point \mathbf{x}_i ,
 - If the point \mathbf{x}_i is misclassified,
 - Update weights: $\mathbf{w} \leftarrow \mathbf{w} + \eta y_i \mathbf{x}_i$
 - Repeat until no more points are misclassified
- **Notes:**
 - The effect of the update step is to rotate \mathbf{w} towards the misclassified point \mathbf{x}_i .
 - This is called *Stochastic Gradient Descent*.
 - useful because we only need to look at a little bit of data at a time.
 - less computing/memory requirement in each iteration.

Example

- Iteration 1
 - \mathbf{w} rotates towards the misclassified point (bold circle)

In [8]: `figs[0]`


Out[8]:



- Iteration 2

In [9]: `figs[1]`


Out[9]:



- Iteration 3

In [10]: `figs[2]`


Out[10]:



- Iteration 4
 - No more errors


In [11]: `figs[3]`

Out[11]:



- Final classifier

In [12]: `plt.figure(figsize=(4,4))`
`plot_perceptron((w,0),X,Y,axbox)`




Perceptron Algorithm

- Fails to converge if data is not linearly separable
- Rosenblatt proved that the algorithm will converge if the data is linearly separable.
 - the number of iterations is inversely proportional to the separation (margin) between classes.
 - *This was one of the first machine learning results!*
- Different initializations can yield different weights
 - There are multiple decision boundaries with 0 loss.

```
In [14]: pfig
```

```
Out[14]:
```



Outline

- History
- Perceptron
- **Multi-class logistic regression**
- Multi-layer perceptron (MLP)

Revisiting Multiclass logistic regression

- Consider a multi-class classification problem with C classes
 - class labels $y \in \{1, \dots, C\}$
 - equivalently, class vectors:

$$\mathbf{y} \in \left\{ \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix}, \dots, \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix} \right\} = \{\mathbf{e}_1, \mathbf{e}_2, \dots, \mathbf{e}_C\}$$

- \mathbf{e}_j is the canonical vector.

Linear functions

- Construct C linear functions, one for each class
 - $g_j(\mathbf{x}) = \mathbf{w}_j^T \mathbf{x}$, for $j = \{1, \dots, C\}$
 - \mathbf{w}_j is the weight vector for the j -th class.
 - (to reduce clutter, we implicitly include the bias term)
- Combine into a vector-valued function (\mathbb{R}^C):
 - $\mathbf{g}(\mathbf{x}) = \begin{bmatrix} g_1(\mathbf{x}) \\ \vdots \\ g_C(\mathbf{x}) \end{bmatrix} = \mathbf{W}^T \mathbf{x}$,
 - Weight matrix: $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_C]$

Mapping to probabilities

- output of $\mathbf{g}(\mathbf{x})$ is a real vector in \mathbb{R}^C .
- How to map it to set of class probabilities?
 - require $0 \leq p(y = j|\mathbf{x}) \leq 1$
 - and $\sum_{j=1}^C p(y = j|\mathbf{x}) = 1$.

Softmax function

- Given a real vector $\mathbf{a} \in \mathbb{R}^C$
- Let $s_j(\mathbf{a}) = \frac{\exp(a_j)}{\sum_{k=1}^C \exp(a_k)}$
 - if $a_j \gg a_i$, then the exponent will cause $s_j(\mathbf{a}) \rightarrow 1$.
 - denominator ensures $\sum_{j=1}^C s_j(\mathbf{a}) = 1$.
- Let $\mathbf{s}(\mathbf{a}) = [s_1(\mathbf{a}) \cdots s_C(\mathbf{a})]^T$
 - the output vector is ~ 1 in the dimension of \mathbf{a} with largest value, and 0 elsewhere.
 - called the **softmax** function ("soft" because the values can be between 0 and 1)

Mapping to probabilities

- Define the probability of the j -th class $p(y = j|\mathbf{x})$ as:
 - $p(y = j|\mathbf{x}) = f_j(\mathbf{x}) = s_j(\mathbf{g}(\mathbf{x})) = \frac{\exp(g_j(\mathbf{x}))}{\sum_{k=1}^C \exp(g_k(\mathbf{x}))}$
 - if $g_j(\mathbf{x}) \gg g_i(\mathbf{x})$, then the exponent will cause numerator to be very large, and thus $p(y = j|\mathbf{x}) \rightarrow 1$.
 - the class with largest response $g_j(\mathbf{x})$ will have highest probability.
 - denominator ensures probabilities sum to 1 over classes.
- Finally, define the posterior probability vector:


$$\begin{bmatrix} p(y = 1|\mathbf{x}) \\ \vdots \\ p(y = C|\mathbf{x}) \end{bmatrix} = \mathbf{f}(\mathbf{x}) = \mathbf{s}(\mathbf{g}(\mathbf{x}))$$

Example

- linear functions and mapped probabilities

In [16]: `sfig`

Out[16]:



Learning with MLE

- let \mathbf{y} be the *class vector* representation of the class
 - i.e. $y_j = 1$ indicates class $y = j$, and 0 otherwise.
- **Log-likelihood function**
 - categorical distribution

$$\begin{aligned}\log p(\mathbf{y}|\mathbf{x}) &= \log \prod_{j=1}^C f_j(\mathbf{x})^{y_j} \\ &= \sum_{j=1}^C y_j \log f_j(\mathbf{x}) = \mathbf{y}^T \log \mathbf{f}(\mathbf{x})\end{aligned}$$

◦ Note: the log of a vector is element-wise log.

- Maximum Likelihood Estimation (MLE)
 - Let $\mathcal{D} = \{(\mathbf{y}_i, \mathbf{x}_i)\}$ be the training set.
 - MLE goal:

$$\begin{aligned}
\mathbf{W}^* &= \operatorname{argmax}_{\mathbf{W}} \sum_{i=1}^N \log p(\mathbf{y}_i | \mathbf{x}_i) \\
&= \operatorname{argmax}_{\mathbf{W}} \sum_{i=1}^N \mathbf{y}_i^T \log \mathbf{f}(\mathbf{x}_i) \\
&= \operatorname{argmax}_{\mathbf{W}} \sum_{i=1}^N \sum_{j=1}^C y_{ij} \log f_j(\mathbf{x}_i)
\end{aligned}$$


- Equivalently, turn maximization problem into minimization

$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} \sum_{i=1}^N \left\{ - \sum_{j=1}^C y_{ij} \log f_j(\mathbf{x}_i) \right\} = \sum_{i=1}^N L(\mathbf{y}_i, \mathbf{f})$$

- Called the **cross-entropy loss** between ground-truth \mathbf{y}_i and prediction $\mathbf{f}(\mathbf{x}_i)$
 - $L(\mathbf{y}, \mathbf{f}) = - \sum_{j=1}^C y_j \log f_j(\mathbf{x})$

How to optimize?


- Use gradient descent:
 - $\mathbf{W}^{(t)} = \mathbf{W}^{(t-1)} - \eta \frac{dL}{d\mathbf{W}} \Big|_{\mathbf{W}^{(t-1)}}$
 - gradient evaluated at current parameters $\mathbf{W}^{(t-1)}$.
- How do we compute the gradient?
 - We have a composition of functions:
 - $\mathbf{g}(\mathbf{x}) = \mathbf{W}^T \mathbf{x}$
 - $\mathbf{f}(\mathbf{x}) = \mathbf{s}(\mathbf{g}(\mathbf{x}))$
 - $L(\mathbf{y}, \mathbf{f}) = -\mathbf{y}^T \log \mathbf{f}(\mathbf{x})$



- Use the chain rule!
 - in one-dimension:
 - suppose $f(x) = s(g(x))$
 - by the chain rule: $\frac{df}{dx} = \frac{df}{dg} \frac{dg}{dx}$
 - our case is more complicated because of vector-valued functions.

Applying the Chain rule

- Work backwards to compute gradient



- Gradient of loss wrt \mathbf{f} :

$$\frac{dL}{d\mathbf{f}} = \begin{bmatrix} \frac{dL}{df_1} \\ \vdots \\ \frac{dL}{df_C} \end{bmatrix}$$

- Gradient of loss wrt \mathbf{g} :
 - First, look at individual g_j
 - changes in g_j affect all f_k , so sum over derivatives of f_k .

$$\frac{dL}{dg_j} = \sum_{k=1}^C \frac{dL}{df_k} \frac{df_k}{dg_j} = \frac{d\mathbf{f}^T}{dg_j} \frac{dL}{d\mathbf{f}}$$

◦ where $\frac{d\mathbf{f}^T}{dg_j} = \left[\frac{df_1}{dg_j} \cdots \frac{df_C}{dg_j} \right]$.

- Gradient of loss wrt \mathbf{g} :
 - For all linear functions \mathbf{g}

$$\frac{dL}{d\mathbf{g}} = \begin{bmatrix} \frac{dL}{dg_1} \\ \vdots \\ \frac{dL}{dg_C} \end{bmatrix} = \begin{bmatrix} \frac{d\mathbf{f}^T}{dg_1} \frac{dL}{d\mathbf{f}} \\ \vdots \\ \frac{d\mathbf{f}^T}{dg_C} \frac{dL}{d\mathbf{f}} \end{bmatrix} = \frac{d\mathbf{f}^T}{d\mathbf{g}} \frac{dL}{d\mathbf{f}}$$

- where the Jacobian (transpose) is:

$$\frac{d\mathbf{f}^T}{d\mathbf{g}} = \begin{bmatrix} \frac{df_1}{dg_1} & \cdots & \frac{df_C}{dg_1} \\ \vdots & \ddots & \vdots \\ \frac{df_1}{dg_C} & \cdots & \frac{df_C}{dg_C} \end{bmatrix}$$


- (how each input dimension of affects each output dimension of $\mathbf{f}(\mathbf{g})$)

- Gradient of loss wrt \mathbf{w}_j :
 - $\frac{dL}{d\mathbf{w}_j} = \sum_{k=1}^C \frac{dg_k}{d\mathbf{w}_j} \frac{dL}{dg_k} = \frac{d\mathbf{g}^T}{d\mathbf{w}_j} \frac{dL}{d\mathbf{g}}$
 - since weight vector \mathbf{w}_j only appears in g_j
 - $\frac{dL}{d\mathbf{w}_j} = \frac{dg_j}{d\mathbf{w}_j} \frac{dL}{dg_j}$

Summary:

- Chain rule:

- compute the gradient by working backwards from \mathbf{f} to \mathbf{w}_j
 - 1) Gradient of loss wrt \mathbf{f} : $\frac{dL}{d\mathbf{f}}$
 - 2) Gradient of loss wrt \mathbf{g} : $\frac{dL}{d\mathbf{g}} = \frac{d\mathbf{f}^T}{d\mathbf{g}} \frac{dL}{d\mathbf{f}}$
 - 3) Gradient of loss wrt \mathbf{w}_j : $\frac{dL}{d\mathbf{w}_j} = \frac{d\mathbf{g}^T}{d\mathbf{w}_j} \frac{dL}{d\mathbf{g}}$



- Final result after some derivation of gradients:
 - $\frac{dL}{d\mathbf{w}_j} = \mathbf{x}(f_j(\mathbf{x}) - y_j)$

Example

In [18]:

```
# learn logistic regression classifier
mlogreg = linear_model.LogisticRegression(C=10,
                                          multi_class='multinomial', solver='lbfgs')
# use multi-class and corresponding solver
mlogreg.fit(trainX, trainY)

# now contains 3 hyperplanes and 3 bias terms (one for each class)
print("w=", mlogreg.coef_)
print("b=", mlogreg.intercept_)

# predict from the model
predY = mlogreg.predict(testX)


# calculate accuracy
acc = metrics.accuracy_score(testY, predY)
print("test accuracy=", acc)
```

```
w= [[-4.13092437  1.30718735]
     [-0.71717021  0.23609022]
     [ 4.84809458 -1.54327757]]
b= [ 11.46078594   5.40723484 -16.86802078]
test accuracy= 0.9733333333333334
```

- class probabilities from softmax function.

```
In [21]: lr3classm
```

```
Out[21]:
```



Summary

- Two related linear classification models: Perceptron, Multi-class logistic regression
 - compute a linear function of input
 - non-linear output function (threshold or soft-max)
- We have assumed the inputs are feature vectors already.
 - What if we want to extract the features vectors too?