

CS5489 - Machine Learning

Lecture 5a - Supervised Learning - Regression

Dr. Antoni B. Chan

Dept. of Computer Science, City University of Hong Kong

Outline

1. Linear Regression
2. Selecting Features
3. Removing Outliers
4. Non-linear regression


Regression

- Supervised learning
 - Input observation \mathbf{x} , typically a vector in \mathbb{R}^d .
 - Output $y \in \mathbb{R}$, a real number.
- **Goal:** predict output y from input \mathbf{x} .
 - i.e., learn the function $y = f(\mathbf{x})$.

In [3]:

```
linfig
```


Out[3]:



Examples:

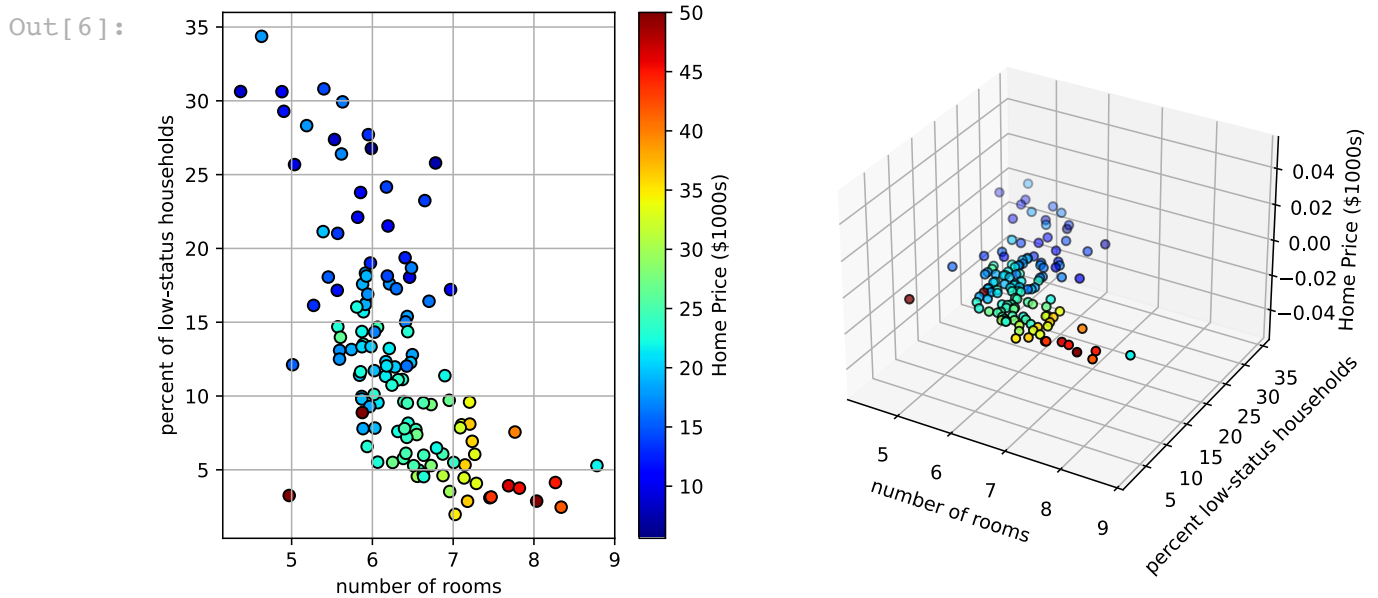
- Predict Boston house price from number of rooms, or percentage of low-status households in neighborhood.

In [5]: `boston1dfig`



- predict from both features

In [6]: `boston2dfig`




Linear Regression

- **1-d case:** the output y is a linear function of input feature x
 - $y = w * x + b$
 - w is the slope, b is the intercept.

In [8]: `linfig`

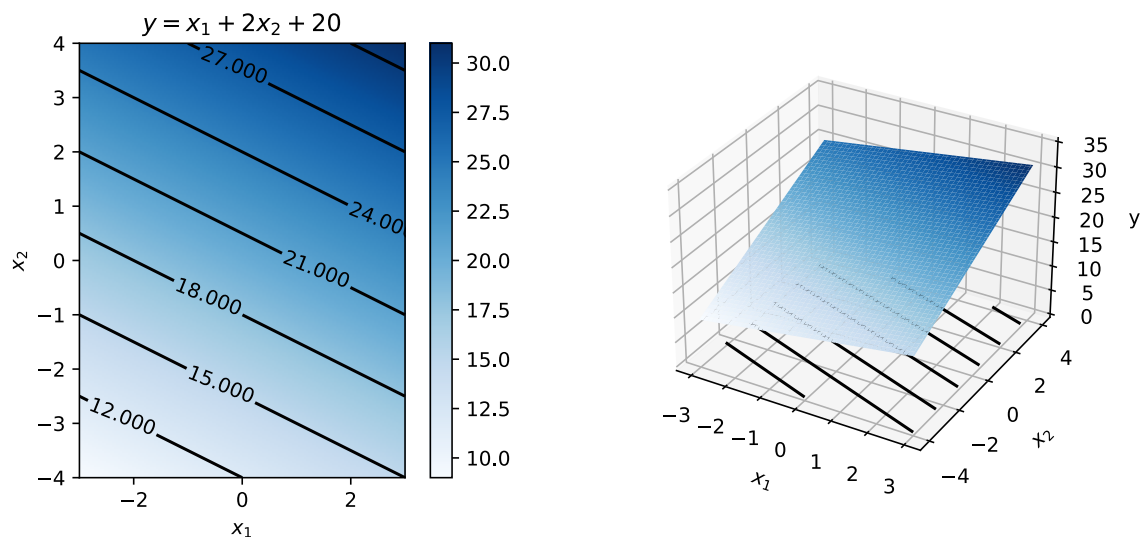
Out[8]:



- **d-dim case:** the output y is a linear combination of d input variables x_1, \dots, x_d :
 - $y = w_0 + w_1x_1 + w_2x_2 + \dots + w_dx_d$
- Equivalently,
 - $y = w_0 + \mathbf{w}^T \mathbf{x} = w_0 + \sum_{j=1}^d w_j x_j$
 - $\mathbf{x} \in \mathbb{R}^d$ is the vector of input values.
 - $\mathbf{w} \in \mathbb{R}^d$ are the weights of the linear function, and w_0 is the intercept (bias term).

In [10]: `lin2dfig`

Out[10]:



Ordinary Least Squares (OLS)

- The linear function has form $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$.
- *How to estimate the parameters (\mathbf{w}, b) from the data?*
- Fit the parameters by minimizing the squared prediction error on the training set $\{(\mathbf{x}_i, y_i)\}_{i=1}^N$:

$$\min_{\mathbf{w}, b} \sum_{i=1}^N (y_i - f(\mathbf{x}_i))^2 = \min_{\mathbf{w}, b} \sum_{i=1}^N (y_i - (\mathbf{w}^T \mathbf{x}_i + b))^2$$

- The bias term b can be absorbed into \mathbf{w} by redefining as follows:

$$\mathbf{w} \leftarrow \begin{bmatrix} \mathbf{w} \\ b \end{bmatrix}, \mathbf{x} \leftarrow \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix}$$

- We can write the minimization problem as:

$$\min_{\mathbf{w}} \|\mathbf{y} - \mathbf{X}^T \mathbf{w}\|^2$$

- where $\mathbf{X} = [\mathbf{x}_1 \cdots \mathbf{x}_N]$ is the data matrix,
- and $\mathbf{y} = [y_1, \dots, y_N]^T$ is vector of outputs.

- To obtain the solution:

- 1) Expand the norm term:

$$\begin{aligned} \|\mathbf{y} - \mathbf{X}^T \mathbf{w}\|^2 &= (\mathbf{y} - \mathbf{X}^T \mathbf{w})^T (\mathbf{y} - \mathbf{X}^T \mathbf{w}) \\ &= \mathbf{y}^T \mathbf{y} - 2\mathbf{y}^T \mathbf{X}^T \mathbf{w} + \mathbf{w}^T \mathbf{X} \mathbf{X}^T \mathbf{w} \end{aligned}$$

- Find the minimum by taking the derivative and setting to 0:

$$\begin{aligned} \frac{d}{d\mathbf{w}} (\mathbf{y}^T \mathbf{y} - 2\mathbf{y}^T \mathbf{X}^T \mathbf{w} + \mathbf{w}^T \mathbf{X} \mathbf{X}^T \mathbf{w}) &= -2\mathbf{X} \mathbf{y} + 2\mathbf{X} \mathbf{X}^T \mathbf{w} = 0 \\ \Rightarrow \mathbf{X} \mathbf{X}^T \mathbf{w} &= \mathbf{X} \mathbf{y} \\ \Rightarrow \mathbf{w}^* &= (\mathbf{X} \mathbf{X}^T)^{-1} \mathbf{X} \mathbf{y} \end{aligned}$$


- closed-form solution!

- Note: $(\mathbf{X} \mathbf{X}^T)^{-1} \mathbf{X}$ is also called the *pseudo-inverse* of \mathbf{X} .

Examples: 1-d

```
In [12]: # fit using ordinary least squares
ols = linear_model.LinearRegression()
ols.fit(linX, linY)

# show plot
axbox = [-3, 3, -60, 100]
plt.figure()
plot_linear_1d(ols, axbox, linX, linY)
plt.xlabel('feature x'); plt.ylabel('output y');
```




Boston housing price (1d)

- learn regression function for each feature separately

```
In [13]: ols = [None]*2
for i in range(2):
    ols[i] = linear_model.LinearRegression()
    tmpX = bostonX[:,i][:,newaxis]
    ols[i].fit(tmpX, bostonY)
```

```
In [15]: ofig
```

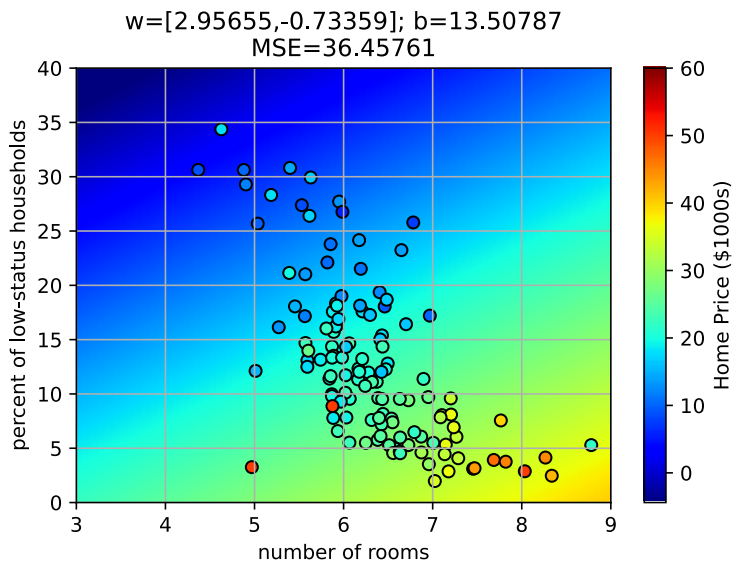


- for both features together

```
In [17]: # learn with both dimensions
ols = linear_model.LinearRegression()
ols.fit(bostonX, bostonY);
```

```
In [19]: ofig
```

```
Out[19]:
```




- interpretation from the linear model parameters:
 - each room increases home price by \$2956 (w_1)
 - each percentage of low-status households decreases home price by \$733 (w_2)
 - the "starting" price is \$13,508 (b).

Selecting Features

- The Boston housing data actually has 12 features.
 - plots of feature vs. housing price

In [21]: `bostonffig`

Out[21]:




- Use `pandas` to view pairwise relationships
 - diagonal shows the histogram
 - off-diagonal shows plots for two features at a time

In [22]: `import pandas as pd`
`boston_feature_names = [x[0] for x in bostonAttr]`

```
boston_df = pd.DataFrame(bostonX, columns=boston_feature_names)

tmp=pd.plotting.scatter_matrix(boston_df, c=bostonY, figsize=(9, 9),
                              marker='o', hist_kwds={'bins': 20}, s=10,
                              alpha=.5)
```



- Can we select a few features that are good for predicting the price?
 - This will provide some insight about our data and what is important.

Shrinkage

- Add a *regularization* term to "shrink" some linear weights to zero.
 - features associated with zero weight are not important since they aren't used to calculate the function output.
 - $y = w_0 + w_1x_1 + w_2x_2 + \dots + w_dx_d$

Ridge Regression

- Add regularization term to OLS:

$$\min_{\mathbf{w}, b} \alpha \|\mathbf{w}\|^2 + \sum_{i=1}^N (y_i - f(\mathbf{x}_i))^2$$

- the first term is the *regularization term*
 - $\|\mathbf{w}\|^2 = \sum_{j=1}^d w_j^2$ penalizes large weights (aka L2-norm)
 - α is the hyperparameter that controls the amount of shrinkage
 - larger α means more shrinkage.
 - $\alpha = 0$ is the same as OLS.
- the second term is the *data-fit term*
 - sum-squared error of the prediction, same as linear regression.
- Also has a closed-form solution (similar derivation to linear regression):
 - $\mathbf{w}^* = (\mathbf{X}\mathbf{X}^T + \alpha I)^{-1} \mathbf{X}\mathbf{y}$
 - Similar to the solution for linear regression: $\mathbf{w}^* = (\mathbf{X}\mathbf{X}^T)^{-1} \mathbf{X}\mathbf{y}$
- What is the effect of the scaled identity matrix αI ?
- with linear regression, if \mathbf{X} does not span the input space \mathbb{R}^d , then $\mathbf{X}\mathbf{X}^T$ could be ill-conditioned or non-invertible.
 - i.e., we don't know how the data varies in the space orthogonal to \mathbf{X} .
- with ridge regression, the scaled identity conditions the matrix $\mathbf{X}\mathbf{X}^T$ so that the inverse can be computed.
 - (The term "ridge regression" comes from the closed-form solution, where a "ridge" is added to the diagonal of the covariance matrix)

Example on Boston data

```
In [23]: # randomly split data into 80% train and 20% test set
trainX, testX, trainY, testY = \
    model_selection.train_test_split(bostonX, bostonY,
                                     train_size=0.8, test_size=0.2, random_state=4487)

# normalize feature values to zero mean and unit variance
# this makes comparing weights more meaningful
# feature value 0 means the average value for that features
# feature value of +1 means one standard deviation above average
# feature value of -1 means one standard deviation below average
scaler = preprocessing.StandardScaler()
trainXn = scaler.fit_transform(trainX)
testXn = scaler.transform(testX)

print(trainXn.shape)
print(testXn.shape)

(101, 12)
(26, 12)
```

- vary α from 10^{-3} (little shrinkage) to 10^6 (lots of shrinkage)

```
In [24]: # alpha values to try
alphas = logspace(-3, 6, 50)

MSEs = empty(len(alphas))
```




```
ws = empty((len(alphas), trainXn.shape[1]))
for i,alpha in enumerate(alphas):
    # learn the RR model
    rr = linear_model.Ridge(alpha=alpha)
    rr.fit(trainXn, trainY)
    ws[i,:] = rr.coef_ # save weights

    MSEs[i] = metrics.mean_squared_error(testY, rr.predict(testXn))
```

- Effect...
 - for small α , all weights are non-zero.
 - for large α , all weights shrink to 0.
 - somewhere in between is the best model...

In [26]: rfig

Out[26]:



Selecting α using cross-validation

- built-in cross-validation (RidgeCV)

```
In [27]: # train RR with cross-validation
rr = linear_model.RidgeCV(alphas=alphas, cv=5)
rr.fit(trainXn, trainY)

MSE = metrics.mean_squared_error(testY, rr.predict(testXn))
print("MSE =", MSE)
print("alpha =", rr.alpha_)
print("w =", rr.coef_)

MSE = 55.24072984514399
alpha = 25.595479226995383
w = [-0.29070744  0.39475303 -0.87840328  0.67118967 -0.62982053  2.448381
      0.6
      -0.80391822 -1.41878524  0.27755795 -0.69521816 -1.75392601 -2.48098356]
```

Interpretation

- Which weights are most important?
 - look at weights with large magnitude.

```
In [28]: # print out sorted coefficients with descriptions
def print_coefs(coefs, bostonAttr):
```

```
# sort coefficients from smallest to largest, then reverse it
inds = argsort(abs(coefs))[:, :-1]
# print out
print("weight : feature description")
for i in inds:
    print("{: .3f} : {:7s} {}".format(coefs[i], bostonAttr[i][0], bostonAttr[i][1]))
```

- Which weights are most important?
 - negative weights indicate factors that decrease the house price
 - *Examples:* LSTAT (having higher percentage of lower status population), DIS (distance to business areas), PTRATIO (higher student-teacher ratio)
 - positive weights indicate factors that increase the house price
 - *Examples:* RM (having more rooms), RAD (proximity to highways)

In [29]:

```
print_coefs(rr.coef_, bostonAttr)
```

```
weight : feature description
-2.481 : LSTAT    % lower status of the population
 2.448 : RM       average number of rooms per dwelling
-1.754 : PTRATIO  pupil-teacher ratio by town
-1.419 : DIS      weighted distances to five Boston employment centres
-0.878 : INDUS    proportion of non-retail business acres per town
-0.804 : AGE      proportion of owner-occupied units built prior to 1940
-0.695 : TAX      full-value property-tax rate per $10,000
 0.671 : CHAS     Charles River dummy variable (= 1 if tract bounds river;
0 otherwise)
-0.630 : NOX      nitric oxides concentration (parts per 10 million)
 0.395 : ZN       proportion of residential land zoned for lots over 25,000
sq.ft.
-0.291 : CRIM     per capita crime rate by town
 0.278 : RAD      index of accessibility to radial highways
```

Better shrinkage

- With ridge regression, some weights are small but still non-zero.
 - these are less important, but somehow still necessary.
- To get better shrinkage to zero, we can change the regularization term to encourage more weights to be 0.
 - also called "sparse" weights, or encouraging "sparsity".

LASSO

- LASSO = "Least absolute shrinkage and selection operator"
- keep the same data fit term, but change the regularization term:
 - sum of absolute weight values: $\sum_{j=1}^d |w_j|$
 - also called L1-norm: $\|\mathbf{w}\|_1$
 - when a weight is close to 0, the regularization term can move the weight to be equal to 0.

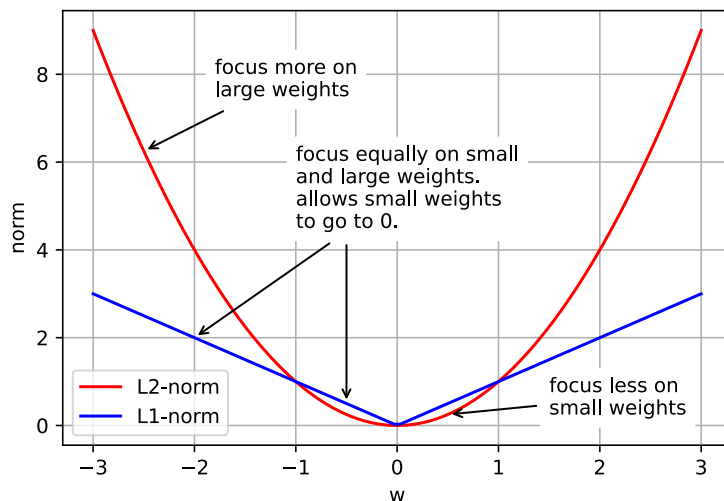
$$\min_{\mathbf{w}, b} \alpha \sum_{j=1}^d |w_j| + \sum_{i=1}^N (y_i - f(\mathbf{x}_i))^2$$

Comparison of L2 and L1 norms.

- L2 focuses more on large weights.
- L1 treats all weights equally.

In [31]: normfig


Out[31]:




Comparison of L2 and L1 norms

- During optimization with L1 norm:
 - for a given value of L1 norm, the minimal objective is usually in a "corner" of the L1 norm contour.
 - The "corner" corresponds to some weights equal to 0.

L1



L2



In [32]:

```
lasalphas = logspace(-3,2,50)

lassoMSEs = empty(len(alphas))
lassows = empty((len(alphas), trainXn.shape[1]))
```

```


for i,alpha in enumerate(lasalphas):
    # learn the LASSO model
    las = linear_model.Lasso(alpha=alpha)
    las.fit(trainXn, trainY)
    lassows[i,:] = las.coef_ # save weights

    lassoMSEs[i] = metrics.mean_squared_error(testY, las.predict(testXn))

```

In [34]: lfig

Out[34]:



Feature selection


- Select α to obtain a given number of features

```

In [35]: # count the number of non-zero weights
nzweights = sum(abs(lassows)>1e-6, axis=1)

plt.semilogx(lasalphas, nzweights, 'b.-')
plt.grid(True)
plt.xlabel('$\alpha$'); plt.ylabel('# non-zero weights');

```



```

In [36]: # get alpha where non-zero weights = 5
myi = where(nzweights==5)[0][0]
print("alpha=", lasalphas[myi])
print("MSE =", lassoMSEs[myi])
print("w =", lassows[myi,:])

```

```

alpha= 0.9102981779915218
MSE = 55.579375077684396
w = [-0.          0.         -0.19049149  0.         -0.         2.264558

```

```
08
-0.          -0.          -0.          -0.11944939 -1.76235357 -3.53224554]
```

Interpretation

- weights for unimportant features are set to 0
 - RAD, DIS, AGE, ...
- important features have non-zero weights
 - LSTAT, RM, PTRATIO, INDUS, TAX

```
In [37]: print_coefs(lassows[myi,:], bostonAttr)

weight : feature description
-3.532 : LSTAT    % lower status of the population
 2.265 : RM      average number of rooms per dwelling
-1.762 : PTRATIO pupil-teacher ratio by town
-0.190 : INDUS   proportion of non-retail business acres per town
-0.119 : TAX     full-value property-tax rate per $10,000
-0.000 : RAD     index of accessibility to radial highways
-0.000 : DIS     weighted distances to five Boston employment centres
-0.000 : AGE     proportion of owner-occupied units built prior to 1940
-0.000 : NOX     nitric oxides concentration (parts per 10 million)
 0.000 : CHAS    Charles River dummy variable (= 1 if tract bounds river;
0 otherwise)
 0.000 : ZN      proportion of residential land zoned for lots over 25,000
sq.ft.
-0.000 : CRIM    per capita crime rate by town
```

Cross-validation to select α

- Use built-in CV function
 - selects α with lowest error.

```
In [38]: # fit with cross-validation (alpha range is determined automatically)
las = linear_model.LassoCV()
las.fit(trainXn, trainY)

MSE = metrics.mean_squared_error(testY, las.predict(testXn))
print("MSE =", MSE)
print("alpha =", las.alpha_)
print("w =", las.coef_)

MSE = 56.05514450375403
alpha = 0.6426533625838364
w = [-0.          0.          -0.27751219  0.13119465 -0.          2.416023
96
-0.          -0.          -0.          -0.20770521 -1.89713883 -3.53302188]
```

Interpretation

- RAD, DIS, AGE, NOX, ZN, CRIM are unimportant features.

```
In [39]: print_coefs(las.coef_, bostonAttr)
```

```
weight : feature description
-3.533 : LSTAT    % lower status of the population
```

```

2.416 : RM      average number of rooms per dwelling
-1.897 : PTRATIO pupil-teacher ratio by town
-0.278 : INDUS  proportion of non-retail business acres per town
-0.208 : TAX    full-value property-tax rate per $10,000
0.131 : CHAS    Charles River dummy variable (= 1 if tract bounds river;
0 otherwise)
-0.000 : RAD    index of accessibility to radial highways
-0.000 : DIS    weighted distances to five Boston employment centres
-0.000 : AGE    proportion of owner-occupied units built prior to 1940
-0.000 : NOX    nitric oxides concentration (parts per 10 million)
0.000 : ZN      proportion of residential land zoned for lots over 25,000
sq.ft.
-0.000 : CRIM   per capita crime rate by town

```

Sparsity Constraints

- In previous formulations, LASSO and Ridge Regression only encourage sparsity using a regularizer.
- We can also formulate the regression problem with *explicit* sparsity constraints:

$$\min_{\mathbf{w}, b} \sum_{i=1}^N (y_i - f(\mathbf{x}_i))^2, \text{ s. t. } \|\mathbf{w}\|_0 \leq K$$

- L0-norm: $\|\mathbf{w}\|_0$ = the number of non-zero entries in \mathbf{w} .
 - (not really a norm)
- K is a hyperparameter - how many non-zero coefficients are desired.

Serious problem...

- LASSO and Ridge Regression are convex problems
 - Ridge Regression - closed-form solution
 - LASSO - efficient optimization algorithms to get exact solution
- Optimization problems with L0-norm constraints are NP-hard.
 - Combinatorial problem - all combinations of features need to be tried.

Orthogonal Matching Pursuit (OMP)

- **Idea:** greedy algorithm that iteratively selects the feature that is most correlated with the current residual error.
- Algorithm
 - Initialize the residual: $\mathbf{r} = \mathbf{y}$
 - For t in 1 to K
 - Find the most correlated feature: $j = \operatorname{argmax}_j |\mathbf{r}^T \mathbf{x}_j|$, where \mathbf{x}_j is the j -th row of \mathbf{X} (the j -th features).
 - Compute the weight: $w_j = \operatorname{argmin}_{w_j} \|\mathbf{r} - \mathbf{x}_j w_j\|^2$
 - Update the residual: $\mathbf{r} \leftarrow \mathbf{r} - \mathbf{x}_j w_j$

```
In [40]: # Example
omp = linear_model.OrthogonalMatchingPursuit(n_nonzero_coefs=2)
omp.fit(trainXn, trainY)

MSE = metrics.mean_squared_error(testY, omp.predict(testXn))
print("MSE =", MSE)
print(omp.coef_)
print(omp.intercept_)
```

```
MSE = 53.86974967354984
[ 0.          0.          0.          0.          0.          0.
  0.          0.          0.          0.         -2.62819004 -5.96007042]
22.85940594059405
```

```
In [41]: print_coefs(omp.coef_, bostonAttr)
```

```
weight : feature description
-5.960 : LSTAT    % lower status of the population
-2.628 : PTRATIO  pupil-teacher ratio by town
 0.000 : TAX      full-value property-tax rate per $10,000
 0.000 : RAD      index of accessibility to radial highways
 0.000 : DIS      weighted distances to five Boston employment centres
 0.000 : AGE      proportion of owner-occupied units built prior to 1940
 0.000 : RM       average number of rooms per dwelling
 0.000 : NOX      nitric oxides concentration (parts per 10 million)
 0.000 : CHAS     Charles River dummy variable (= 1 if tract bounds river;
0 otherwise)
 0.000 : INDUS    proportion of non-retail business acres per town
 0.000 : ZN       proportion of residential land zoned for lots over 25,000
sq.ft.
 0.000 : CRIM     per capita crime rate by town
```

- Note that LASSO selects different features, and also has worse MSE.

```
In [42]: # get alpha where non-zero weights = 2
myi = where(nzweights==2)[0][0]
print("MSE =", lassoMSEs[myi])
print_coefs(lassows[myi,:], bostonAttr)
```

```
MSE = 65.57831140414783
weight : feature description
-2.605 : LSTAT    % lower status of the population
 0.765 : RM       average number of rooms per dwelling
-0.000 : PTRATIO  pupil-teacher ratio by town
-0.000 : TAX      full-value property-tax rate per $10,000
-0.000 : RAD      index of accessibility to radial highways
 0.000 : DIS      weighted distances to five Boston employment centres
-0.000 : AGE      proportion of owner-occupied units built prior to 1940
-0.000 : NOX      nitric oxides concentration (parts per 10 million)
 0.000 : CHAS     Charles River dummy variable (= 1 if tract bounds river;
0 otherwise)
-0.000 : INDUS    proportion of non-retail business acres per town
 0.000 : ZN       proportion of residential land zoned for lots over 25,000
sq.ft.
-0.000 : CRIM     per capita crime rate by town
```