

What is Multimodality?

Letitia Parcalabescu
Computational Linguistics
Department
Heidelberg University
parcalabescu@cl.*

Nils Trost
Center for Molecular
Biology (ZMBH)
Heidelberg University
trost@zmbh.*

Anette Frank
Computational Linguistics
Department
Heidelberg University
frank@cl.*

Abstract

The last years have shown rapid developments in the field of multimodal machine learning, combining e.g., vision, text or speech. In this position paper we explain how the field uses outdated definitions of *multimodality* that prove unfit for the *machine learning* era. We propose a new *task-relative* definition of (multi)modality in the context of multimodal machine learning that focuses on representations and information that are *relevant* for a given machine learning task. With our new definition of multimodality we aim to provide a missing foundation for multimodal research, an important component of language grounding and a crucial milestone towards NLU.

1 Introduction

The holy grail of NLP is natural language understanding (NLU). As previously argued, NLU cannot be achieved by learning from text alone (Bender and Koller, 2020). Instead, an important step towards NLU is grounding language, especially in sight and sounds (Bisk et al., 2020). There is thus great interest in the field of NLP to go beyond the textual modality and to conduct multimodal machine learning (ML) research.

Multimodal ML has made great progress during the last years. Neural architectures are employed in tasks that go beyond single modalities. E.g., language is integrated with vision in Visual Question Answering (Antol et al., 2015), Visual Commonsense Reasoning (Zellers et al., 2019), Visual Dialogue (Das et al., 2017), or Phrase Grounding (Plummer et al., 2015). Audio signal processing has made advances in speech recognition (Nassif et al., 2019) and (visual) speech synthesis (Alam et al., 2020). But ML applications may reach beyond modalities that are familiar to us: Astronomi-



Figure 1: Are these examples instances of the same modality? = the same; \neq different. Depending on perspective, input data can be judged differently. Human- and machine-centered views would agree for (a) speech and text \neq , (b) images and text \neq . For (c) an image of text and text, the opinions could differ, while for (d) a visible light vs. infrared picture, humans could not even judge the infrared data, since it is not in their sensory capability.

cal and medical imaging techniques record wavelengths outside of what we call visible light. Genetic research measures signals alien to human perception, like activity and structure of molecules. Hence, we argue that current definitions of multimodality fall short of capturing the full space of multimodality in the ML era, and – even more seriously – that the field of multimodal ML, including vision and language integration, is lacking a proper definition of *multimodality*.

The point we aim to make in this position paper can be crystallized by asking the following questions regarding the input types shown on the left of Figure 1: *Are the data types shown on the left different from those that appear on the right, or are they instances of the same modality?* This question leads to different answers, depending on whether we take a human-centered as opposed to a machine-centered standpoint. In this position pa-

*uni-heidelberg.de

per we reason that either standpoint is insufficient, therefore we develop a definition of multimodality that allows researchers to judge cases as displayed in Figure 1, as involving or not different modalities from a *task-specific perspective*.

Our contributions are three-fold:

- i) We discuss existing definitions of multimodality and argue that they are insufficient to define multimodal tasks in the era of machine learning.
- ii) We propose a task-relative definition of multimodality and illustrate it for a variety of (multimodal) tasks and data. We argue that a task-relative definition better suits the current age of diversified data, where the field attempts to create multimodally enabled intelligent machines.
- iii) By providing a novel definition of multimodality, we offer a foundation for building a roadmap towards multimodality in NLP and beyond.

Multimodality is *a new challenge in NLP*. For building a roadmap towards Multimodality in NLP, we must establish common ground for what multimodal ML is, what are possible instantiations or specializations of multimodality, and what research questions it gives rise to: Pinning down what makes a task multimodal in terms of information diversity is one important research question; from deeper understanding of what makes a task multimodal, we may – conversely – derive novel tasks. Similarly, if we aim to extend the range of modalities in future multimodal tasks, how can we sensibly choose other modalities to ground language in, if we are not clear about what kinds of modalities language can represent itself, and about what constitutes a modality in the first place?

2 Multimodality and multimedia

The use of the terms *multimodality* and *multimedia* in academic literature and beyond is confusing at best. Works attempting to compare these two focus on different aspects further increase confusion of the matter. Lauer (2009) argues that in the context of classrooms and scholarship there is no difference in meaning between the terms, and claims that different terms are used for different audiences: multimedia being appreciated by the public and industry, while multimodality being preferred by teachers, since it emphasises the composing process. Anastopoulou et al. (2001) relates

multimedia to the ways of conveying information and multimodality to the sensory experience of the learning individual.

However, these works mentioned above and other notorious works on multimodality (Bezemer and Kress, 2008) have been focusing on the concepts of media and modality in relation to human society and its learning environments. For multimodal machine learning, the information receiver and processor is an ML system. To decide whether *multimodality* as used so far is the right term for the field of multimodal ML, we find the following definitions for “medium” in the Oxford Advanced Learner’s Dictionary:

- (a) “a way of communicating information, etc. to people” comprising text, (moving) images, graphs, sounds.
- (b) “a substance that something exists or grows in or that it travels through” being mainly applied in natural sciences.

“Modality” on the other hand is defined as follows:

- (A) “the particular way in which something exists, is experienced or is done”
- (B) “the kind of senses that the body uses to experience things” being applied in biology.

We argue that the term *multimodality* should be preferred in the context of ML, since multimodal ML aims to enhance world perception through ML systems (cf. def. (A)). Finding new ways of presenting information to humans (cf. def. (a)) is an endeavour that can benefit from multimodal ML while not being the main focus of multimodal ML research. For the rest of the paper, we will use the term “medium” or “media” in different contexts, referring to the meaning defined in (b).

In the following Section 3, we discuss how *multimodality* has been interpreted by the multimodal ML field and highlight the shortcomings of the corresponding definitions in Section 4.

3 How multimodality is defined – so far

In the multimodal ML literature and beyond, we find three general ways of defining “modality” or “multimodality”: not at all, or etymologically (bypassing the problem), or by way of a human-centered or a machine-centered definition¹.

¹For the scope of this paper, we disregard the statistical sense of “multimodality”, which describes a distribution with

Not at all or etymological Especially recent publications, as in Lu et al., 2020; Tan and Bansal, 2019; Gao et al., 2019, bypass a definition, assuming that the term is generally understood. Others offer an etymological definition: multimodal research involves not one, but *multiple* modalities (Zhang et al., 2020). Clearly, this definition leaves the notion of modality itself unexplained.

Human-centered Popular definitions of multimodality rely on the human **perceptual** experience as found in Baltrusaitis et al. (2019); Lyons (2016); Ngiam et al. (2011); Kress (2010). From this literature, we chose the following illustrative example because it focuses specifically on multimodality for ML, as is the interest of this paper:²

“Our experience of the world is multimodal – we see objects, hear sounds, feel texture, smell odors, and taste flavors. Modality refers to the way in which something happens or is experienced”. Baltrusaitis et al. (2019)

This view appeals to humans, who are bound to their senses when experiencing the world. It is thus an intuitive explanation of the concept of multimodality, focusing on the propagation channels that the human communication is adapted to (e.g., vision, sound).

Using this definition, one can agree for Fig. 1.a that speech (hearing) and text (seeing) are different modalities. But decisions are less clear for images and text as in Fig. 1.(b,c), as humans perceive both of them with their visual apparatus. Hence, as for written and depicted language, the human-centered definition *contradicts* the common conception in the community, that vision and language are different modalities, as in Lu et al. (2019); Su et al. (2019).

Machine-centered Another accepted perspective for defining multimodality is a machine-centered one, that focuses on the state in which information is transferred or encoded before being processed by a ML system:

“In the representation learning area, the word ‘modality’ refers to a particular way or mechanism of encoding information.” (Guo et al., 2019)

more than one peak. Such distributions can occur with any kind of data, unimodal or multimodal in the sense of “modality” we use for this paper

²Cf. section 4.1 for considerations about other human-centered formulations.

This definition is practical, focuses on the technical aspects of data representation, and captures how different types of inputs usually require specific programming solutions. For example, neural architectures typically use CNNs to encode images (exploiting $2d$ patches) and LSTMs to encode text (modeling sequential reading), exploiting the respective architecture’s inductive bias. From this viewpoint, the machine-centered definition naturally regards images vs. text as different modalities (cf. Fig. 1). However, recent developments in neural architectures are challenging this view, since multimodal transformers are processing both images and text with transformer blocks (Lu et al., 2020; Tan and Bansal, 2019; Su et al., 2019).

4 Why we need a better definition

In this section, we indicate shortcomings of the existing definitions and motivate why time is ripe for a new definition of multimodality.

4.1 Human-centered

The human-centered definition (Section 3) is rooted in research on how information is communicated to humans through, e.g., GUIs (Bernsen, 2008; Jovanovic, 2015; Lyons, 2016). However, this definition does not cover the large gap between the rich sensorial capabilities of humans as opposed to machines, and leaves open in which ways the signals that can be perceived by the variety of human senses will be converted to specific types of inputs that can be given to the ML system. Moreover, while we think that the human biology and psychology are and should be a valuable inspiration for ML research, we contend that ML should look far beyond the human, into other organisms³, and even purely physical phenomena.

We will discuss three aspects of the gap between humans and machines: (a) the types of inputs their sensors can detect, (b) the range of these inputs, and (c) the importance of the processor – a human or a machine – when processing the inputs.

(a) Input types Human senses are not prescriptive for the sensorial apparatus of machines. Humans, on one side, are limited to the senses nature has gifted them with: sight, hearing, touch, etc. Machines, on the other hand, are only limited by the creativity of physicists or engineers: machines read tapes by measuring magnetization, or DVDs

³e.g., a self-driving car imitating a nematode’s nervous system, as in Lechner et al., 2020

by detecting microscopic engravings. What modalities would the human-centered definition assign to these signals? Also, machines can surpass the human sensorial capabilities, by performing more exact measurements of e.g., humidity or pressure. This may not seem too relevant when restricting multimodality to the context of NLU, as language was created by humans to fit our experience of the world. In other fields however, ML systems are used on signals that are completely outside of the perception of humans, e.g., predicting gene regulation based on chromatin accessibility and transcription in biology (Minnoye et al., 2020).

(b) Input range Biological sensory detection systems are restricted to specific ranges of signals, and thus impose unnecessary limitations when applied to machines (Bernsen, 2008). Machines are limited only by the borders of human engineering⁴ and use manifold materials and physical phenomena that biological organisms adapted to their specific environment do not. For example, humans can detect and interpret only a tiny part of the electromagnetic spectrum (380 to 700 nanometers, Starr et al., 2010) and call it *visible light*. But machines can detect and (if programmed) process the whole electromagnetic spectrum, far beyond the visible light. Humans cannot perceive infrared – and hence this modality is non-existent if we adopt the human-centered definition. This again mainly impacts fields apparently remote to NLP, e.g., medicine, where imaging techniques are employed that measure wavelengths far outside of the perceptive range of humans. But we desire that future systems can combine their experience of e.g., both the visible light and other wavelengths with natural language (to aid medical diagnose, for example).

(c) The processor: a human or an ML system? Humans have the innate ability of seeing objects, hearing sounds and delivering *some* interpretation of these signals. Machines, by contrast, can be clever sensors and count photons on a semiconductor or measure air pressure oscillations – yet without special programming they cannot interpret or derive information from the inputs. Detected physical quantities are then mapped to a voltage sequence interpreted as 0 if the voltage does not surpass a threshold, 1 if it does.

Since ultimately, behind all data encodings, there are just 0s and 1s waiting to be interpreted by a

⁴Engineering solutions which can be biologically inspired.

program, we argue that multimodal ML research should focus on these programs and that a definition of multimodality should answer the question: What are the challenges that a program needs to address when it is exposed to a new modality rather than more unimodal data? In humans, evolution has already addressed this question, by specializing sensory organs and brain areas to the processing and interpretation of various input types (Schlosser, 2018). Multimodal ML is not there yet.

Language – human’s stroke of genius It is generally accepted in the multimodal ML literature that vision is one modality and language the other, typically given in form of text (Kafle et al., 2019). But humans *hear* speech i.e., spoken language, *read* written language with their visual apparatus, *see* signed languages, or *feel* Braille. The upside of the human-centered definition is that it captures the plurality of media that support the transmission of language and accordingly, it can assign different modalities to such different manifestations of language.

However, we are concerned with multimodality in ML and there are important edge cases for which the human-centered definition is inadequate: Does a screenshot of a text editor that displays the content of a .txt file containing a sentence *s* represent a different modality than the data encoding in the text file? For a human, it is the same visual modality, since in both cases *s* is visually perceived. But a machine needs very different programming in order to extract the same information about sentence *s* from an *image* vs. the *ASCII encoding* of *s* stored in a .txt file.

Having raised this criticism of the human-centered view on multimodality, it seems like the machine-centered view, focusing on the encoding of information, can offer a more viable interpretation of multimodality.

4.2 Machine-centered

Data representation While data representations are a challenge (Bengio et al., 2013), we argue that representations themselves should not be the defining trait for multimodal learning: For example, if we follow the machine-centered definition, an undirected unlabeled graph and its adjacency matrix have to be considered different modalities because they are different ways of representing a network. Similarly, PNG and JPEG are possible encodings of the same image and hence would be

considered different modalities. This interpretation seems unintuitive for tasks like image recognition, where the image format does not play a big role and is usually homogenized in pre-processing.

Still, there are applications where the 4th dimension of PNGs is useful for encoding transparency or depth maps, by using the additional channel that PNG has compared to JPEG. Here we stand with a puzzle: Should data representation matter or not?

When does representation impact information?

Speech and handwriting are propagated through different media⁵ and are represented differently in a computer⁶. Does this make them *different* modalities? Both speech and writing can be propagated through the same digital medium (fiber-optic cable) and if speech has been transcribed to text in pre-processing, they can have the same (ASCII) encoding. Does this make them the *same* modality?

When converting between media and representations, information can get lost. But without knowing the multimodal ML task and the information it requires, we cannot decide whether the loss was noteworthy or not. Below, we will argue that crucial factors for defining multimodality are not only an efficient encoding of information, but also the ML task itself.

5 A task-relative definition

We argue that a definition of multimodality for multimodal ML should relate to the task an ML system needs to solve, since it is the task that determines what **information** is relevant and in which **representation** this information can be efficiently stored. The human- and machine-centered definitions try to capture the essence of multimodality in a task-agnostic manner, relating it to categories of human experience, media, representations, and encodings. As shown in Section 4, these definitions turn out to be insufficient in view of the plurality of physical sensors, tasks and data types of the present.

Instead, our task-relative definition aims to answer the question: Under which conditions does a multimodal setting reveal crucial insight *for the given task* – and do we need multiple modalities? In our view, (i) different inputs can contribute specific information, but (ii) what is relevant information can only be determined in relation to the task at hand; and only by taking the task into account

⁵e.g. air for speech, ink trail on paper for handwriting

⁶time series of amplitudes for speech, images or textual transcriptions for handwriting




Figure 2: Our definition of *multimodality* determines the modalities of input channels by considering i) how each input channel is *represented*, ii) whether the *information* units each input carries is complementary to each other iii) *in relation to the ML task*.

(iii) we can determine the status of the inputs as (possibly complementary) modalities.

5.1 Task-relative definition of multimodality

We propose the following **task-relative definition** of multimodality in ML that relates *representation*, *information* and *task* as depicted in Figure 2.

*A machine learning **task** is multimodal when inputs or outputs are **represented** differently or are composed of distinct types of **atomic units of information**.*

Since the definition focuses on multimodal ML, the input representation of interest is the direct input to the learning system; any prior representational steps that are not seen, or filtered by pre-processing, are irrelevant for our definition.

Atomic units of information *Information* is one of the three dimensions we use to decide multimodality. Hence we ask: What information does the input carry and is relevant for the task? As a defining criterion on whether two input channels contribute information from different modalities, we consider what types of atomic units of information they provide: they are of different types if they cannot be captured by a 1-to-1 mapping between their domains. In this case, we speak of multiple modalities.

In other words, if we find that **i)** after pre-processing, inputs x_1 and x_2 are represented the same, we examine whether they live in the same modality or not by asking whether x_1 and x_2 are formed of different atomic units: If **ii)** we can't establish a bijective mapping between the domains of x_1 and x_2 , then x_1 and x_2 must be composed of different atomic units, and if **i)** is the case, they represent different modalities.

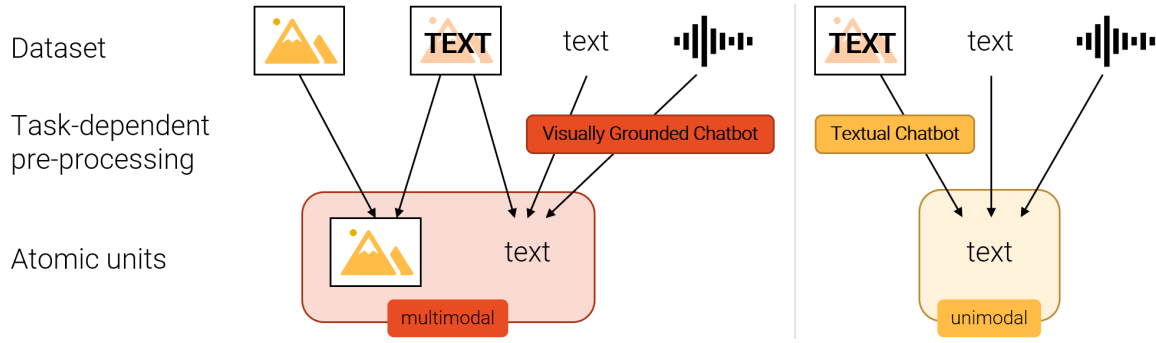


Figure 3: The task defines the required atomic units of information for solving it. A task that requires multiple atomic units after pre-processing is multimodal.

Similarly, if **iii**) after pre-processing, x_1 and x_2 are represented *differently*, the task is multimodal.

The bijectivity criterion is lenient towards a task-dependent error for real-world applicability. Regardless of multimodal characteristics, low data quality and compression can cause loss of task-relevant information.

We thus speak of a new modality when it contributes information that cannot be delivered by larger (but not infinite) amounts of unimodal data. Note that the same information captured in one modality may be encoded in a different modality, however, not necessarily with the same efficiency: We can, in infinite time, describe every minute detail of a landscape unimodally through language. But it is clearly more efficient to capture the minute details of a landscape in a different modality, e.g., a photograph. In general, any kind of information can be reduced to a string of 1s and 0s, yet, depending on the information source and the given task, another representation might be more convenient.

Task-specificity The *task* is the second crucial characteristic of multimodality because the task determines (i) what information of the input channels is necessary to solve the task as discussed above and (ii) determines two crucial components: *input and output* and how the latter depends on the former. Regarding inputs and their possible (multi)modality, we critically focus on the nature of the data that forms the *direct input* to the ML system, *disregarding* any data representations that the input may take during any pre-processing stages. Hence, we draw a sharp line between input encodings and representations that are part of the learning process, and that will be continuously refined – as opposed to input formats that are external to the learning process.

For our task-dependent definition of multimodal-

ity it will therefore not matter whether the original input was speech, but was transcribed to text *if*, for example, the task at hand does not require information that is specific to spoken language and can be conveyed by a textual representation of spoken input (see Figure 3). Also for image recognition it is not crucial whether images come as PNG or JPEG, and harmonizing the data is a matter of pre-processing.

Atomic units vs. data representation The term *data representation*, as we use it here, refers to the encoding of information and data formats. However, *atomic units* are not bound by their technical implementation – they constitute the informational content in *relation* to the task. For example, PNG and JPEG are undeniably different data representations which can, depending on the task, represent (a) two *different* types of atomic units of information, if the additional dimension in PNG is important for the task (e.g., in view of encoding transparency, or depth), or (b) a *single* type of atomic units, because a task-specific bijective mapping can be established that does not lose (task-relevant) information.

Not sensor-specific Our definition is especially robust to changes of medium and representation of information via data transfer or storage. We can thus neglect the physical or biological sensors that capture the data, and the encoding, transmission or storage of the data until it reaches the processor. This property makes our task-relative definition a good working definition: the constant change of representation that information may undergo, does not immediately span a new modality.

5.2 Applying the task-relative definition

We now apply our definition of multimodality to various examples – including edge cases – in order to demonstrate its breadth, flexibility and robustness.

Images vs. text By our definition, tasks working with images (stacks of intensity matrices) and text in ASCII format are multimodal, because they consist of different atomic units of information: An image can be truthfully described by multiple textual descriptions, and similarly, text translates equivocally to images (e.g., pictures of different hand writings). The decision becomes less clear-cut, if we are given two inputs: natural images and images of text. Both inputs are intensity matrices and therefore unimodal. However, if the task does not consider the differences in hand writing style and applies Optical Character Recognition to obtain a uniform text representation, the *images of text* turn into *text* in pre-processing, and the task becomes multimodal.

Images of infrared vs. visible light Infrared images and visible light images are represented uniformly (a stacked grid of photon counts of different energy intervals). Additionally, we can define a bijective mapping between the two, e.g. by adding/subtracting a fixed frequency and thereby shifting the infrared image into a visible light image (as used for night vision). Therefore, by our definition, they represent the same modality. In this way, our definition enables us to define modalities for information not directly perceptible by humans because of sensory limitations. For example, the photons of infrared light do not carry the correct amount of energy to be perceptible by our eyes, as it is the case with most of the electromagnetic spectrum.

Language Finally, our definition captures two key traits of the multimodal nature of language: (i) coming in many forms (speech, handwriting, signed language, ASCII-code), language constantly switches representations and media. In cases where (ii) after pre-processing different language representations cannot be converted one to the other without losing task-relevant information (e.g., intonation, hesitation, modulation), they become multiple modalities, like speech–text, or handwriting–text, etc.

For example, using our definition, languages like English and Japanese are considered to be unimodal, if after pre-processing both are represented

in Unicode. If not, handling them becomes a multimodal task. This behavior relates the essence of multimodal ML and multilingual NLP, in terms of their complementarity: There are concepts in some languages that cannot be efficiently translated to other languages; much like humans cannot conceive how a bee sees ultra-violet light (Chittka and Wells, 2004).

5.3 Relevant aspects of multimodality

Similarly to previous definitions of multimodality, we focus on the representation and content of inputs, but we consider them in *relation* to the *task* (see Figure 2):

- (i) Starting from the *task*, we capture *what information is put to use* and relevant for solving it. Hence, despite the presence of e.g. video and language in a dataset, the task may be such that one of them is not relevant and could be fully ignored.
- (ii) We focus on the *representation* of the information that forms the direct input to the ML system after pre-processing. Unlike previous definitions, we reduce the importance of the medium in multimodal ML. The medium may change without modifying the nature of the *information* that is represented, used and needed for the task.
- (iii) Finally and crucially, we determine the uni- vs. multi-modality of the information joined to solve a task, by analyzing its lack of bijectivity, i.e. its *complementarity*.

Hereby we conclude that it is only in the context of a task and data used to solve it that we can analyze the issue of information relevance and bijectivity/complementarity, i.e. of *multimodality*. In this sense we argue that there is no multimodal input or data *per se*; it is only through a task that requires multimodal information sources to be solved that the corresponding inputs can be truly considered multimodal.

6 Conclusion

In this paper we have shown how human-centered definitions of multimodality are too restrictive for current and future developments of ML. Also, machine-centered definitions focusing on representations only do not capture the crucial trait of multimodal machine learning. We instead propose a new definition of multimodality that focuses on the

relation between representations, information and the given task, and that – through the novel dimension of the *task* – enables us to make much sharper distinctions compared to current standards, while covering a much wider space of multimodal inputs.

With this position paper, we (a) offer a working definition on how to use the term *multimodality*, (b) aim to raise awareness that defining multimodality is harder than expected, and (c) invite the community to *discuss* these challenges and (why not?) to provide a better definition.

7 Ethical considerations

In the present paper, we portray the opinion that the human experience of the world should not be normative for defining the multimodal characteristic of a machine learning system. Instead, we claim that machine learning should draw inspiration from human biology and psychology, but not limit itself by imitation. Independently of (a) where the design of machine learning systems is inspired from, and (b) their capabilities, which may extend or exceed those of humans, we strongly believe that machine learning should be done for humans' service, following the ethical considerations developed and accepted by human society.

References

- Mahbubul Alam, Manar D Samad, Lasitha Vidyaratne, Alexander Glandon, and Khan M Iftekharuddin. 2020. Survey on deep neural networks in speech and vision systems. *Neurocomputing*, 417:302–321.
- S Anastopoulou, C Baber, and M Sharples. 2001. Multimedia and multimodal systems: commonalities and differences. In *5th Human Centred Technology Postgraduate Workshop, University of Sussex*.
- Stanislaw Antol, Aishwarya Agrawal, Jiasen Lu, Margaret Mitchell, Dhruv Batra, C Lawrence Zitnick, and Devi Parikh. 2015. Vqa: Visual question answering. In *Proceedings of the IEEE international conference on computer vision*, pages 2425–2433.
- Tadas Baltrusaitis, Chaitanya Ahuja, and Louis-Philippe Morency. 2019. [Multimodal machine learning: A survey and taxonomy](#). *IEEE Trans. Pattern Anal. Mach. Intell.*, 41(2):423–443.
- Emily M. Bender and Alexander Koller. 2020. [Climbing towards NLU: On meaning, form, and understanding in the age of data](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5185–5198, Online. Association for Computational Linguistics.
- Yoshua Bengio, Aaron Courville, and Pascal Vincent. 2013. Representation learning: A review and new perspectives. *IEEE transactions on pattern analysis and machine intelligence*, 35(8):1798–1828.
- Niels Ole Bernsen. 2008. Multimodality theory. In *Multimodal User Interfaces*, pages 5–29. Springer.
- Jeff Bezemer and Gunther Kress. 2008. Writing in multimodal texts: A social semiotic account of designs for learning. *Written communication*, 25(2):166–195.
- Yonatan Bisk, Ari Holtzman, Jesse Thomason, Jacob Andreas, Yoshua Bengio, Joyce Chai, Mirella Lapata, Angeliki Lazaridou, Jonathan May, Aleksandr Nisnevich, et al. 2020. Experience grounds language. *arXiv preprint arXiv:2004.10151*.
- Lars Chittka and Harrington Wells. 2004. Color vision in bees: mechanisms, ecology and evolution. *Complex worlds from simpler nervous systems*, pages 165–191.
- Abhishek Das, Satwik Kottur, Khushi Gupta, Avi Singh, Deshraj Yadav, José M.F. Moura, Devi Parikh, and Dhruv Batra. 2017. Visual Dialog. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*.
- Peng Gao, Haoxuan You, Zhanpeng Zhang, Xiaogang Wang, and Hongsheng Li. 2019. Multi-modality latent interaction network for visual question answering. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 5825–5835.
- Wenzhong Guo, Jianwen Wang, and Shiping Wang. 2019. Deep multimodal representation learning: A survey. *IEEE Access*, 7:63373–63394.
- Danica Jovanovic. 2015. Key terms in multimodality: Definitions, issues, discussions: Entries on: Important people in multimodality theory; information value; connotation/denotation; perspective.
- Kushal Kafle, Robik Shrestha, and Christopher Kanan. 2019. Challenges and prospects in vision and language research. *Frontiers in Artificial Intelligence*, 2:28.
- Gunther R Kress. 2010. *Multimodality: A social semiotic approach to contemporary communication*. Taylor & Francis.
- Claire Lauer. 2009. [Contending with terms: “multimodal” and “multimedia” in the academic and public spheres](#). *Computers and Composition*, 26(4):225–239.
- Mathias Lechner, Ramin Hasani, Alexander Amini, Thomas A Henzinger, Daniela Rus, and Radu Grosu. 2020. Neural circuit policies enabling auditable autonomy. *Nature Machine Intelligence*, 2(10):642–652.

- Jiasen Lu, Dhruv Batra, Devi Parikh, and Stefan Lee. 2019. Vilbert: Pretraining task-agnostic visiolinguistic representations for vision-and-language tasks. In *Advances in Neural Information Processing Systems*, pages 13–23.
- Jiasen Lu, Vedanuj Goswami, Marcus Rohrbach, Devi Parikh, and Stefan Lee. 2020. 12-in-1: Multi-task vision and language representation learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 10437–10446.
- Agnieszka Lyons. 2016. 18 multimodality. *Research Methods in Intercultural Communication*, page 268.
- Liesbeth Minnoye, Ibrahim Ihsan Taskiran, David Mauduit, Maurizio Fazio, Linde Van Aerschot, Gert Hulselmans, Valerie Christiaens, Samira Makhzami, Monika Seltenshammer, Panagiotis Karras, et al. 2020. Cross-species analysis of enhancer logic using deep learning. *Genome research*, 30(12):1815–1834.
- Ali Bou Nassif, Ismail Shahin, Imtinan Attili, Mohammad Azzeh, and Khaled Shaalan. 2019. Speech recognition using deep neural networks: A systematic review. *IEEE Access*, 7:19143–19165.
- Jiquan Ngiam, Aditya Khosla, Mingyu Kim, Juhan Nam, Honglak Lee, and Andrew Y Ng. 2011. Multimodal deep learning. In *ICML*.
- Bryan A Plummer, Liwei Wang, Chris M Cervantes, Juan C Caicedo, Julia Hockenmaier, and Svetlana Lazebnik. 2015. Flickr30k entities: Collecting region-to-phrase correspondences for richer image-to-sentence models. In *Proceedings of the IEEE international conference on computer vision*, pages 2641–2649.
- Gerhard Schlosser. 2018. A short history of nearly every sense—the evolutionary history of vertebrate sensory cell types. *Integrative and comparative biology*, 58(2):301–316.
- Cecie Starr, Christine Evers, and Lisa Starr. 2010. *Biology: concepts and applications*. Cengage Learning.
- Weijie Su, Xizhou Zhu, Yue Cao, Bin Li, Lewei Lu, Furu Wei, and Jifeng Dai. 2019. Vi-bert: Pre-training of generic visual-linguistic representations. In *International Conference on Learning Representations*.
- Hao Tan and Mohit Bansal. 2019. Lxmert: Learning cross-modality encoder representations from transformers. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 5103–5114.
- Rowan Zellers, Yonatan Bisk, Ali Farhadi, and Yejin Choi. 2019. From recognition to cognition: Visual commonsense reasoning. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*.
- Chao Zhang, Zichao Yang, Xiaodong He, and Li Deng. 2020. Multimodal intelligence: Representation learning, information fusion, and applications. *IEEE Journal of Selected Topics in Signal Processing*.