



# Chapter 3: Introduction to SQL

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Outline

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database



## 3.1 SQL查询语言概览

### ■ History

- IBM Sequel language developed as part of System R project at the IBM San Jose Research Laboratory
- Renamed Structured Query Language (SQL)
- ANSI and ISO standard SQL:
  - ▶ SQL-86
  - ▶ SQL-89
  - ▶ SQL-92
  - ▶ SQL:1999 (language name became Y2K compliant!)
  - ▶ SQL:2003
- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.



# SQL语言有以下几个部分:

- 数据定义语言(Data- Definition Language, DDL): SQL DDL提供定义关系模式、删除关系以及修改关系模式的命令。
- 数据操纵语言( Data- Manipulation Language, DML): SQL DML提供从数据库中查询信息, 以及在数据库中插入元组、删除元组、修改元组的能力。



- 完整性(integrity): SQL DDL包括**定义完整性约束**的命令, 保存在数据库中的数据**必须满足**所定义的完整性约束。破坏完整性约束的更新是**不允许**的。
- 视图定义(view definition): SQL DDL包括**定义视图**的命令。
- 事务控制(transaction control): SQL包括定义事务的**开始和结束**的命令。



- 嵌入式SQL和动态SQL( embedded SQL and dynamic SQL): 嵌入式和动态SQL定义SQL语句如何嵌入到通用编程语言, 如C、C++和Java中。
- 授权(authomation): SQL DDL包括定义对关系和视图的访问权限的命令。



## 3.2 SQL Data Definition Language

- 数据库中的关系集合必须由数据定义语言(DDL)指定给系统。
- The SQL data-definition language (DDL) allows the specification of information about relations, including:
  - The schema for each relation.
  - The domain of values associated with each attribute.
  - Integrity constraints



- also other information such as
  - ▶ The set of indices to be maintained for each relations.
  - ▶ Security and authorization information for each relation.
  - ▶ The physical storage structure of each relation on disk.





## 3.2.1 Domain Types in SQL

- **char(*n*)**. Fixed length character string, with user-specified length *n*.
- **varchar(*n*)**. Variable length character strings, with user-specified maximum length *n*.
- **int**. Integer (a finite subset of the integers that is machine-dependent).
- **smallint**. Small integer (a machine-dependent subset of the integer domain type).
- **numeric(*p*,*d*)**. Fixed point number, with user-specified precision of *p* digits, with *d* digits to the right of decimal point. (ex., **numeric**(3,1), allows 44.5 to be stored exactly, but not 444.5 or 0.32)



- **real, double precision.** Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(*n*).** Floating point number, with user-specified precision of at least *n* digits.



- **空值Null**表示一个**缺失**的值，该值可能**存在**但**并不为人所知**，或者可能根本**不存在**。
- **char**数据类型存放**固定长度**的字符串。
  - 例如，属性A的类型是**char(10)**。
    - ▶ 如果我们为此属性存入字符串“**Avi**”，那么该字符串后会**追加7个空格**来使其达到10个字符的串长度。
    - ▶ **varchar(10)**，则**不会**增加空格。
- 比较一个char类型和一个varchar类型的时候，如果**长度不同**，比较结果将是**假**。



- 议始终使用varchar类型而不是char类型来避免这样的问题。



## 3.2.2 Create Table Construct

- An SQL relation is defined using the **create table** command:

**create table**  $r$  ( $A_1 D_1, A_2 D_2, \dots, A_n D_n,$   
     $\langle \text{integrity-constraint}_1 \rangle,$   
     $\dots,$   
     $\langle \text{integrity-constraint}_k \rangle$ );

- $r$  is the name of the relation
- each  $A_i$  is an attribute name in the schema of relation  $r$
- $D_i$  is the data type of values in the domain of attribute  $A_i$



■ Example:

```
create table instructor (  
    ID           char(5),  
    name        varchar(20),  
    dept_name   varchar(20),  
    salary     numeric(8,2))
```



# Integrity Constraints in Create Table

- **not null**
- **primary key**  $(A_1, \dots, A_n)$
- **foreign key**  $(A_m, \dots, A_n)$  **references**  $r$
- **primary key** declaration on an attribute automatically ensures **not null**



■ *Example:*

■ **create table** *instructor* (  
                          *ID*                  **char**(5),  
                          *name*             **varchar**(20) **not null**,  
                          *dept\_name* **varchar**(20),  
                          *salary*          **numeric**(8,2),  
                          **primary key** (*ID*),  
                          **foreign key** (*dept\_name*)  
          **references** *department*);

■ 书P35页





## And a Few More Relation Definitions

■ **create table** *student* (  
    *ID*                    **varchar(5),**  
    *name*                **varchar(20) not null,**  
    *dept\_name*        **varchar(20),**  
    *tot\_cred*           **numeric(3,0),**  
    **primary key** (*ID*),  
    **foreign key** (*dept\_name*) **references**  
    *department*);



- **create table** *takes* (  
    *ID*                    **varchar**(5),  
    *course\_id*        **varchar**(8),  
    *sec\_id*            **varchar**(8),  
    *semester*        **varchar**(6),  
    *year*             **numeric**(4,0),  
    *grade*            **varchar**(2),  
    **primary key** (*ID*, *course\_id*, *sec\_id*, *semester*,  
*year*) ,  
    **foreign key** (*ID*) **references** *student*,  
    **foreign key** (*course\_id*, *sec\_id*, *semester*, *year*)  
**references** *section*);
- Note: *sec\_id* can be dropped from primary key above, to ensure a student cannot be registered for two sections of the same course in the same semester



# Updates to tables

## ■ Insert

- **insert into** *instructor* **values** ('10211', 'Smith', 'Biology', 66000);

## ■ Delete

- Remove all tuples from the *student* relation
  - ▶ **delete from** *student*

## ■ Drop Table

- **drop table** *r*



## ■ Alter

### ● **alter table $r$ add $A$ $D$**

- ▶ where  $A$  is the name of the attribute to be added to relation  $r$  and  $D$  is the domain of  $A$ .
- ▶ All exiting tuples in the relation are assigned *null* as the value for the new attribute.

### ● **alter table $r$ drop $A$**

- ▶ where  $A$  is the name of an attribute of relation  $r$
- ▶ Dropping of attributes not supported by many databases.



## 3.3 Basic Query Structure

- A typical SQL query has the form:

**select**  $A_1, A_2, \dots, A_n$   
**from**  $r_1, r_2, \dots, r_m$   
**where**  $P$

- $A_i$  represents an attribute
  - $R_i$  represents a relation
  - $P$  is a predicate.
- The result of an SQL query is a relation.



# The select Clause

- The **select** clause **lists** the attributes desired in the result of a query
  - corresponds to the **projection** operation of the relational algebra
- Example: find the names of all instructors:  

**select** *name*  
**from** *instructor*
- NOTE: SQL names are **case insensitive** (i.e., you may use upper- or lower-case letters.)
  - E.g., *Name*  $\equiv$  *NAME*  $\equiv$  *name*
  - Some people use upper case wherever we use bold font.



## The select Clause (Cont.)

- SQL allows **duplicates** in relations as well as in query results.
- To force the **elimination** of duplicates, insert the keyword **distinct** after select.
- Find the department names of all instructors, and remove duplicates

```
select distinct dept_name  
from instructor
```

- The keyword **all** specifies that **duplicates** should **not be removed**.

```
select all dept_name  
from instructor
```



## The select Clause (Cont.)

- An **asterisk** in the select clause denotes “**all attributes**”

**select \***  
**from** *instructor*

- An attribute can be a literal with **no from clause**

**select** ‘437’

- Results is a table with one column and a single row with value “437”
- Can give the column a name using:

**select** ‘437’ **as** *FOO*





- An attribute can be a literal with **from** clause

**select** 'A'  
**from** *instructor*

- Result is a table with one column and  $N$  rows (number of tuples in the *instructors* table), each row with value “A”



## The select Clause (Cont.)

- The **select** clause can contain **arithmetic expressions** involving the operation, **+**, **-**, **\***, **and** **/**, and operating on **constants** or **attributes** of tuples.

- The query:

```
select ID, name, salary/12  
from instructor
```

would return a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12.

- Can **rename** “*salary/12*” using the **as clause**:

```
select ID, name, salary/12 as monthly_salary
```



# The where Clause

- The **where** clause specifies **conditions** that the result must satisfy
  - Corresponds to the **selection predicate** of the relational algebra.
- To find all instructors in *Comp. Sci. dept*

```
select name  
from instructor  
where dept_name = 'Comp. Sci.'
```



- Comparison results can be combined using the logical connectives **and, or, and not**

- To find all instructors in Comp. Sci. dept with salary > 80000

```
select name
from instructor
where dept_name = 'Comp. Sci.' and salary
> 80000
```

- Comparisons can be applied to **results of arithmetic expressions.**
- 逻辑连词的**运算对象**可以是包含**比较运算符**<、<=、>、>=、=和<>的**表达式**。



## 3.3.2 The from Clause

- The **from** clause lists the **relations** involved in the query
  - Corresponds to the **Cartesian product operation** of the relational algebra.
- Find the Cartesian product *instructor X teaches*

**select \***

**from** *instructor, teaches*

- generates **every possible** instructor – teaches pair, with all attributes from both relations.



- For **common** attributes (e.g., *ID*), the attributes in the **resulting** table are **renamed** using the relation name (e.g., *instructor.ID*)
- Cartesian product **not very useful** directly, but **useful** combined with ***where-clause*** condition (selection operation in relational algebra).



# Cartesian Product

*instructor*

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

*teaches*

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

Inst.ID	name	dept_name	salary	teaches.ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	12121	FIN-201	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	15151	MU-199	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
12121	Wu	Finance	90000	10101	CS-347	1	Fall	2009
12121	Wu	Finance	90000	12121	FIN-201	1	Spring	2010
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2010
12121	Wu	Finance	90000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...



# Examples

- Find the names of all instructors who have taught some course and the course\_id
  - **select** *name, course\_id*  
**from** *instructor, teaches*  
**where** *instructor.ID = teaches.ID*
- Find the names of all instructors in the Art department who have taught some course and the course\_id
  - **select** *name, course\_id*  
**from** *instructor, teaches*  
**where** *instructor.ID = teaches.ID and instructor.dept\_name = 'Art'*





### 3.3.3 自然连接

- **自然连接**(natural join) 运算作用于两个关系，并产生一个关系作为结果。
- **不同于**两个关系上的笛卡儿积，
  - **笛卡儿积**将第一个关系的**每个**元组与第二个关系的所有元组都**进行连接**；
  - 自然连接**只考虑**那些在**两个关系**模式中都出现的属性上**取值相同**的元组对。



- 自然连接没有重复列出那些在两个关系模式中都出现的属性，这样的属性只出现一次。
- 自然连接列出属性的顺序：
  - 先是两个关系模式中的共同属性，
  - 然后是那些只出现在第一个关系模式中的属性，
  - 最后是那些只出现在第二个关系模式中的属性。



● 例如：instructor关系和teaches关系的自然连接

n     **select** name, course\_id

n     **from** instructor, teaches

n     **where** instruclor.ID=teaches.ID:

● 该查询可以用SQL的自然连接运算更简洁地写作：

n     **select** name, course\_id

n     **from** instructor **natural join** teaches;

● 以上两个查询产生相同的结果。



- 在一个SQL查询的from子句中，可以用自然连接将多个关系结合在一起，如下所示：
- **select A,, Az,... A.**
- **from r1 natural join r2 natural join...natural join rm**
- **where P;**
- 更为一般地，from子句可以为如下形式：
  - **from E1, E2,..., E。**
  - 其中每个E可以是单个关系或一个包含自然连接的表达式。



- 例如：查询“列出教师的名字以及他们所讲授课程的名称”
  - **select** name, title
  - **from** instructor **natural join** teaches, course
  - **where** teaches.course id= course.course id;
- 注意： 下面的SQL查询不会计算出相同的结果：
  - **select** name, title
  - **from** instructor **natural join** teaches **natural join** course;



## ■ 原因是：

- instructor和teaches的自然连接包括属性（ID， name， *dept\_name*， salary. *course\_id*， sec\_id），
- 而course关系包含的属性是（*course\_id*， title， *dept\_name*， credits）。
- 二者自然连接的结果，是需要来自这两个输入的元组既要在属性*dept\_name*上取值相同，还要在*course\_id*上取值相同。



- SQL提供了一种自然连接的构造形式，允许用户来指定需要哪些列相等。
  - **select** name, title
  - **from** (instructor **natural join** teaches) **join** course **using** (course\_id);
- **join... using**运算中需要给定一个属性名列表，其中，两个输入中都必须具有指定名称的属性。



- 运算  $r1 \text{ join } r2 \text{ using}(A1, A2)$ ,
- 它与  $r1$  和  $r2$  的自然连接类似, 只不过
  - $r1.A1 = t2.A1$ , 并且  $t1.A2 = t2.A2$
- 即使  $r$  和  $r2$  都具有名为  $A3$  的属性, 也不需要  $r1$  .  $A3 = r2.A3$  成立。





## 3.4 附加的基本运算

### 3.4.1 The Rename Operation

- The SQL allows renaming **relations** and **attributes** using the **as** clause:

*old-name as new-name*

- Find the names of all instructors who have a higher salary than some instructor in ‘Comp. Sci.’.

- **select distinct** *T.name*  
**from** *instructor as T, instructor as S*  
**where** *T.salary > S.salary* **and** *S.dept\_name =*  
*‘Comp. Sci.’*

- Keyword **as** is **optional** and may be **omitted**  
*instructor as T*  $\equiv$  *instructor T*



- SQL提供了一个重命名结果关系中属性的方法
  - from子句的两个关系中可能存在同名属性，在这种情况下，结果中就会出现重复的属性名；
  - 其次，如果我们在select子句中使用算术表达式，那么结果属性就没有名字；
  - 再次，属性名可以从基关系导出，但我们也许想要改变结果中的属性名字。



- 在from子句中，**重命名关系**的原因
  - 是把一个**长的关系名**替换成**短的**，这样，在查询的其他地方使用起来就更为**方便**
  - 另一个原因是为了适用于需要**比较同一个关系**中的元组的情况。
- 像T和S那样，被用来**重命名关系的标识符**在SQL标准中被称作**相关名称(correlation name)**，**表别名(table alias)**，或者**相关变量(correlation variable)**，或者**元组变量(tuple variable)**。



## 3.4.2 String Operations

- SQL使用一对单引号来标示字符串
- 如果单引号是字符串的组成部分，那就用两个单引号字符来表示
- 在SQL标准中，字符串上的相等运算是区分大小写的
- 然而，一些数据库管理系统，如MySQL和SQL Server，在匹配字符串时并不区分大小写
- 这种默认方式，可以在数据库级或特定属性级被修改的。



- SQL includes a *string-matching operator* for comparisons on character strings.
- The operator **like** uses patterns that are described using two special characters:
  - percent ( % ). The % character matches any substring.
  - underscore ( \_ ). The \_ character matches any character.



- Find the names of all instructors whose name includes the **substring** “dar” .

- select name  
      from instructor  
      where name **like** '%dar%'

- Match the string “100%”

**like** '100 \%' **escape** '\'

in that above we use **backslash** (\) as the **escape** character.

- escape character, **转义字符**直接放在**特殊字符**前面，表示该**特殊字符**被当成**普通字符**。



# String Operations (Cont.)

- Patterns are case sensitive.
- Pattern matching examples:
  - ‘Intro%’ matches any string beginning with “Intro”.
  - ‘%Comp%’ matches any string containing “Comp” as a substring.
  - ‘\_\_\_’ matches any string of exactly three characters.
  - ‘\_\_\_%’ matches any string of at least three characters.



- SQL supports a variety of string operations such as
  - concatenation (using “||”)
  - converting from upper to lower case (and vice versa), **upper (s)** , **lower (s)**
  - 去掉字符串后面的空格(使用**trim(s)**)
  - finding string length,
  - extracting substrings, etc.





### 3.4.3 select子句中的属性说明

- 星号“\*”可以用在select子句中，表示“**所有的属性**”，
- 因而，如下查询的select子句中使用instructor.\*
  - select instructor.\*
  - from instructor, teaches
  - where instructor. ID= teaches.ID;
- 表示instructor中的所有属性都被选中。



## 3.4.4 Ordering the Display of Tuples

- **order by**子句就可以让查询结果中元组**按排列顺序**显示
- List in **alphabetic order the names** of all instructors
- select distinct name  
from instructor  
**order by** name
- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; **ascending** order is the **default**.



- Example:
  - order by name **desc**
- Can sort on **multiple attributes**
- Example:
  - order by dept\_name, name



## 3.4.5 Where Clause Predicates

- SQL includes a **between...and...** comparison operator
- Example:
  - Find the names of all instructors with salary between \$90,000 and \$100,000
  - **select** name  
**from** instructor  
**where** salary **between** 90000 **and** 100000



## ■ Tuple comparison

● **select** name, course\_id  
**from** instructor, teaches  
**where** (instructor.ID, dept\_name) =  
(teaches.ID, 'Biology');

- SQL允许我们用记号( $v_1, v_2, \dots, v_n$ )来表示一个分量值分别为  $v_1, v_2, \dots, v_n$  的  $n$  维元组。
- 在元组上可以运用比较运算符，按字典顺序进行比较运算。



# Duplicates

- In relations with **duplicates**, SQL can **define** how many **copies** of tuples appear in the **result**.
- **Multiset (多重集合)** versions of some of the relational algebra operators – given multiset relations  $r1$  and  $r2$ :
  - 1.  $\sigma_{\theta}(r1)$ : If there are  **$c1$  copies** of tuple  $t1$  in  $r1$ , and  $t1$  **satisfies** selections  $\sigma_{\theta}$ , then there are  $c1$  copies of  $t1$  in  $\sigma_{\theta}(r1)$ .
  - 2.  $\Pi_A(r1)$ : For each copy of tuple  $t1$  in  $r1$ , there is a copy of tuple  $\Pi_A(t1)$  in  $\Pi_A(r1)$ , where  $\Pi_A(t1)$  denotes the projection of the single tuple  $t1$ .



- 3.  $r_1 \times r_2$ : If there are **c1 copies** of tuple  $t_1$  in  $r_1$  and **c2 copies** of tuple  $t_2$  in  $r_2$ , there are  $c_1 \times c_2$  copies of the tuple  $t_1 \cdot t_2$  in  $r_1 \times r_2$

- Example: Suppose multiset relations  $r_1 (A, B)$  and  $r_2 (C)$  are as follows:

$$r_1 = \{(1, a) (2, a)\} \quad r_2 = \{(2), (3), (3)\}$$

- Then  $\Pi_B(r_1)$  would be  $\{(a), (a)\}$ , while  $\Pi_B(r_1) \times r_2$  would be

$$\{(a, 2), (a, 2), (a, 3), (a, 3), (a, 3), (a, 3)\}$$



## Duplicates (Cont.)

- SQL duplicate semantics:

**select**  $A_1, A_2, \dots, A_n$   
**from**  $r_1, r_2, \dots, r_m$   
**where**  $P$

is equivalent to the *multiset* version of the expression:

$$\prod_{A_1, A_2, \dots, A_n} (\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$





## 3.5 Set Operations

### 3.5.1~3.5.3 并、交、差运算

- SQL作用在关系上的`union`、`intersect`和`except`运算对应于数学集合论中的 $\cup$ ,  $\cap$ , 和 $-$ 运算。
- Find courses that ran in Fall 2009 or in Spring 2010
- `(select course_id from section where sem = 'Fall' and year = 2009)`  
`union`  
`(select course_id from section where sem = 'Spring' and year = 2010)`



- Find courses that ran in Fall 2009 and in Spring 2010
- **(select *course\_id* from *section* where *sem* = 'Fall' and *year* = 2009)**  
**intersect**  
**(select *course\_id* from *section* where *sem* = 'Spring' and *year* = 2010)**



- Find courses that ran in Fall 2009 but not in Spring 2010
- **(select *course\_id* from *section* where *sem* = 'Fall' and *year* = 2009)**  
**except**  
**(select *course\_id* from *section* where *sem* = 'Spring' and *year* = 2010)**
- **union**、**intersect**和**except**运算自动去除重复
- 如果想保留所有重复，就必须用**union all**、**intersect all**和**except all**代替**union**、**intersect**和**except**



## Set Operations (Cont.)

- Suppose a tuple occurs **m times** in r and **n times** in s, then, it occurs:
  - **m + n** times in r union all s
  - **min(m,n)** times in r intersect all s
  - **max(0, m - n)** times in r except all s



## 3.6 Null Values

- It is possible for **tuples** to have a **null value**, denoted by *null*, for some of their attributes
- *null* signifies an **unknown** value or that a value does **not exist**.
- The result of any **arithmetic expression** involving *null* is *null*
  - Example:  $5 + \textit{null}$  returns **null**



- The **predicate is null** can be used to check for **null** values.

- Example: Find all instructors whose salary is **null**.

```
select name  
from instructor  
where salary is null
```

- SQL将涉及**空值**的任何**比较运算**的结果视为**unknown**

- 既不是谓词**is null**，也不是**is not null**



# Null Values and Three Valued Logic

- **Three values** – *true*, *false*, ***unknown***
- Any comparison with *null* returns ***unknown***
  - Example:  $5 < null$  or  $null <> null$  or  $null = null$
- Three-valued logic using the value ***unknown***:
  - OR:  $(unknown \text{ or } true) = true$ ,  
 $(unknown \text{ or } false) = unknown$   
 $(unknown \text{ or } unknown) = unknown$
  - AND:  $(true \text{ and } unknown) = unknown$ ,  
 $(false \text{ and } unknown) = false$ ,  
 $(unknown \text{ and } unknown) = unknown$



- NOT: (**not** *unknown*) = *unknown*
- “*P* **is unknown**” evaluates to **true** if predicate *P* evaluates to *unknown*
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*
- 某些SQL，还允许我们使用子句**is unknown**和**is not unknown**来测试一个表达式的结果是否为**unknown**，而不是**true**或**false**。





- 当一个查询使用 **select distinct** 子句时，**重复元组** 将被去除。
- 为了达到这个目的，当比较两个元组对应的属性值时，
  - 如果这两个值都是**非空并且值相等**，或者**都是空**，那么它们是**相同的**。
- 这里，对待空值的方式与谓词中对待空值的方式是不同的，
  - 在谓词中“`null = null`”会返回**unknown**，而不是**true**。



- 如果元组在**所有属性**上的**取值**相等，那么它们就被当作**相同元组**，即使某些值为**空**。
- 上述方式还应用于集合的**并、交和差**运算



## 3.7 Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value

**avg:** average value

**min:** minimum value

**max:** maximum value

**sum:** sum of values

**count:** number of values



## 3.7.1 Aggregate Functions (Cont.)

- Find the **average** salary of instructors in the Computer Science department

- ```
select avg (salary)
  from instructor
 where dept_name= 'Comp. Sci.';
```

- Find the **total number** of instructors who teach a course in the Spring 2010 semester

- ```
select count (distinct ID)
  from teaches
 where semester = 'Spring' and year =
2010;
```



- Find the **number of tuples** in the course relation
  - **select count (\*)**  
from course;
- SQL不允许在用count(\*)时使用distinct。
- 在用max和min时使用distinct是合法的，尽管结果并无差别。



## 3.7.2 Aggregate Functions – Group By

- Find the average salary of instructors in each department
  - **select** *dept\_name*, **avg** (*salary*) **as** *avg\_salary*  
**from** *instructor*  
**group by** *dept\_name*;

ID	name	dept_name	salary
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

dept_name	avg_salary
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000



## Aggregation (Cont.)

- **Attributes** in **select** clause outside of aggregate functions must appear in **group by** list
  - **/\* erroneous query \*/**  
**select** *dept\_name*, **ID**, **avg** (*salary*)  
**from** *instructor*  
**group by** *dept\_name*;



## 3.7.3 Having Clause

- Find the names and average salaries of all departments whose average salary is **greater than** 42000
  - **select** *dept\_name*, **avg** (*salary*)
  - **from** *instructor*
  - **group by** *dept\_name*
  - **having** **avg** (*salary*) > 42000;
- Note: **predicates** in the **having** clause are applied **after** the formation of groups, whereas predicates in the **where** clause are applied before forming groups





- 任何出现在**having**子句中，但没有被聚集的属性必须出现在**group by**子句中，否则查询就被当成是错误的。
- 包含**聚集**、**group by**或**having**子句的查询的含义可通过下述操作序列来定义：
  1. 最先根据**from**子句来计算出一个**关系**。
  2. 如果出现了**where**子句，**where**子句中的**谓词**将应用到**from**子句的**结果关系**上。
  3. 如果出现了**group by**子句，满足**where**谓词的元组通过**group by**子句形成**分组**。如果没有**group by**子句，满足**where**谓词的**整个元组集**被当作一个**分组**。



4. 如果出现了**having**子句，它将**应用到每个分组上**；**不满足having子句谓词的分组将被抛弃**。
5. **select子句**利用**剩下的分组**产生出查询**结果**中的元组，即在**每个分组上**应用**聚集函数**来得到**单个结果元组**。

- `select course_id, semester, year, sec_id, avg(tot_cred)`
- `from takes natural join student`
- `where year = 2009`
- `group by course_id, semester, year, sec_id`
- `having count (ID) >= 2;`



## 3.7.4 Null Values and Aggregates

### ■ Total all salaries

```
select sum (salary )  
from instructor
```

- Above statement **ignores** null amounts
  - Result is *null* if there is **no non-null** amount
- ### ■ All **aggregate operations** except **count(\*)** **ignore** tuples with **null values** on the aggregated attributes
- ### ■ What if collection has only null values?
- **count** returns 0
  - **all other aggregates** return **null**



## 3.8 Nested Subqueries

- SQL provides a **mechanism** for the nesting of subqueries.
- A **subquery** is a **select-from-where** expression that is nested within **another** query.
- The nesting can be done in the following SQL query

**select**  $A_1, A_2, \dots, A_n$   
**from**  $r_1, r_2, \dots, r_m$   
**where**  $P$



■ as follows:

- $A_i$  can be **replaced** be a subquery that generates a **single** value.
- $r_i$  can be replaced by **any valid subquery**
- $P$  can be replaced with an **expression** of the form:

$B$  <operation> (**subquery**)

Where  $B$  is an **attribute** and <operation> to be defined later.



# Subqueries in the Where Clause

- A common use of **subqueries** is to perform tests:
  - For **set membership**
  - For **set comparisons**
  - For **set cardinality**.



## 3.8.1 Set Membership

- SQL允许测试元组在关系中的成员资格。
  - 连接词**in**测试元组是否是集合中的**成员**，**集合**是由**select**子句产生的一组值构成的。
  - 连接词**not in**则测试元组是否**不是**集合中的**成员**。
- Find courses offered in Fall 2009 and in Spring 2010
  - **select distinct** *course\_id*
  - **from** *section*
  - **where** *semester* = 'Fall' **and** *year* = 2009 **and**  
*course\_id* **in** (**select** *course\_id*
  - **from** *section*
  - **where** *semester* = 'Spring' **and**  
*year* = 2010);



- Find courses offered in Fall 2009 but not in Spring 2010
  - **select distinct** *course\_id*
  - **from** *section*
  - **where** *semester* = 'Fall' **and** *year* = 2009  
**and** *course\_id* **not in** (**select** *course\_id*
  - **from** *section*
  - **where** *semester* =  
'Spring' **and** *year* = 2010);





## Set Membership (Cont.)

- Find the **total number** of (distinct) students who have taken course sections taught by the instructor with *ID* 10101
- **select count (distinct *ID*)**
- **from *takes***
- **where (*course\_id*, *sec\_id*, *semester*, *year*) in**  
**(select *course\_id*, *sec\_id*, *semester*, *year***
- **from *teaches***
- **where *teaches.ID*= 10101);**
- Note: Above query can be written in a much simpler manner.  
The formulation above is simply to **illustrate** SQL features.



- **in**和**not in**操作符也能用于枚举集合。
- **select distinct name**
- **from instructor**
- **where name not in ('Mozart', 'Einstein');**



## 3.8.2 Set Comparison – “some” Clause

- Find names of instructors with salary **greater than** that of some (at least one) instructor in the Biology department.
  - **select distinct** *T.name*
  - **from** *instructor* **as** *T*, *instructor* **as** *S*
  - **where** *T.salary* > *S.salary* **and** *S.dept\_name* = 'Biology';
- Same query using **> some** clause
  - **select** *name*
  - **from** *instructor*
  - **where** *salary* > **some** (**select** *salary*
  - *from* *instructor*
  - **where** *dept name* = 'Biology');



# Definition of “some” Clause

■  $F \langle \text{comp} \rangle \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F \langle \text{comp} \rangle t)$

Where  $\langle \text{comp} \rangle$  can be:  $<, \leq, >, =, \neq$

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$

(read: 5 < some tuple in the relation)

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$

$(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \neq 5)$

$(= \text{some}) \equiv \text{in}$

However,  $(\neq \text{some}) \not\equiv \text{not in}$



# Set Comparison – “all” Clause

- Find the names of all instructors whose salary is **greater than** the salary of **all instructors** in the Biology department.

- **select** *name*
- **from** *instructor*
- **where** *salary* > **all** (**select** *salary*
- *from* *instructor*
- **where** *dept name* =  
'Biology');



# Definition of “all” Clause

■  $F \text{ <comp> all } r \Leftrightarrow \forall t \in r (F \text{ <comp> } t)$

$(5 < \text{all } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$

$(5 < \text{all } \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$

$(5 = \text{all } \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 \neq \text{all } \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6)$

$(\neq \text{ all}) \equiv \text{not in}$

However,  $(= \text{ all}) \equiv \text{in}$



### 3.8.3 Test for Empty Relations

- The **exists** construct returns the value **true** if the argument subquery is **nonempty**.
- **exists**  $r \Leftrightarrow r \neq \emptyset$
- **not exists**  $r \Leftrightarrow r = \emptyset$



## Use of “exists” Clause

- Yet another way of specifying the query “Find all courses taught in **both** the Fall 2009 semester and in the Spring 2010 semester”

```
select course_id
from section as S
where semester = 'Fall' and year = 2009 and
      exists (select *
               from section as T
               where semester = 'Spring' and
year = 2010
               and S.course_id =
T.course_id);
```





- SQL的一个特性，来自外层查询的一个相关名称（上述查询中的S）可以用在where子句的子查询中。
- 使用了来自外层查询相关名称的子查询被称作相关子查询(correlated subquery)。
- **Correlation name** – variable S in the outer query
- **Correlated subquery** – the inner query



# Use of “not exists” Clause

- Find all students who have taken all courses offered in the Biology department.
  - **select distinct** *S.ID, S.name*
  - **from** *student as S*
  - **where not exists** ( (select *course\_id*
  - *from course*
  - *where dept\_name = 'Biology'*)
  - **except**
  - (select *T.course\_id*
  - *from takes as T*
  - *where S.ID = T.ID*));



- First nested query lists all courses offered in Biology
- Second nested query lists all courses a particular student took
- Note that  $X - Y = \emptyset \Leftrightarrow X \subseteq Y$
- Note: *Cannot* write this query using **= all** and its variants



## 3.8.4 Test for Absence of Duplicate Tuples

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
  - “true” , if a given subquery contains no duplicates .
- For example:
- Find all courses that were offered **at most once** in 2009



■ **select** *T.course\_id*  
**from** *course* **as** *T*  
**where unique** (**select** *R.course\_id*  
                  **from** *section* **as** *R*  
                  **where** *T.course\_id*=  
*R.course\_id* **and** *R.year* = 2009);



- 在不使用unique结构的情况下，上述查询的一种等价表达方式是：
- `select T.course id`
- `from course as T`
- `where 1 <= (select count(R.course id)`
  - » `from section as R`
  - » `where T.course id= R.course id and`
  - » `R.year = 2009);`



- **not unique**结构测试在一个子查询结果中是否存在重复元组。
- 对一个关系的**unique**测试结果为**假**的定义是：
  - 当且仅当在关系中存在着两个元组 $t_1$ 和 $t_2$ ，且 $t_1=t_2$ 。
  - 由于在 $t_1$ 或 $t_2$ 的**某个域为空**时，判断 $t_1=t_2$ 为假
- 所以，尽管一个元组有**多个副本**，只要该元组有一个属性**为空**，unique测试就有可能**为真**。



## 3.8.5 Subqueries in the Form Clause

- SQL allows a **subquery** expression to be used in the **from** clause
- 任何select-from-where表达式返回的结果都是**关系**，
- 因而**该关系**可以被**插入**到另一个select- from-where中**任何**关系**可以出现**的位置。





- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000.”

```
select dept_name, avg_salary
from (select dept_name, avg (salary) as
avg_salary
      from instructor
      group by dept_name)
where avg_salary > 42000;
```



- Note that we do **not need** to use the **having** clause
- Another way to write above query

```
select dept_name, avg_salary
from (select dept_name, avg (salary)
      from instructor
      group by dept_name) as dept_avg
(dept_name, avg_salary)
where avg_salary > 42000;
```

- 子查询的**结果关系**被命名为dept\_avg，其属性名是dept\_name和avg\_salary。



- 很多（但并非全部）SQL实现都支持在from子句中嵌套子查询。
- 注意，某些SQL实现要求对每一个子查询结果关系都给一个名字，即使该名字从不被引用；
- Oracle允许对子查询结果关系命名（省略掉关键字as），但是不允许对关系中的属性重命名。



- 在from子句嵌套的子查询中，不能使用来自from子句其他关系的相关变量。
- 然而SQL:2003，允许from子句中的子查询用**关键词lateral**作为前缀，以便访问from子句中在它前面的表或子查询中的属性。
- 例如：
  - select name, salary, avg salary
  - from instructor *I1*, **lateral** (select avg(salary) as avg salary
    - » from instructor I2
    - » where I2.dept name= *I1*.dept name);



## 3.8.6 With Clause

- The **with** clause provides a way of defining a **temporary** relation whose definition is **available only** to the **query** in which the **with** clause **occurs**.
- Find all departments with the maximum budget

```
with max_budget (value) as  
    (select max(budget)  
     from department)  
select department.name  
from department, max_budget  
where department.budget = max_budget.value;
```



# Complex Queries using With Clause

- Find all departments where the total salary is greater than the average of the total salary at all departments
- **with** *dept\_total* (*dept\_name*, *value*) **as**
- (**select** *dept\_name*, **sum**(*salary*)
- **from** *instructor*
- **group by** *dept\_name*),
- ▶ *dept\_total\_avg*(*value*) **as**
- (**select** **avg**(*value*)
- **from** *dept\_total*)
- **select** *dept\_name*
- **from** *dept\_total*, *dept\_total\_avg*
- **where** *dept\_total.value* > *dept\_total\_avg.value*;



### 3.8.7 Scalar Subquery

- **Scalar subquery** is one which is used where a **single value** is expected
- List all departments along with the number of instructors in each department

```
select dept_name,  
       (select count(*)  
        from instructor  
        where department.dept_name =  
        instructor.dept_name)  
       as num_instructors  
from department;
```

- Runtime **error** if **subquery** returns **more than one** result tuple



## 3.9 Modification of the Database

- **Deletion** of **tuples** from a given relation.
- **Insertion** of new **tuples** into a given relation
- **Updating** of **values** in some tuples in a given relation





## 3.9.1 Deletion

- 删除整个元组，而不能只删除某些属性上的值
  - **delete from** r
  - **whert** P;
- 其中P代表一个谓词，r代表一个关系。
- **delete**语句，
  - 首先从r中找出所有使P(t)为真的元组，
  - 然后，把它们从r中删除。



- Delete **all** instructors

**delete from** *instructor*

- Delete all instructors from the **Finance** department

**delete from** *instructor*  
**where** *dept\_name*= 'Finance';



- Delete **all tuples** in the *instructor* relation for those instructors associated with a department located in the **Watson** building.
- **delete from** instructor
- **where** dept\_name **in** (select dept name  
from *department*  
where *building* = 'Watson');



## Deletion (Cont.)

- Delete all instructors whose salary is less than the average salary of instructors
  - ▶ **delete from** *instructor*
  - ▶ **where** *salary* < (**select avg** (*salary*)
  - ▶ **from** *instructor*);
- Problem: as we **delete** tuples from deposit, the **average** salary **changes**
- **Solution** used in SQL:
  1. **First**, compute **avg** (*salary*) and **find all** tuples to delete
  2. Next, **delete** all tuples found above (***without*** ***recomputing*** **avg** or retesting the tuples)



## 3.9.2 Insertion

- Add a new tuple to *course*

**insert into** *course*

**values** ('CS-437', 'Database Systems',  
'Comp. Sci.', 4);

- or equivalently

**insert into** *course* (*course\_id*, *title*, *dept\_name*,  
*credits*)

**values** ('CS-437', 'Database Systems',  
'Comp. Sci.', 4);



- Add a new tuple to *student* with *tot\_creds* set to *null*

```
insert into student  
    values ('3003', 'Green', 'Finance', null);
```

- Add **all instructors** to the *student* relation with *tot\_creds* set to 0

```
insert into student  
    select ID, name, dept_name, 0  
    from instructor
```



## Insertion (Cont.)

- The **select from where** statement is evaluated fully **before** any of its results **are inserted** into the relation.

Otherwise queries like

**insert into *table1* select \* from *table1***  
would cause **problem**



## 3.9.3 Updates

- 不改变整个元组的情况下，改变其部分属性的值
  - **update** r
  - **set** A=new value
  - **where** P;
- **update**语句中可以嵌套的**select**可以引用待更新的关系。
- 同样，SQL首先检查关系中的所有元组，看它们是否应该被更新，然后才执行更新。





- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%

- Write *two* **update** statements:

**update** *instructor*

**set** *salary* = *salary* \* 1.03

**where** *salary* > 100000;

**update** *instructor*

**set** *salary* = *salary* \* 1.05

**where** *salary* <= 100000;

- The *order* is important
- Can be done better using the **case statement**.



# Case Statement for Conditional Updates

- Same query as before but with **case statement**

```
update instructor  
  set salary = case  
    when salary <= 100000 then  
      salary * 1.05  
    else salary * 1.03  
  end
```



## ■ **case**语句的一般格式如下：

### ● **case**

- ▶ **when** pred1 **then** result1
- ▶ **when** pred2 **then** result2
- ▶ ...
- ▶ **when** predn **then** resultn
- ▶ **else** result0

### ● **end**



# Updates with Scalar Subqueries

- **Recompute** and **update** tot\_creds value for all students

**update** *student S*

**set** *tot\_cred* = (select sum(*credits*)

**from** *takes* **natural join**

*course*

**where** *S.ID* = *takes.ID* **and**

*takes.grade* <> 'F' **and** *takes.grade* **is not null**);



- Sets *tot\_creds* to **null** for students who have *not taken any* course

- Instead of **sum(credits)**, use:

**update** *student S*

**set** *tot\_cred* = **case**

**when** **sum(credits)** **is not null** **then**

**sum(credits)**

**else 0**

**end**



# End of Chapter 3

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan  
See [www.db-book.com](http://www.db-book.com) for conditions on re-use