

# Lecture 9: Integrating Learning and Planning

Joseph Modayil

# Outline

- 1 Introduction
- 2 Model-Based Reinforcement Learning
- 3 Model-Based Reinforcement Learning
- 4 Integrated Architectures
- 5 Simulation-Based Search

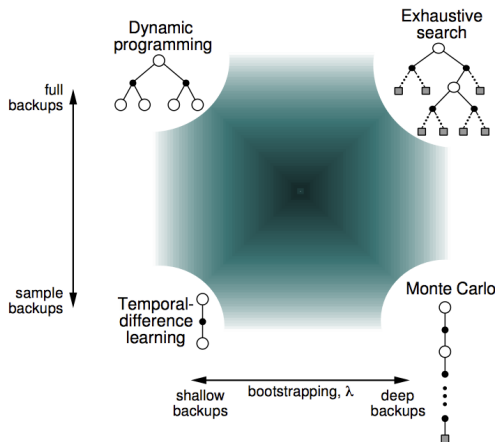
# Model-Based Reinforcement Learning

- *Last lecture*: learn **policy** directly from experience
- *Previous lectures*: learn **value function** directly from experience
- *This lecture*:
  - Learn **model** directly from experience (or be given a model)
  - **Plan** with the model to construct a value function or policy
  - Integrate learning and planning into a single architecture


# Model-Based and Model-Free RL

- Model-Free RL
  - No model
  - **Learn** value function (and/or policy) from experience
- Model-Based RL
  - **Learn** a model from experience OR be given a model
  - **Plan** value function (and/or policy) from model

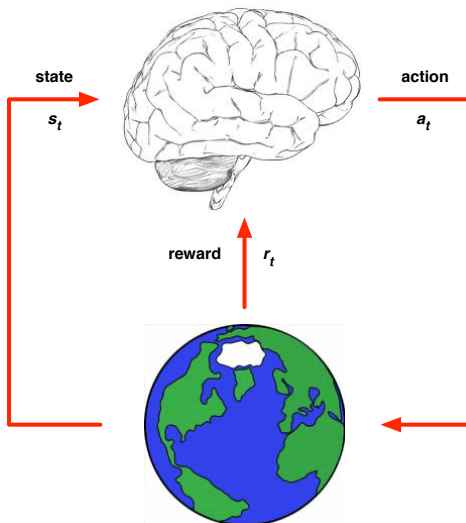
# Filling in the middle of algorithm space




# Model-Free RL



# Model-Based RL

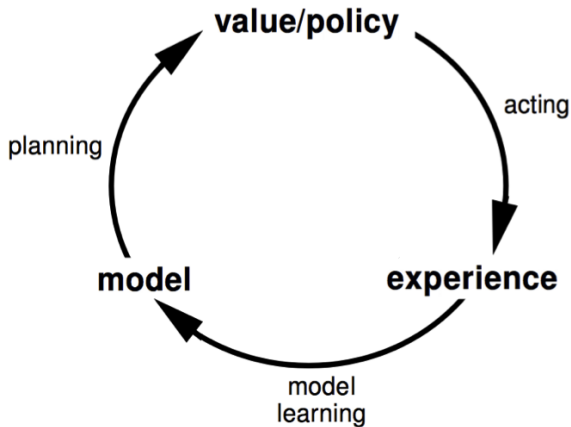


## Model-Based RL





## Model-Based RL



# What is a Model?

- A *model*  $\mathcal{M}$  is a representation of an MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R} \rangle$ , parametrized by  $\eta$
- A model  $\mathcal{M} = \langle \mathcal{P}_\eta, \mathcal{R}_\eta \rangle$  approximates the state transitions  $\mathcal{P}_\eta \approx \mathcal{P}$  and rewards  $\mathcal{R}_\eta \approx \mathcal{R}$ . e.g.

$$S_{t+1} \sim \mathcal{P}_\eta(S_{t+1} \mid S_t, A_t)$$

$$R_{t+1} = \mathcal{R}_\eta(R_{t+1} \mid S_t, A_t)$$

This particular model imposes conditional independence between state transitions and rewards

$$\mathbb{P}[S_{t+1}, R_{t+1} \mid S_t, A_t] = \mathbb{P}[S_{t+1} \mid S_t, A_t] \mathbb{P}[R_{t+1} \mid S_t, A_t]$$

- Conventionally a method is called *model-based* when the transition and reward dynamics are explicitly represented (to support planning), and as *model-free* otherwise. Some new methods lie in-between these extremes.

# Model Learning

- Goal: estimate model  $\mathcal{M}_\eta$  from experience  $\{S_1, A_1, R_2, \dots, S_T\}$
- This is a supervised learning problem

$$S_1, A_1 \rightarrow R_2, S_2$$

$$S_2, A_2 \rightarrow R_3, S_3$$

$$\vdots$$

$$S_{T-1}, A_{T-1} \rightarrow R_T, S_T$$

- Learn a function  $s, a \rightarrow r$  and also learn a function  $s, a \rightarrow s'$
- Pick loss function (e.g. mean-squared error), and find parameters  $\eta$  that minimise empirical loss

# Examples of Models

- Table Lookup Model
- Linear Expectation Model
- Linear Gaussian Model
- Gaussian Process Model
- Deep Belief Network Model
- ...

# Table Lookup Model

- Model is an explicit MDP,  $\hat{\mathcal{P}}, \hat{\mathcal{R}}$
- Count visits  $N(s, a)$  to each state action pair

$$\hat{\mathcal{P}}_{s,s'}^a = \frac{1}{N(s, a)} \sum_{t=1}^T \mathbf{1}(S_t, A_t, S_{t+1} = s, a, s')$$

$$\hat{\mathcal{R}}_s^a = \frac{1}{N(s, a)} \sum_{t=1}^T \mathbf{1}(S_t, A_t = s, a) R_t$$

- Alternatively
  - At each time-step  $t$ , record experience tuple  $\langle S_t, A_t, R_{t+1}, S_{t+1} \rangle$
  - To sample model, randomly pick tuple matching  $\langle s, a, \cdot, \cdot \rangle$

# AB Example

Two states  $A, B$ ; no discounting; 8 episodes of experience

$A, 0, B, 0$

$B, 1$

$B, 1$


$B, 1$

$B, 1$

$B, 1$

$B, 1$

$B, 0$



We have constructed a **table lookup model** from the experience

# Planning with a Model

- Given a model  $\mathcal{M}_\eta = \langle \mathcal{P}_\eta, \mathcal{R}_\eta \rangle$
- Solve the MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}_\eta, \mathcal{R}_\eta \rangle$
- Using favourite planning algorithm
  - Value iteration
  - Policy iteration
  - Tree search
  - ...

# Sample-Based Planning

- A simple but powerful approach to planning
- Use the model **only** to generate samples
- **Sample** experience from model

$$s_{t+1} \sim \mathcal{P}_\eta(s_{t+1} \mid s_t, a_t)$$

$$r_{t+1} = \mathcal{R}_\eta(r_{t+1} \mid s_t, a_t)$$

- Apply **model-free** RL to samples, e.g.:
  - Monte-Carlo control
  - Sarsa
  - Q-learning
- Sample-based planning methods are often more efficient




# Back to the AB Example

- Construct a table-lookup model from real experience
- Apply model-free RL to sampled experience

Real experience

A, 0, B, 0  
 B, 1  
 B, 1  
 B, 1  
 B, 1  
 B, 1  
 B, 1  
 B, 1  
 B, 0



Sampled experience

B, 1  
 B, 0  
 B, 1  
 A, 0, B, 1  
 B, 1  
 A, 0, B, 1  
 B, 1  
 B, 0

e.g. Monte-Carlo learning:  $V(A) = 1, V(B) = 0.75$

# Planning with an Inaccurate Model

- Given an imperfect model  $\langle \mathcal{P}_\eta, \mathcal{R}_\eta \rangle \neq \langle \mathcal{P}, \mathcal{R} \rangle$
- Performance of model-based RL is limited to optimal policy for approximate MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}_\eta, \mathcal{R}_\eta \rangle$
- i.e. Model-based RL is only as good as the estimated model
- When the model is inaccurate, planning process will compute a suboptimal policy
- Approach 1: when model is wrong, use model-free RL
- Approach 2: reason explicitly about model uncertainty over  $\eta$  (e.g. Bayesian methods)
- Approach 3: Combine model-based and model-free methods in a safe way.

# Real and Simulated Experience

We consider two sources of experience

**Real experience** Sampled from environment (true MDP)

$$s' \sim \mathcal{P}_{s,s'}^a$$

$$r = \mathcal{R}_s^a$$

**Simulated experience** Sampled from model (approximate MDP)


$$s' \sim \mathcal{P}_\eta(s' \mid s, a)$$

$$r = \mathcal{R}_\eta(r \mid s, a)$$

# Integrating Learning and Planning

- Model-Free RL
  - No model
  - **Learn** value function (and/or policy) from real experience
- Model-Based RL (using Sample-Based Planning)
  - Learn a model from real experience
  - **Plan** value function (and/or policy) from simulated experience
- Dyna
  - Learn a model from real experience
  - **Learn AND plan** value function (and/or policy) from real and simulated experience
  - Treat real and simulated experience equivalently. Conceptually, the updates from learning or planning are not distinguished.

# Dyna Architecture



# Dyna-Q Algorithm

Initialize  $Q(s, a)$  and  $Model(s, a)$  for all  $s \in \mathcal{S}$  and  $a \in \mathcal{A}(s)$

Do forever:

(a)  $s \leftarrow$  current (nonterminal) state

(b)  $a \leftarrow \varepsilon$ -greedy( $s, Q$ )

(c) Execute action  $a$ ; observe resultant state,  $s'$ , and reward,  $r$

(d)  $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$

(e)  $Model(s, a) \leftarrow s', r$  (assuming deterministic environment)

(f) Repeat  $N$  times:


$s \leftarrow$  random previously observed state

$a \leftarrow$  random action previously taken in  $s$

$s', r \leftarrow Model(s, a)$


$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$

# Dyna-Q on a Simple Maze



# Dyna-Q with an Inaccurate Model


- The changed environment is **harder**





# Dyna-Q with an Inaccurate Model (2)

- The changed environment is **easier**



# Linear Dyna

- What can we do when the actual states are not known?
- Suppose we have features vectors for states  $\phi(s)$ .
- Consider the expectation models of the next reward and of the next feature vector.

$$\hat{\mathcal{R}}_s^a = \mathbb{E}_\pi [R_{t+1} | S_t = s, A_t = a]$$

$$\hat{\mathcal{F}}_s^a = \mathbb{E}_\pi [\phi(S_{t+1}) | S_t = s, A_t = a]$$

- We can make linear approximations to the expectation models.
- Equivalent fixed points with linear function approximation when learning model-free or planning with the linear expectation model.

# Linear Dyna Model Updates

- Consider a single transition  $\langle \phi(s), a, r, \phi(s') \rangle$
- Learn vector  $b_a$  for a reward model.  $b_a^\top \phi(s) \approx \hat{\mathcal{R}}_s^a$

$$b_a \leftarrow b_a + \alpha(r - b_a^\top \phi(s))\phi(s)$$

- Learn matrix  $F_a$  for a transition model.  $F_a^\top \phi(s) \approx \hat{\mathcal{F}}_s^a$

$$F_a \leftarrow F_a + \alpha(\phi(s') - F_a \phi(s))\phi(s)^\top$$

# Linear Dyna Value Function Updates

- We consider a linear approximation for the action value function with  $\theta_a^\top \phi(s) \approx q(s, a)$
- *Learning* uses an observed transition  $\langle \phi(s), a, r, \phi(s') \rangle$ ,

$$\delta \leftarrow \max_{a'} r + \gamma \theta_{a'}^\top \phi(s') - \theta_a^\top \phi(s)$$

$$\theta_a \leftarrow \theta_a + \alpha \delta \phi(s)$$

- *Planning* uses the model for an expected transition from an arbitrary feature vector  $\psi$ ,  $\langle \psi, a, b_a^\top \psi, F_a \psi \rangle$

$$\delta \leftarrow \max_{a'} b_a^\top \psi + \gamma \theta_{a'}^\top F_a \psi - \theta_a^\top \psi$$

$$\theta_a \leftarrow \theta_a + \alpha \delta \psi$$

# Linear Dyna Algorithm

- We consider a linear approximation for the action value function with  $\theta_a^\top \phi(s) \approx q(s, a)$
- *Learning* uses an observed transition  $\langle \phi(s), a, r, \phi(s') \rangle$ ,

$$\delta \leftarrow \max_{a'} r + \gamma \theta_{a'}^\top \phi(s') - \theta_a^\top \phi(s)$$

$$\theta_a \leftarrow \theta_a + \alpha \delta \phi(s)$$

- *Planning* uses the model for an expected transition from an arbitrary feature vector  $\psi$ ,  $\langle \psi, a, b_a^\top \psi, F_a \psi \rangle$


$$\delta \leftarrow \max_{a'} b_a^\top \psi + \gamma \theta_{a'}^\top F_a \psi - \theta_a^\top \psi$$

$$\theta_a \leftarrow \theta_a + \alpha \delta \psi$$

- We have been learning a model and planning with it.
- Now consider that setting where the model is given (fixed), and we want to use it.

# Forward Search


- **Forward search** algorithms select the best action by **lookahead**
- They build a **search tree** with the current state  $s_t$  at the root
- Using a **model** of the MDP to look ahead



- No need to solve whole MDP, just sub-MDP starting from **now**

# Simulation-Based Search

- **Forward search** paradigm using sample-based planning
- **Simulate** episodes of experience from **now** with the model
- Apply **model-free** RL to simulated episodes





## Simulation-Based Search (2)

- **Simulate** episodes of experience from **now** with the model

$$\{S_t^k, A_t^k, R_{t+1}^k, \dots, S_T^k\}_{k=1}^K \sim \mathcal{M}_\nu$$

- Apply **model-free** RL to simulated episodes
  - Monte-Carlo control  $\rightarrow$  Monte-Carlo search
  - Sarsa  $\rightarrow$  TD search

# Monte-Carlo Simulation

- Given a parameterized model  $\mathcal{M}_\nu$  and a **simulation policy**  $\pi$
- Simulate  $K$  episodes from current state  $S_t$

$$\{\textcolor{red}{S}_t^k = S_t, A_t^k, R_{t+1}^k, S_{t+1}^k, \dots, S_T^k\}_{k=1}^K \sim \mathcal{M}_\nu, \pi$$

- Evaluate state by mean return (**Monte-Carlo evaluation**)

$$V(\textcolor{red}{S}_t) = \frac{1}{K} \sum_{k=1}^K G_t^k \overset{P}{\rightsquigarrow} V^\pi(S_t)$$

# Monte-Carlo Tree Search (Evaluation)

- Given a model  $\mathcal{M}_\nu$
- Simulate  $K$  episodes from current state  $S_t$  using current simulation policy  $\pi$

$$\{\textcolor{red}{S}_t^k = S_t, A_t^k, R_{t+1}^k, S_{t+1}^k, \dots, S_T^k\}_{k=1}^K \sim \mathcal{M}_\nu, \pi$$

- Build a search tree containing visited states and actions
- **Evaluate** states  $Q(s, a)$  by mean return of episodes from  $s, a$

$$Q(s, a) = \frac{1}{N(s, a)} \sum_{k=1}^K \sum_{u=t}^T \mathbf{1}(S_u^k, A_u^k = s, a) G_u \overset{P}{\rightsquigarrow} Q^\pi(s, a)$$

- After search is finished, select current (real) action with maximum value in search tree

$$a_t = \operatorname{argmax}_{a \in \mathcal{A}} Q(S_t, a)$$

# Monte-Carlo Tree Search (Simulation)

- In MCTS, the simulation policy  $\pi$  **improves**
- The simulation policy  $\pi$  has two phases (in-tree, out-of-tree)
  - **Tree policy** (improves): pick actions from  $Q(s, a)$   
(e.g.  $\epsilon - \text{greedy}(Q(s, a))$ )
  - **Default policy** (fixed): pick actions randomly
- Repeat (for each simulated episode)
  - **Select** actions in tree according to tree policy.
  - **Expand** search tree by one node
  - **Rollout** to termination with default policy
  - **Update** action-values  $Q(s, a)$  in the tree
- Output best action when simulation time runs out.
- With some assumptions, converges to the optimal values,  
 $Q(s, a) \rightarrow Q^*(s, a)$


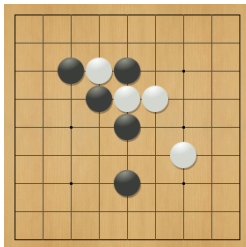
## Case Study: the Game of Go

- The ancient oriental game of Go is 2500 years old
- Considered to be the hardest classic board game
- Considered a grand challenge task for AI (*John McCarthy*)
- Traditional game-tree search has failed in Go



# Rules of Go

- Usually played on 19x19, also 13x13 or 9x9 board
- Simple rules, complex strategy
- Black and white place down stones alternately
- Surrounded stones are captured and removed
- The player with more territory wins the game



# Position Evaluation in Go

- How good is a position  $s$ ?
- Reward function (undiscounted):

$$R_t = 0 \text{ for all non-terminal steps } t < T$$


$$R_T = \begin{cases} 1 & \text{if Black wins} \\ 0 & \text{if White wins} \end{cases}$$

- Policy  $\pi = \langle \pi_B, \pi_W \rangle$  selects moves for both players
- Value function (how good is position  $s$ ):

$$V^\pi(s) = \mathbb{E}_\pi [R_T \mid s] = \mathbb{P}[\text{Black wins} \mid s]$$


$$V^*(s) = \max_{\pi_B} \min_{\pi_W} V^\pi(s)$$

# Monte-Carlo Evaluation in Go

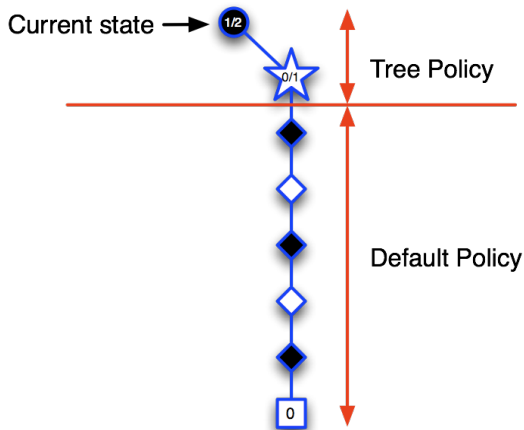




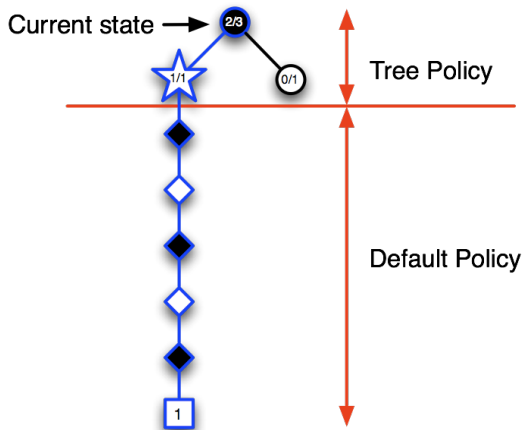
# Applying Monte-Carlo Tree Search (1)




# Applying Monte-Carlo Tree Search (2)



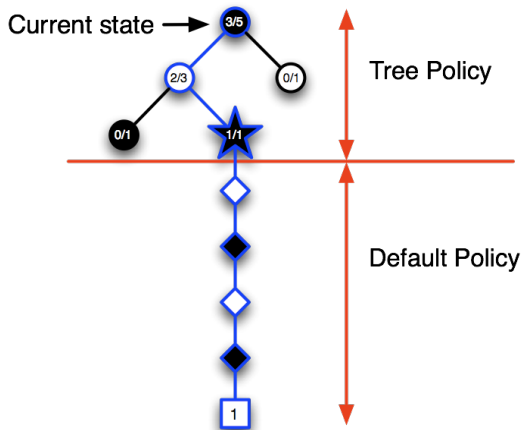
# Applying Monte-Carlo Tree Search (3)



# Applying Monte-Carlo Tree Search (4)



# Applying Monte-Carlo Tree Search (5)



## Advantages of MC Tree Search

- Highly selective best-first search
- Evaluates states *dynamically* (unlike e.g. DP)
- Uses sampling to break curse of dimensionality
- Works for “black-box” models (only requires samples)
- Computationally efficient, anytime, parallelisable

# Example: MC Tree Search in Computer Go

